

# Understanding Linear Regression with a Single Neuron

---



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**Using linear regression for prediction**

**Linear regression using a single neuron**

**Hand-crafting an MSE regression model**

**Hand-crafting a Ridge regression model**

**Comparing to scikit-learn's linear regression estimator**

# Linear Regression

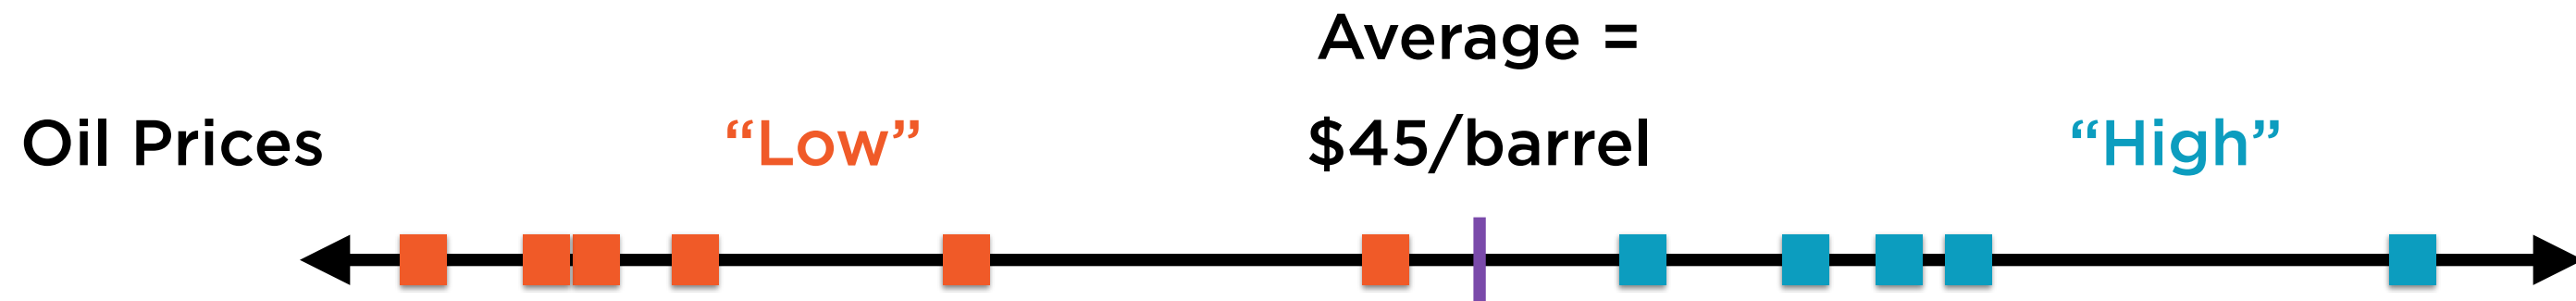
---

# Data in One Dimension



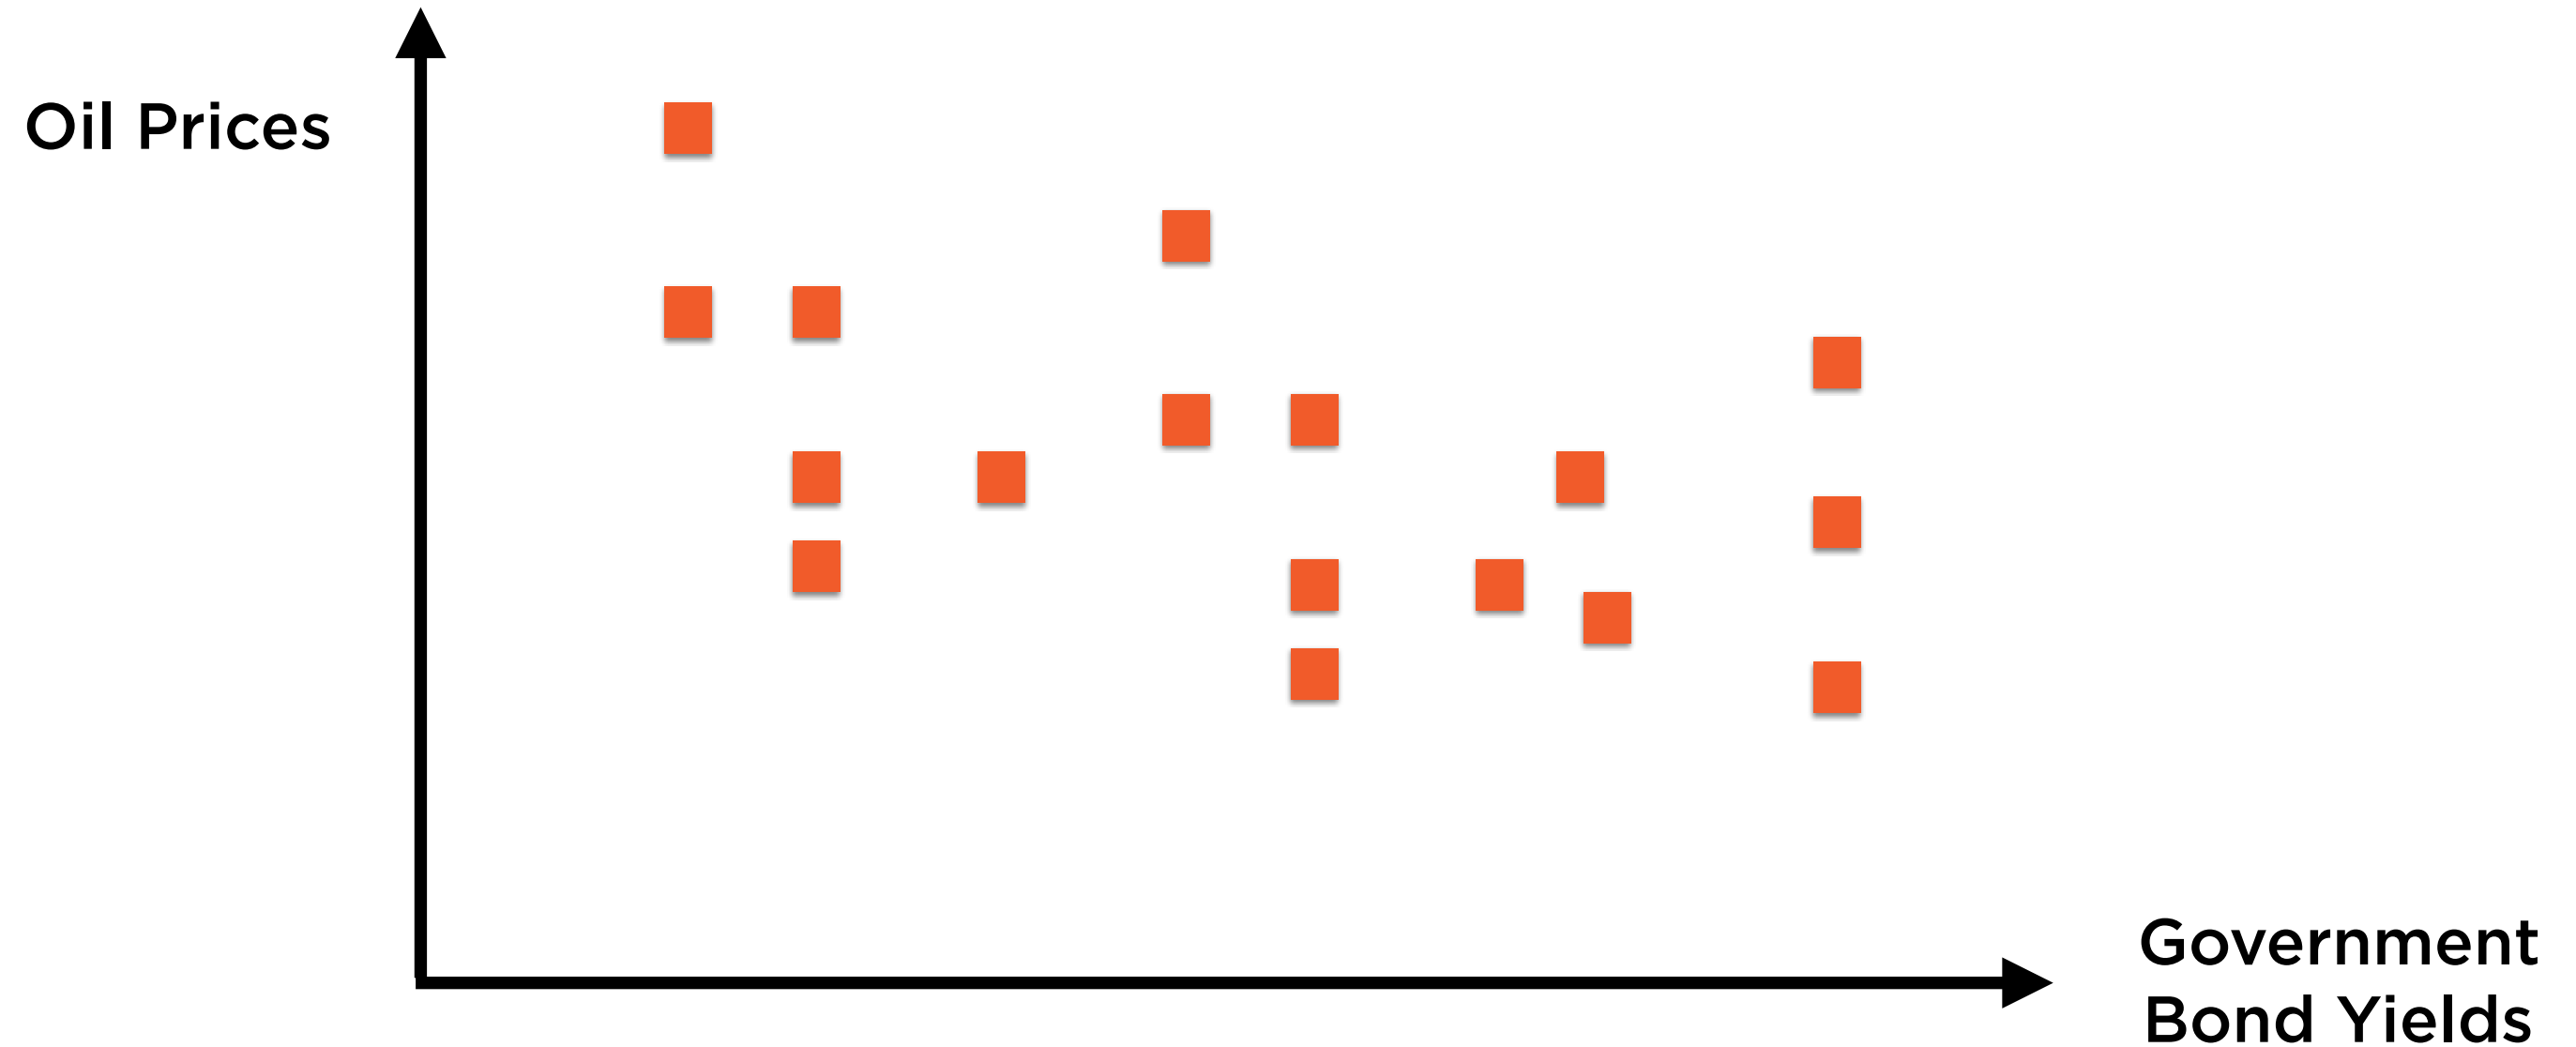
Unidimensional data points can be represented using  
a line, such as a number line

# Data in One Dimension



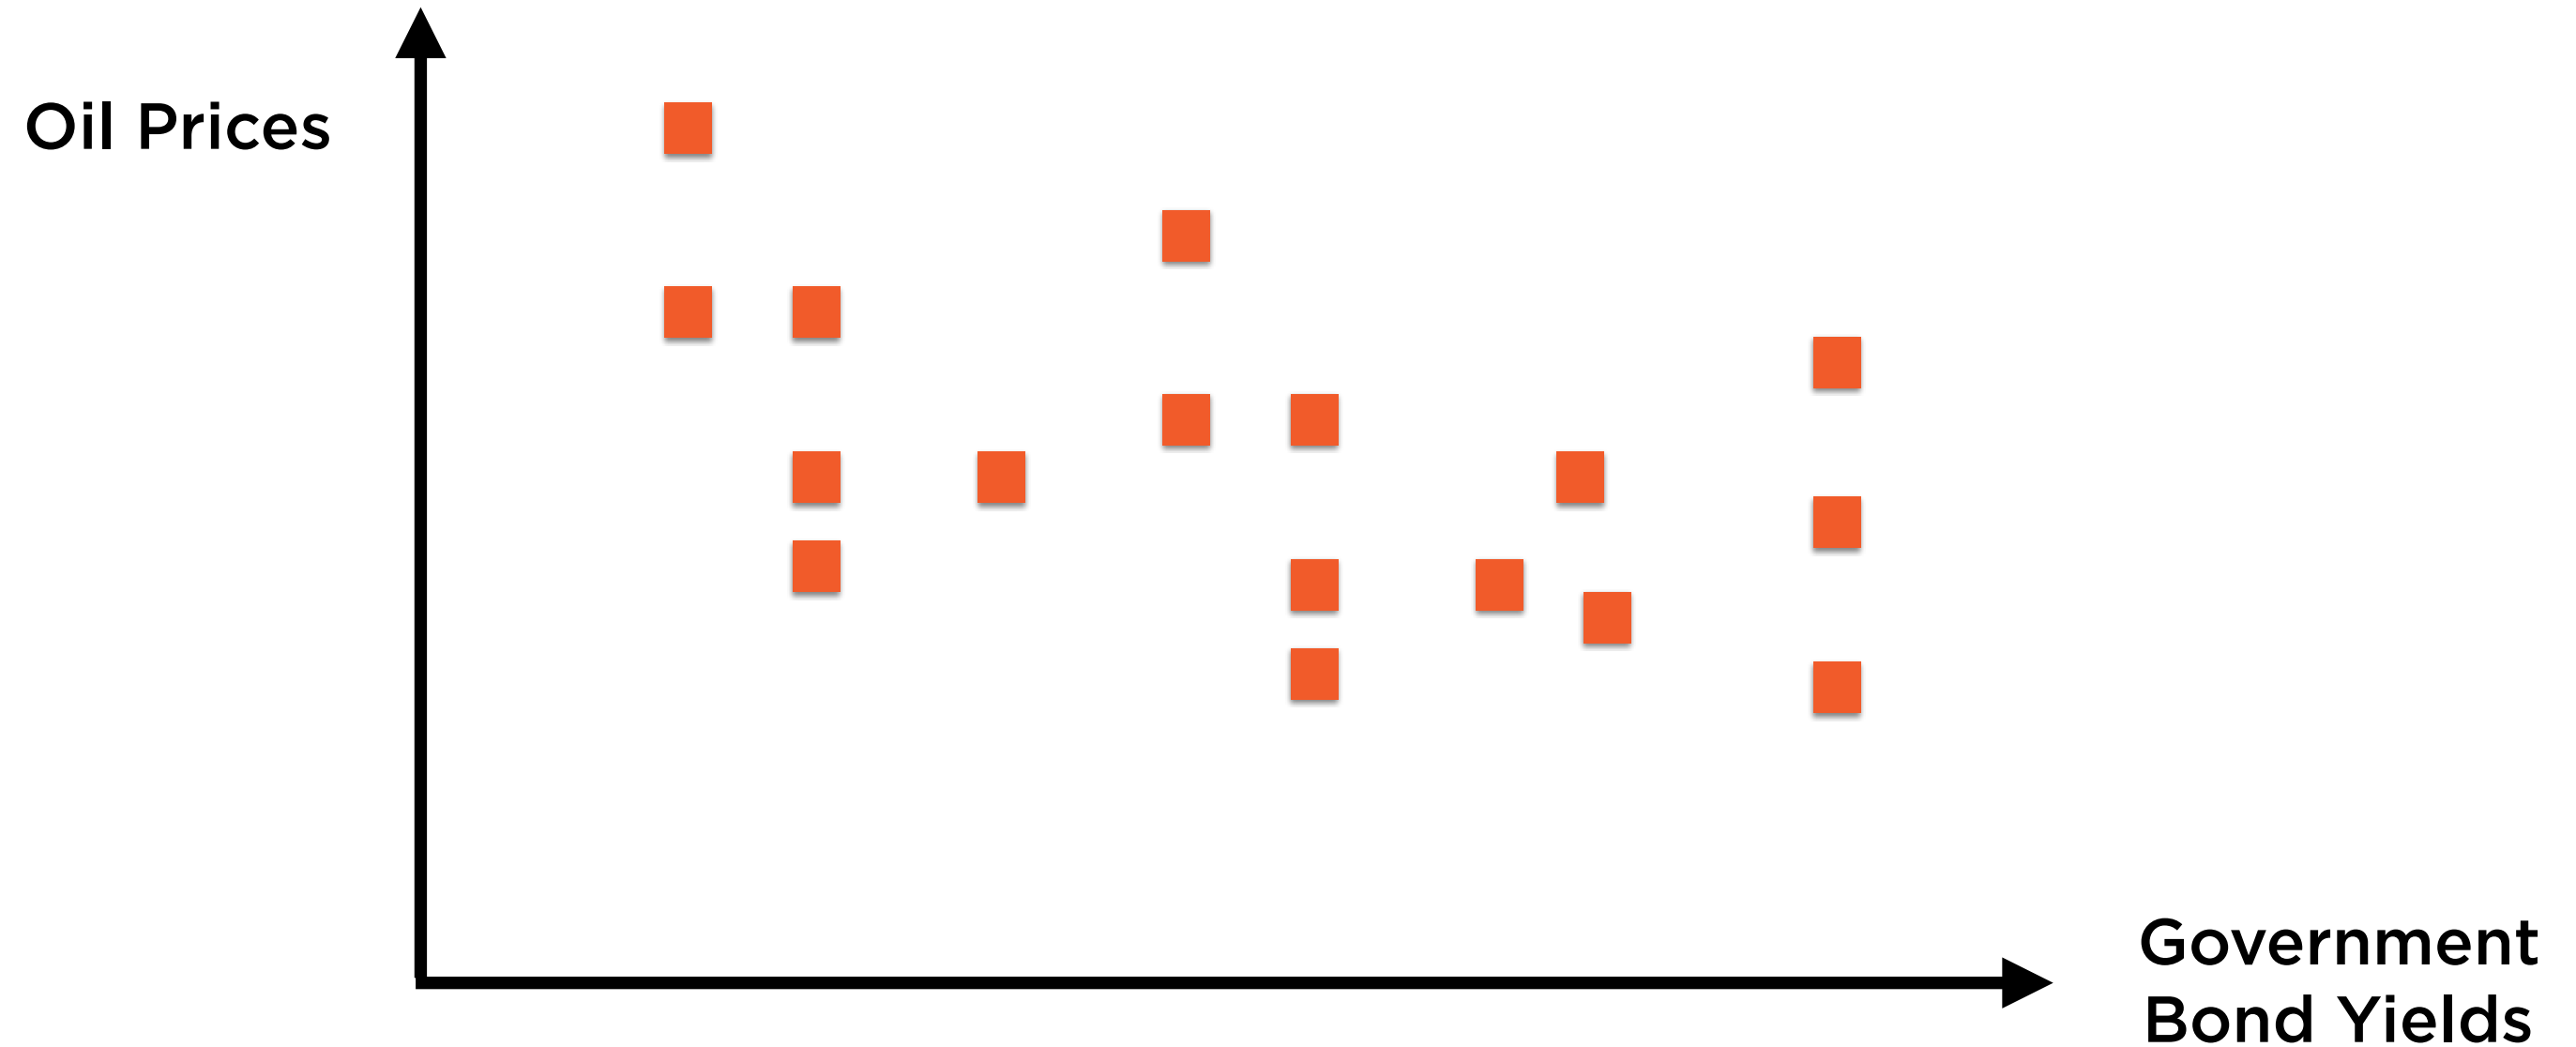
Unidimensional data is analysed using statistics such  
as mean, median, standard deviation

# Data in Two Dimensions



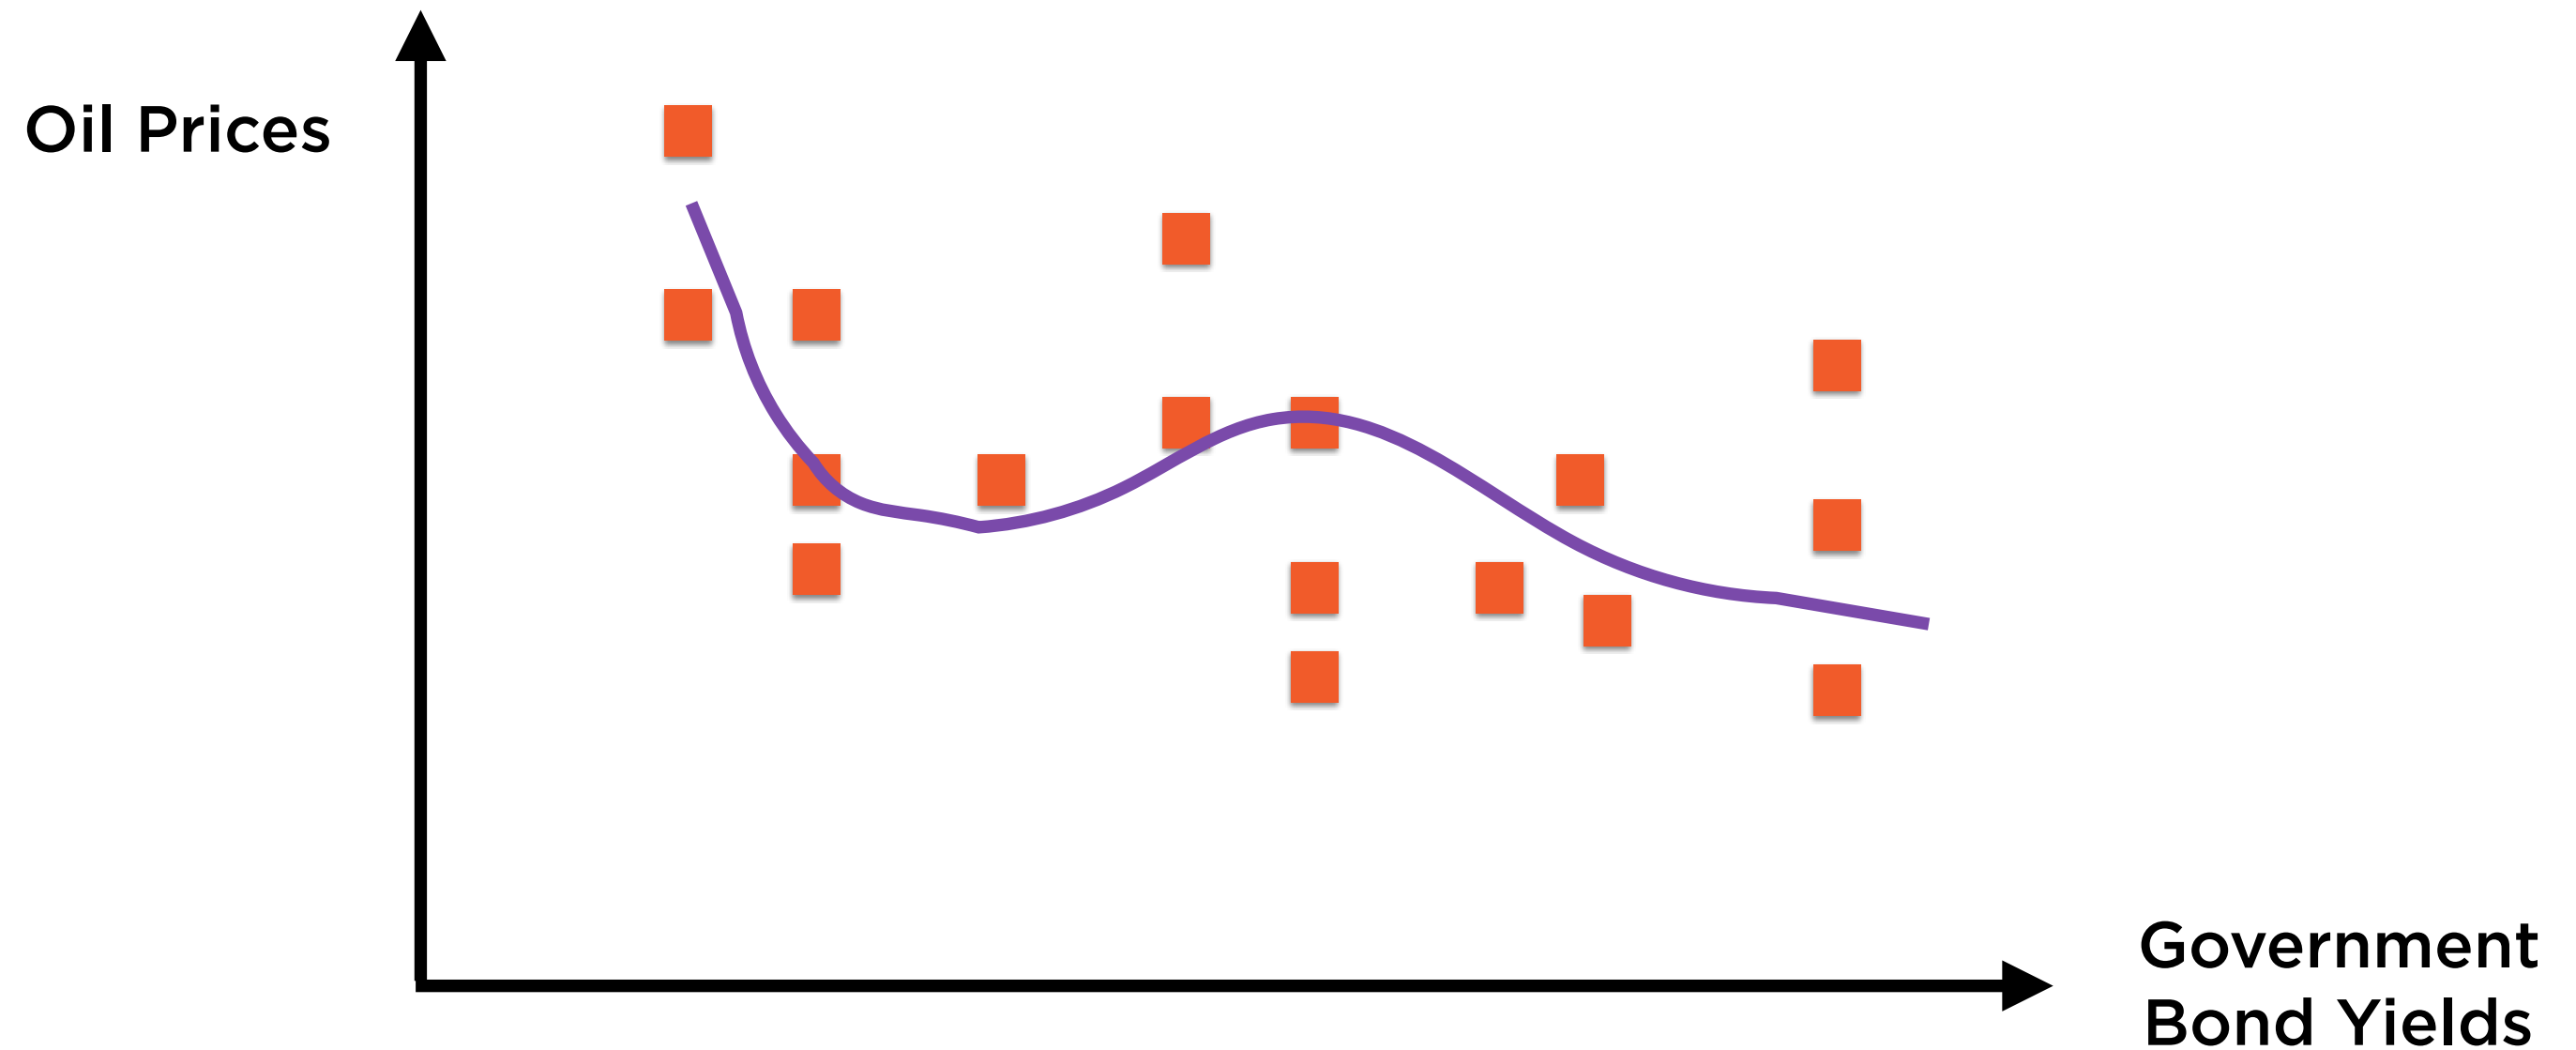
Its often more insightful to view data in relation to  
some other, related data

# Data in Two Dimensions



Bidimensional data can be represented in a plane

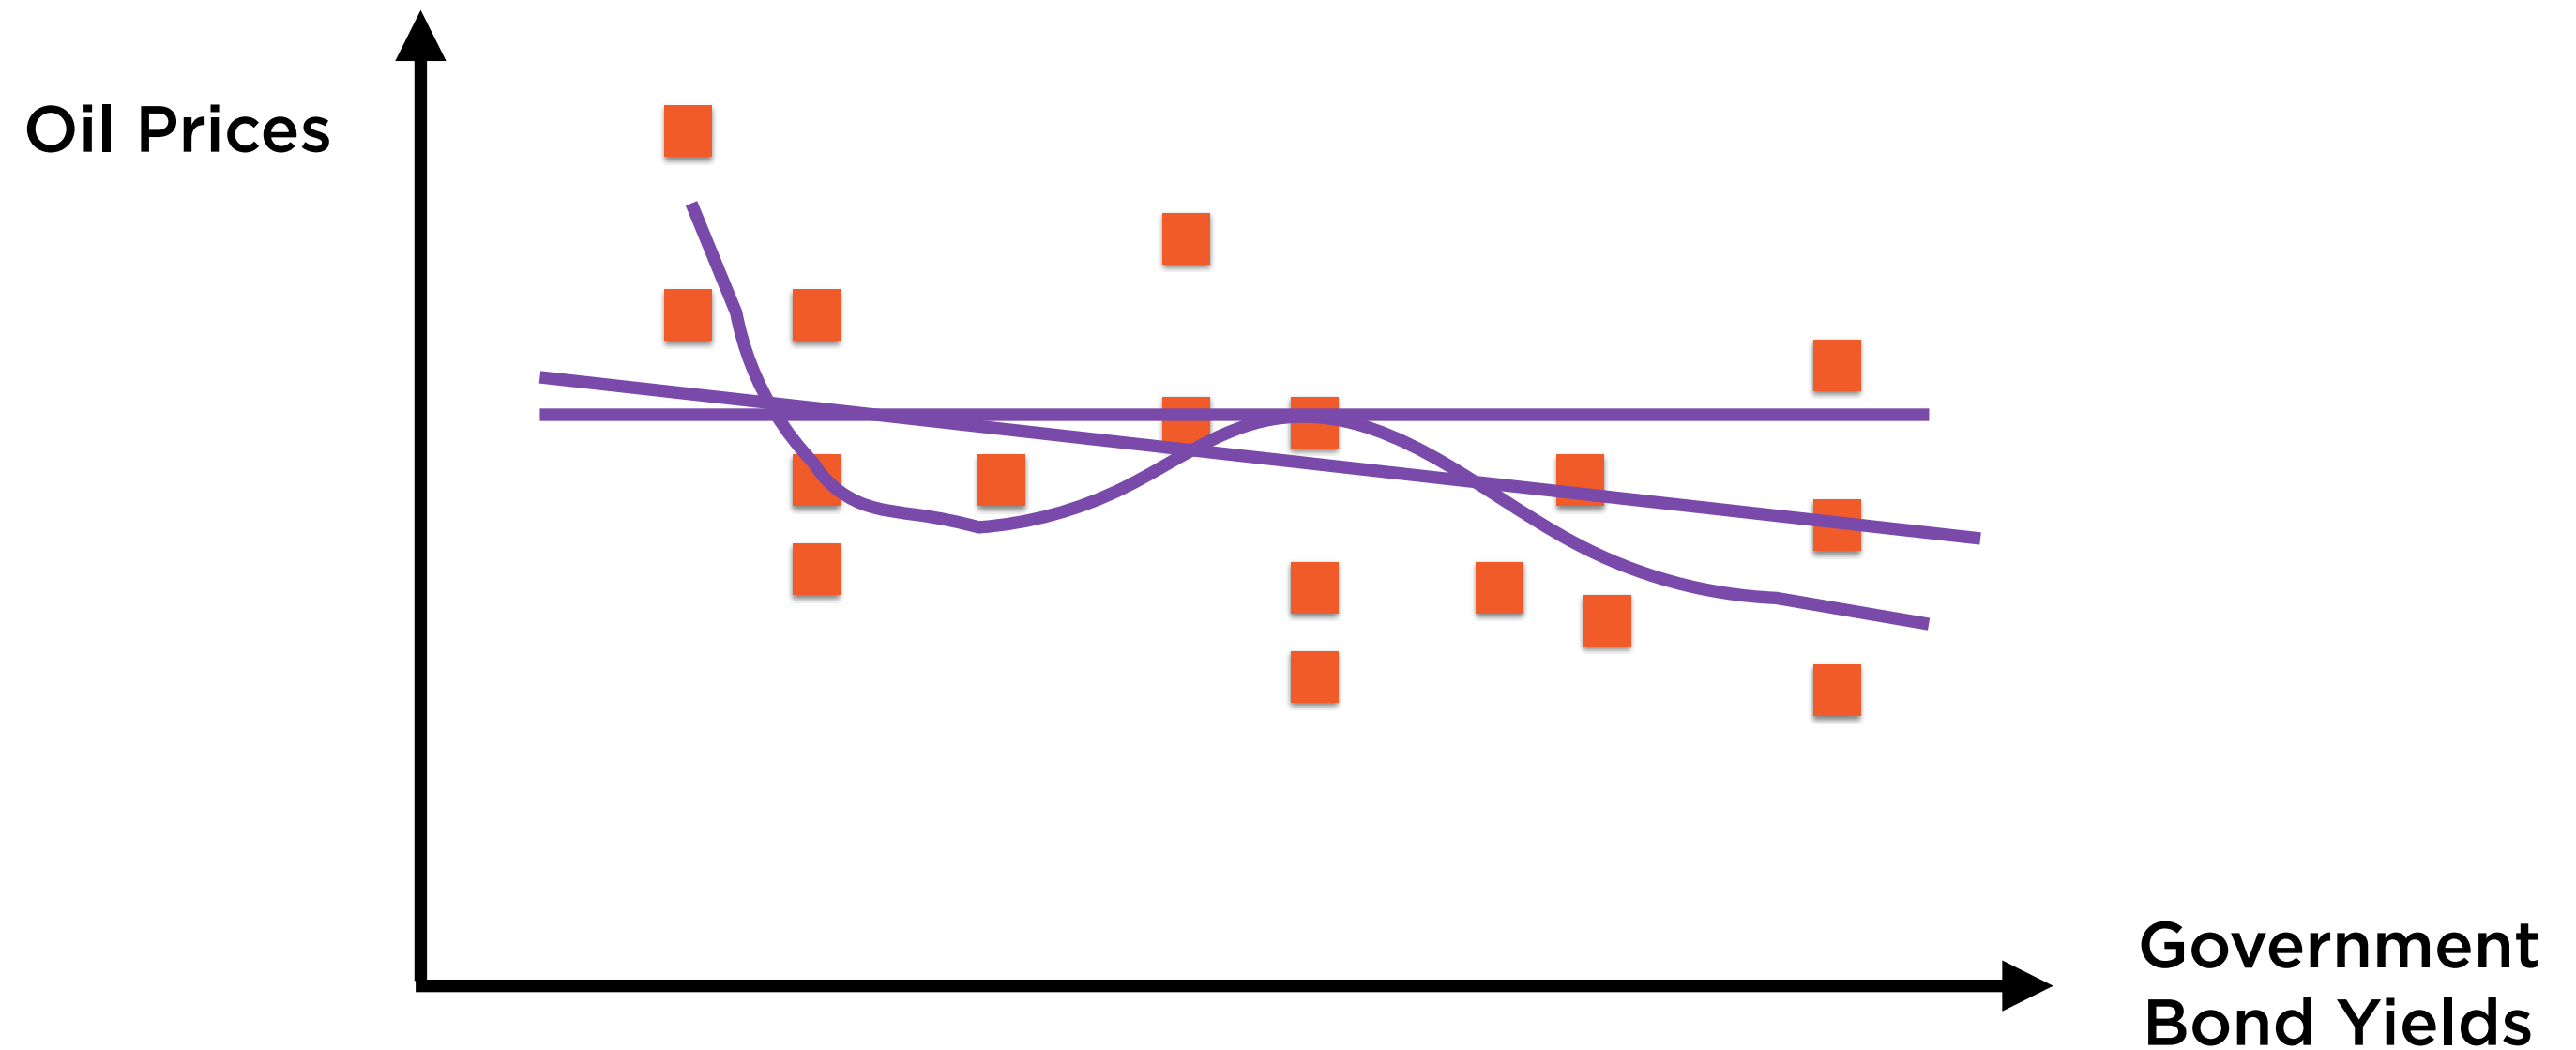
# Data in Two Dimensions



We can draw any number of curves to fit such data

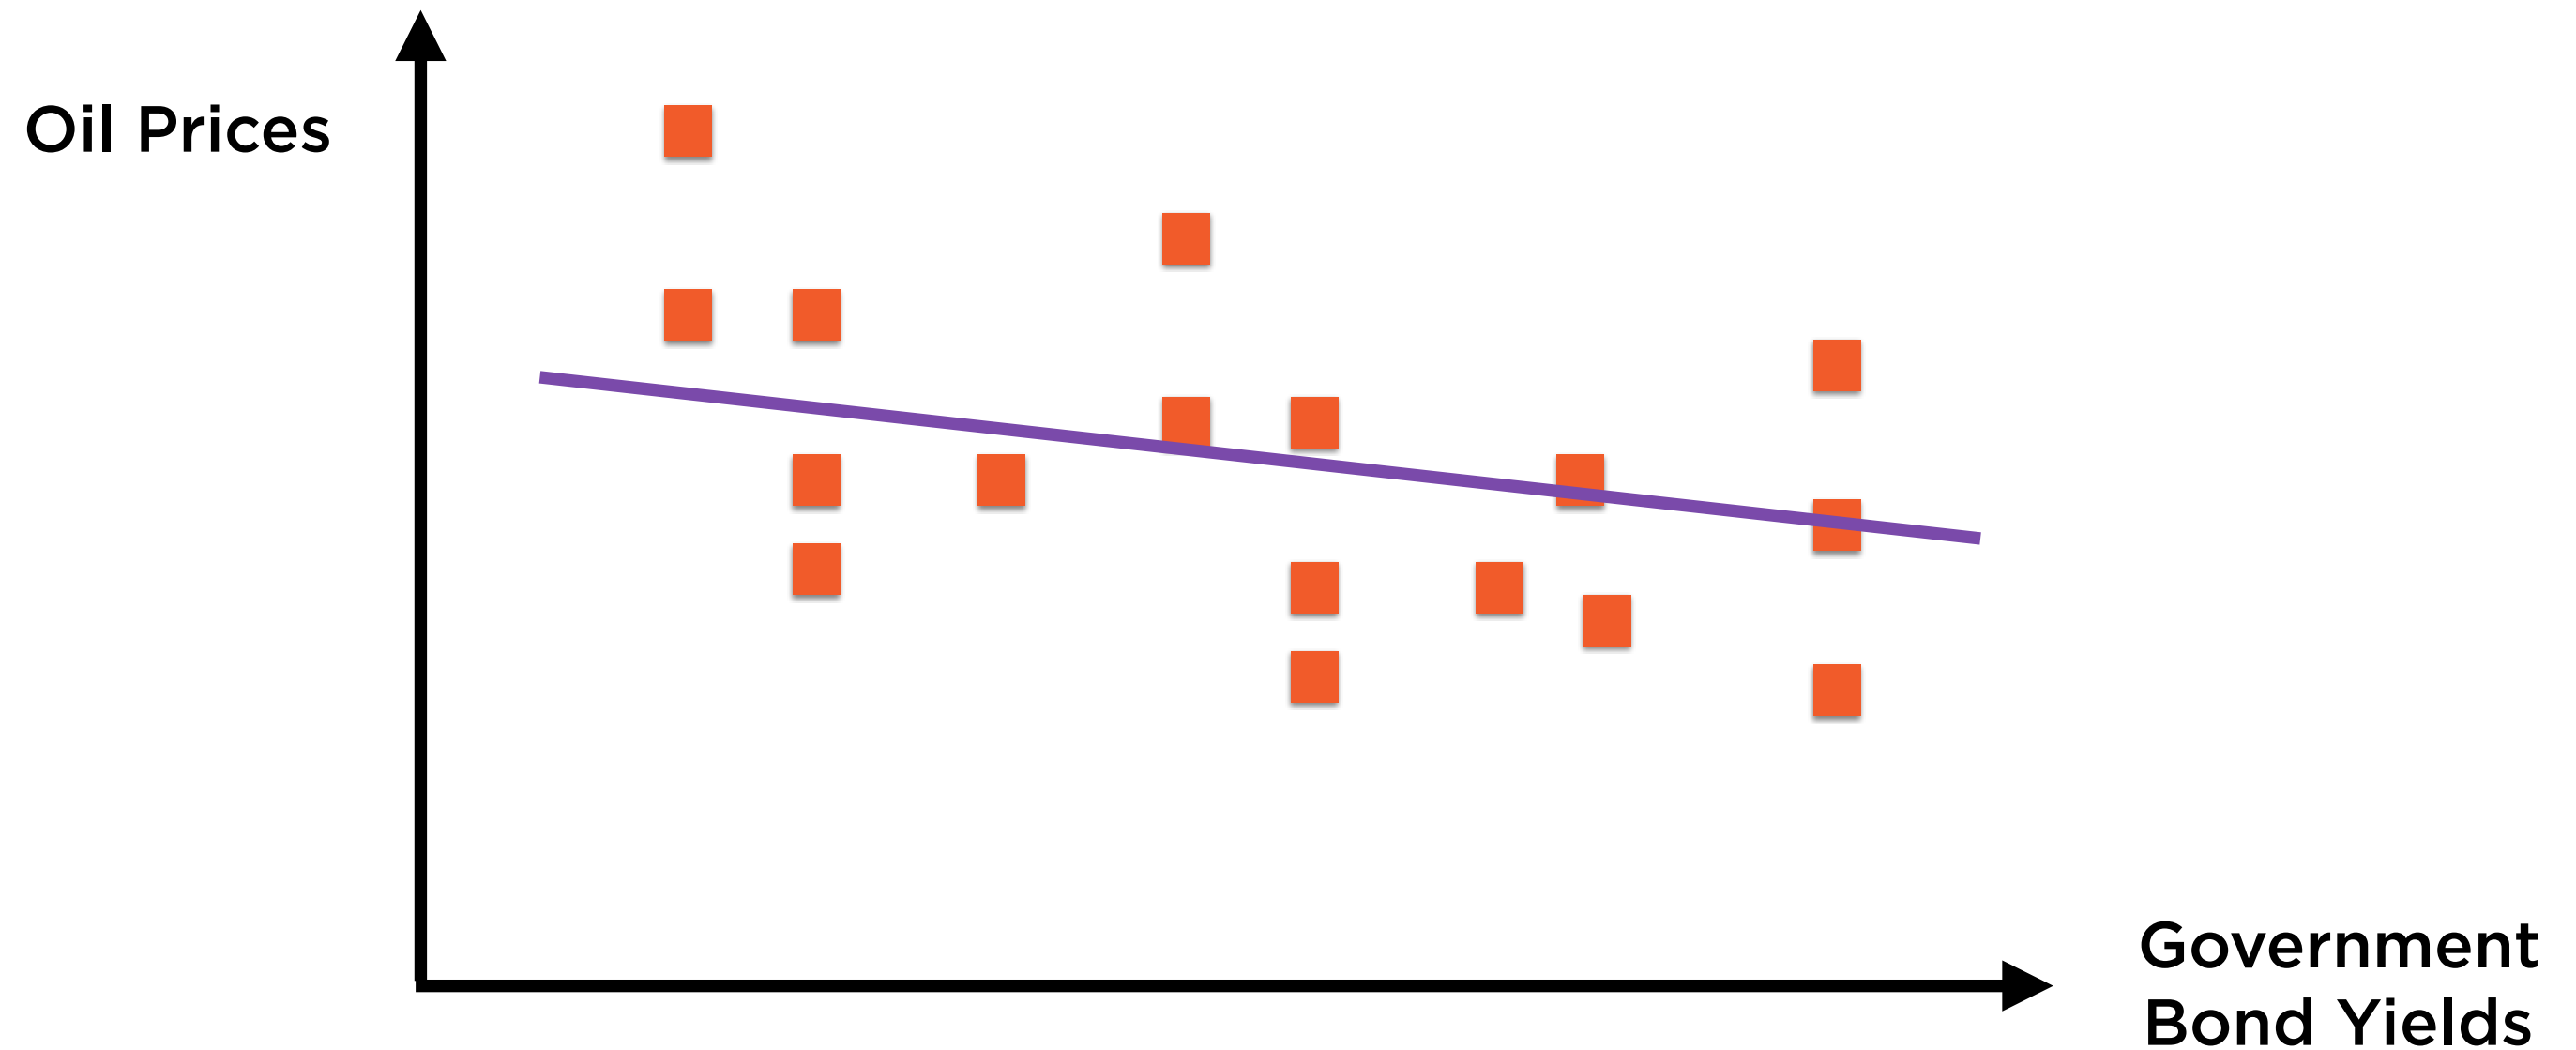


# Data in Two Dimensions



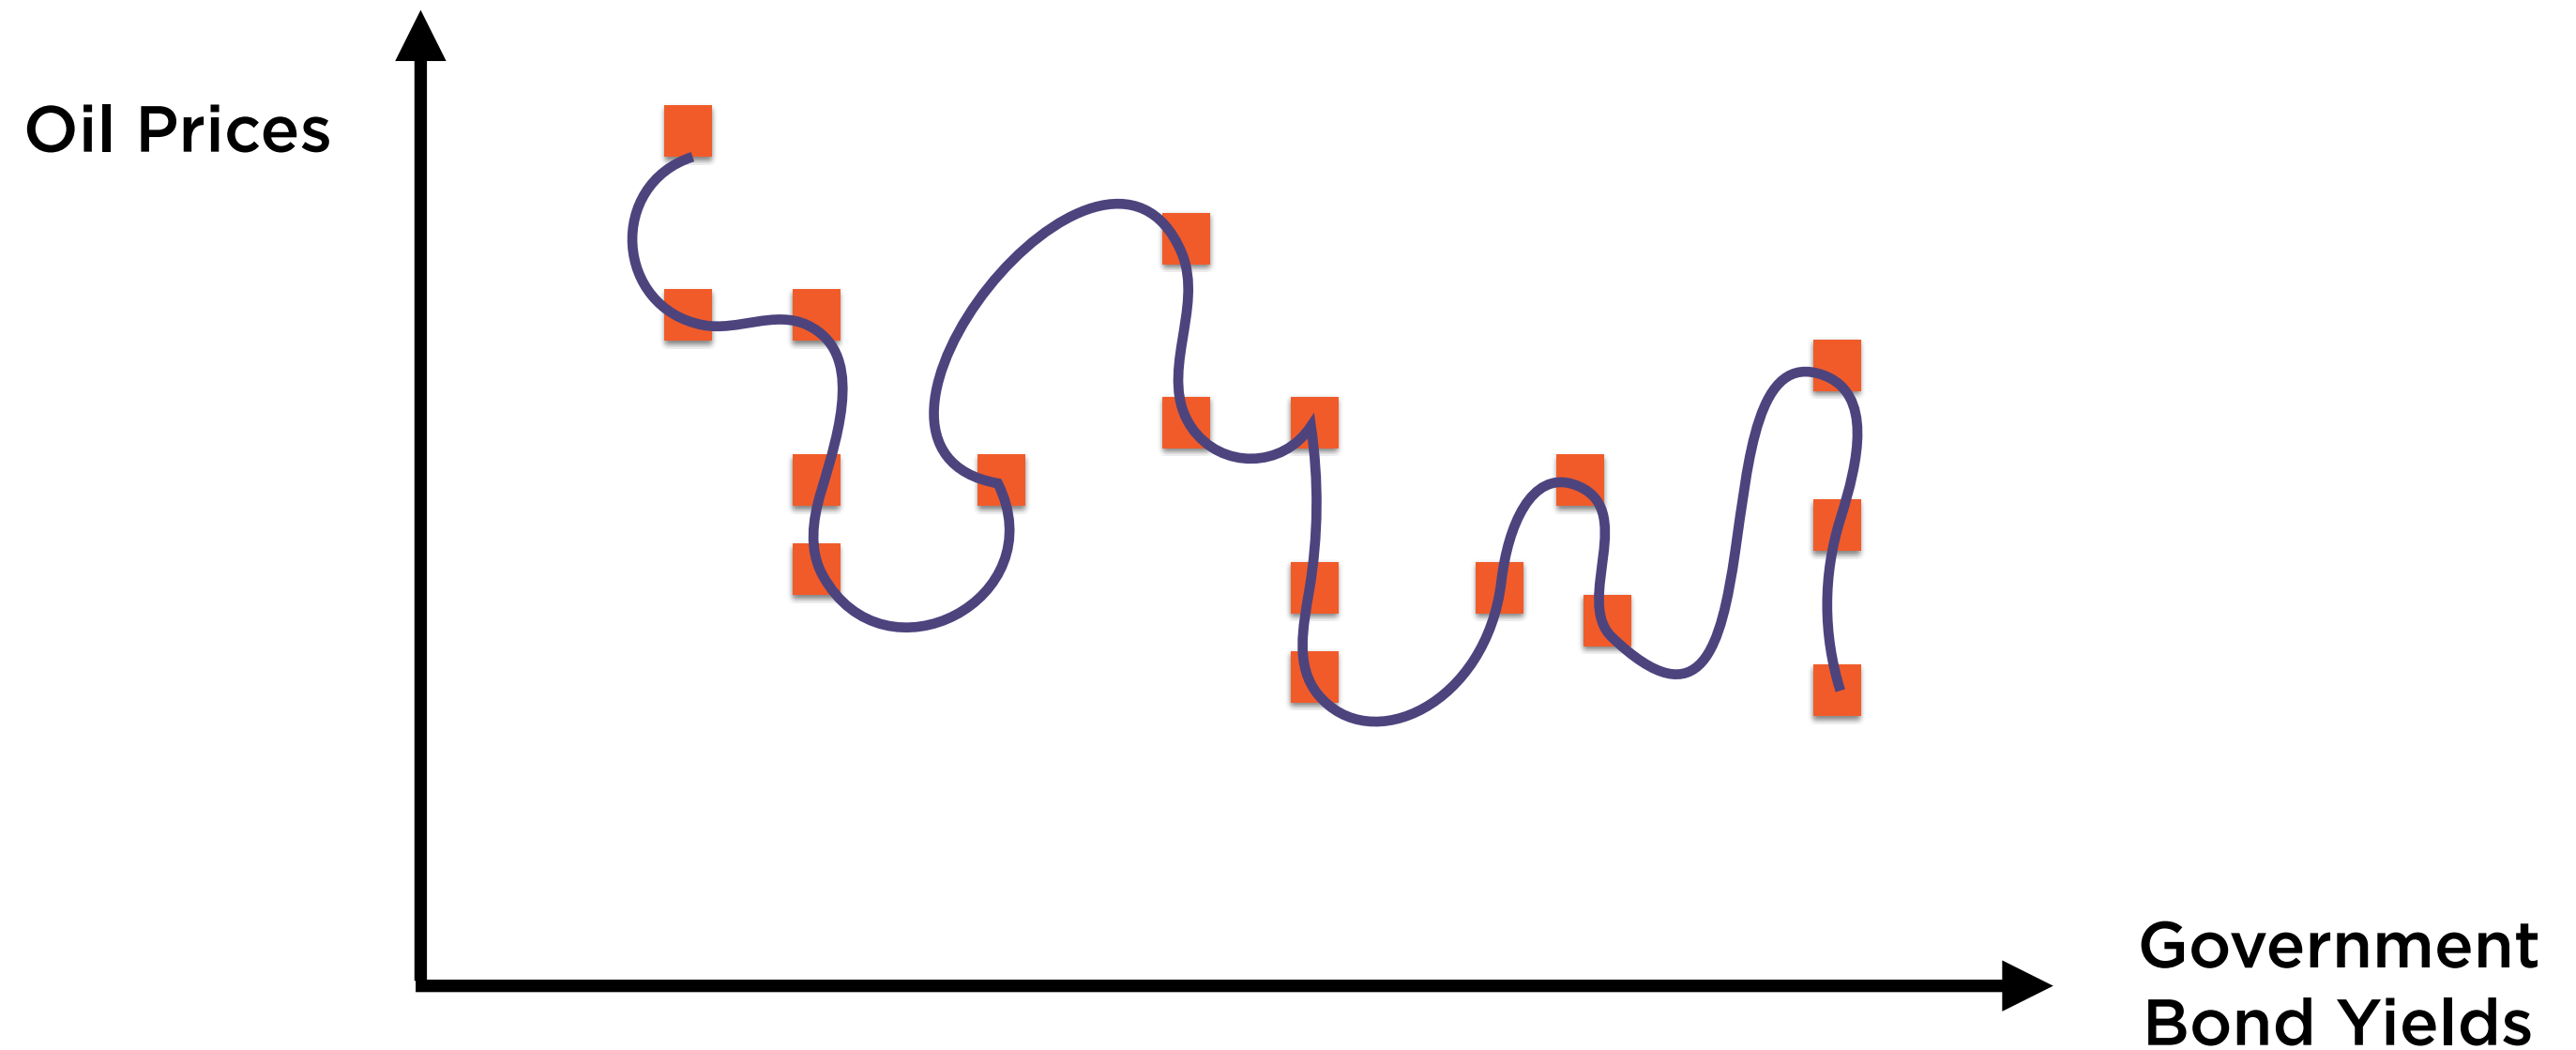
We can draw any number of curves to fit such data

# Data in Two Dimensions



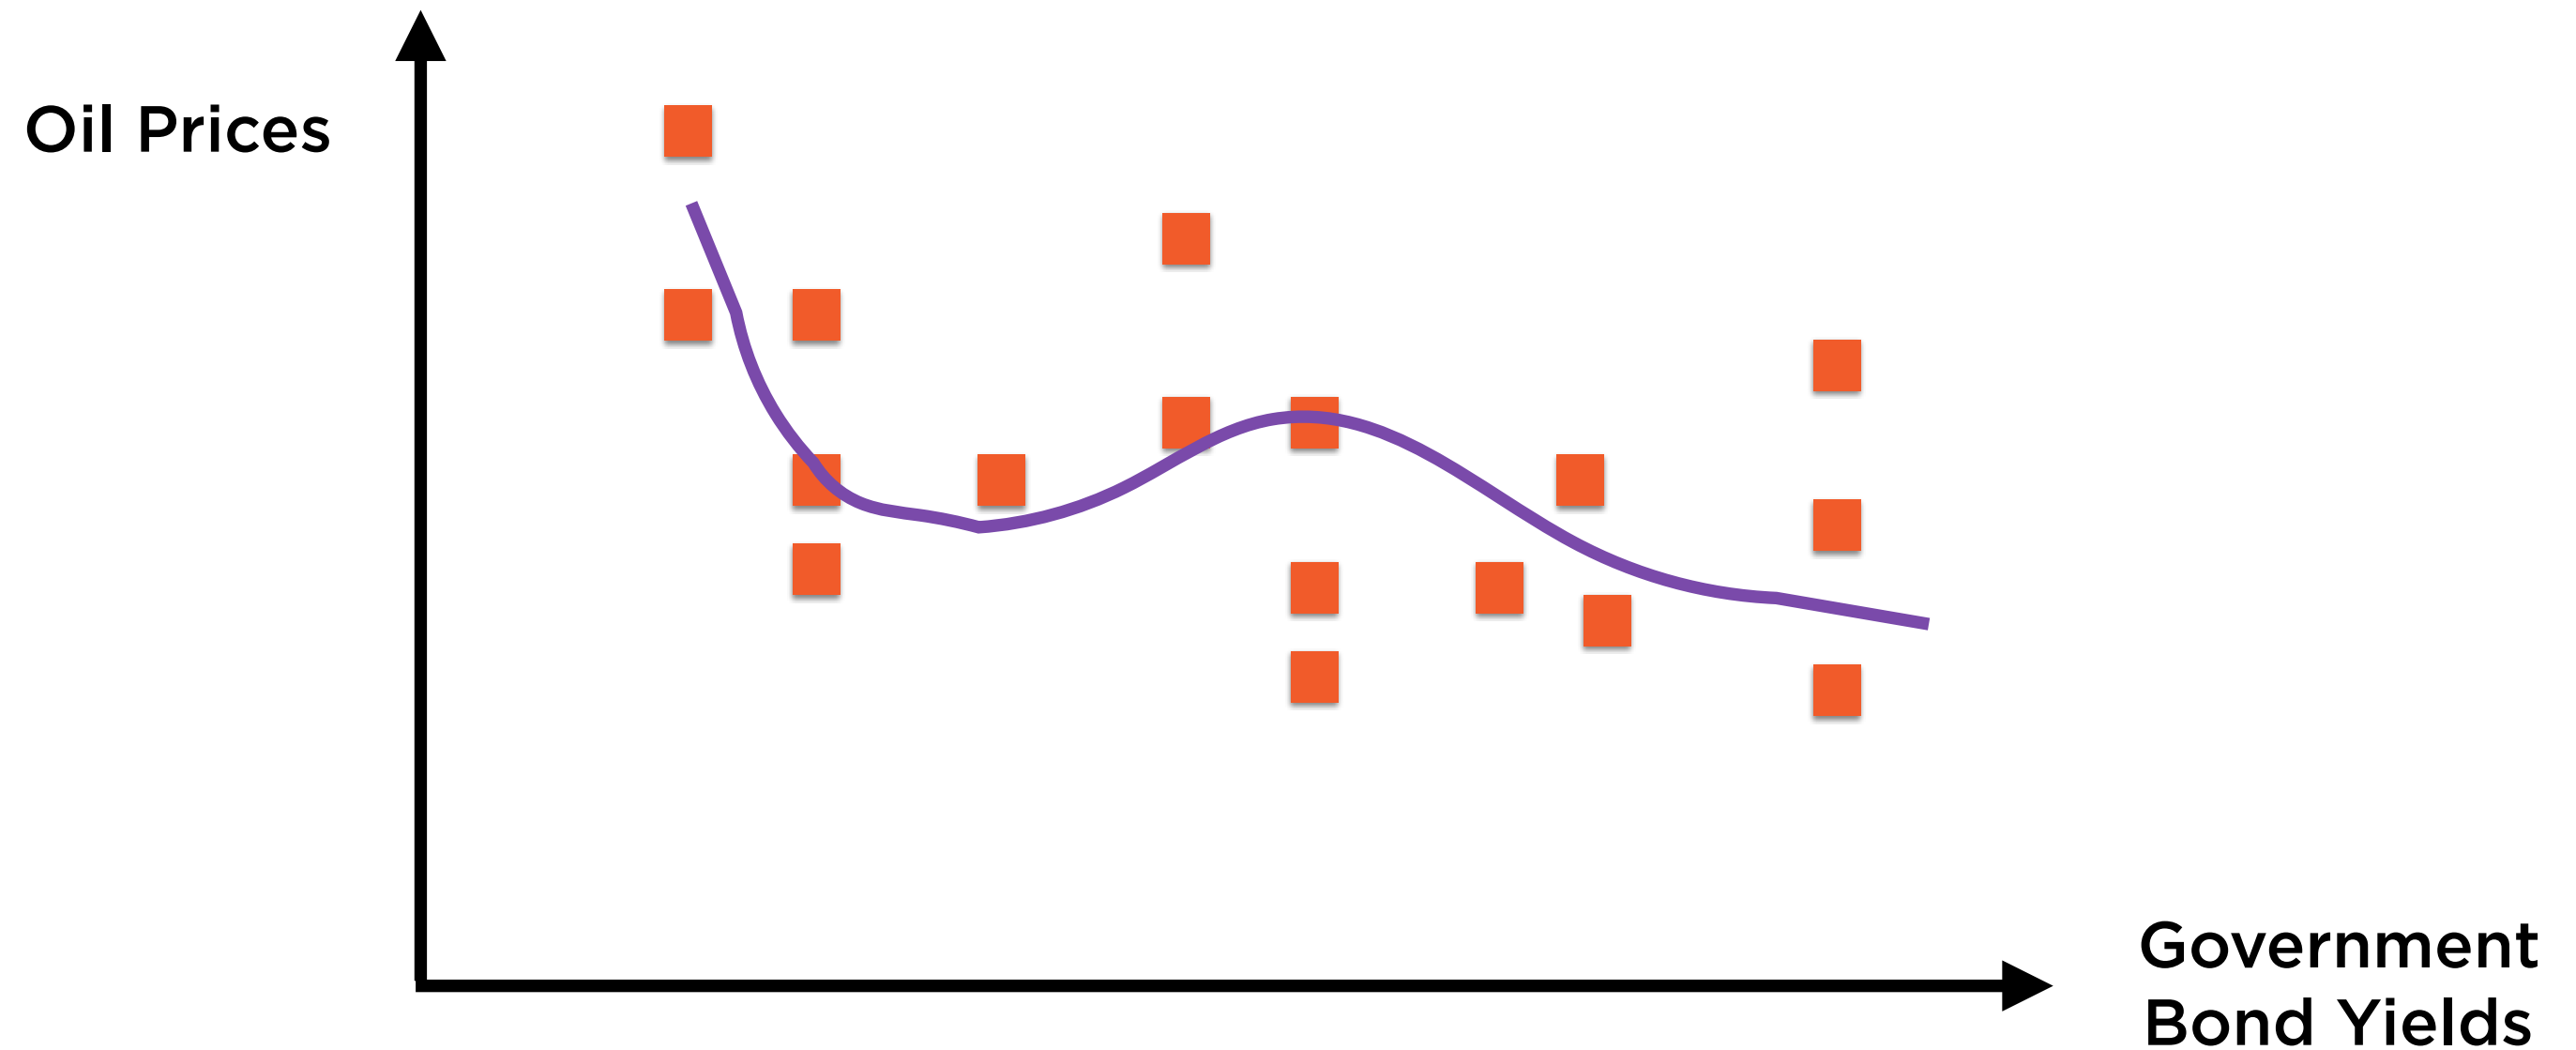
A straight line represents a linear relationship

# Data in Two Dimensions



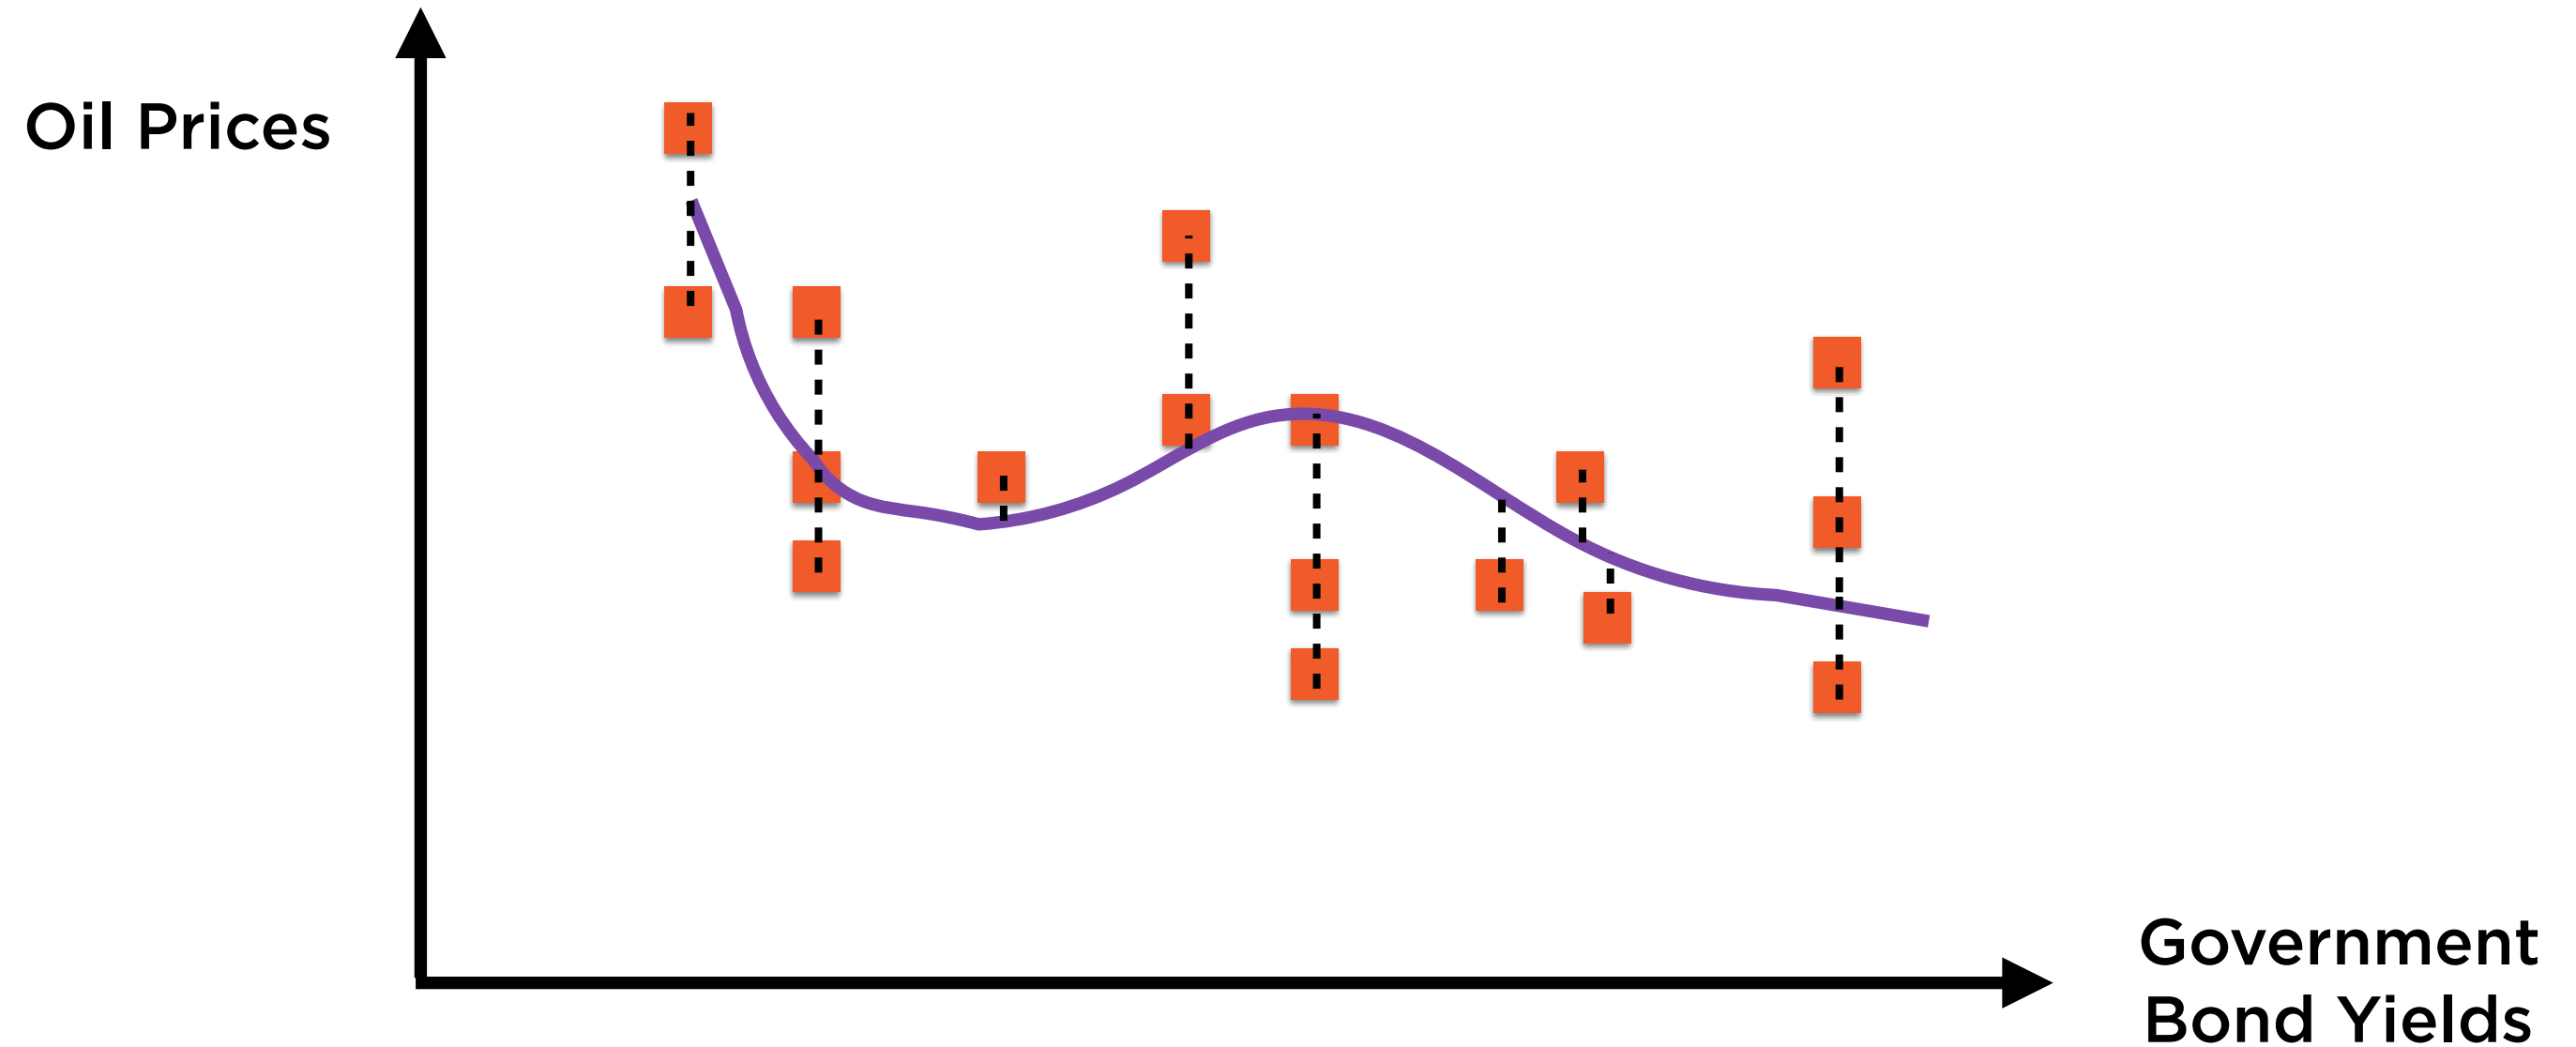
We could either make this curve pass through each point...

# Data in Two Dimensions




...Or in some sense “fit” the data in aggregate

# Data in Two Dimensions



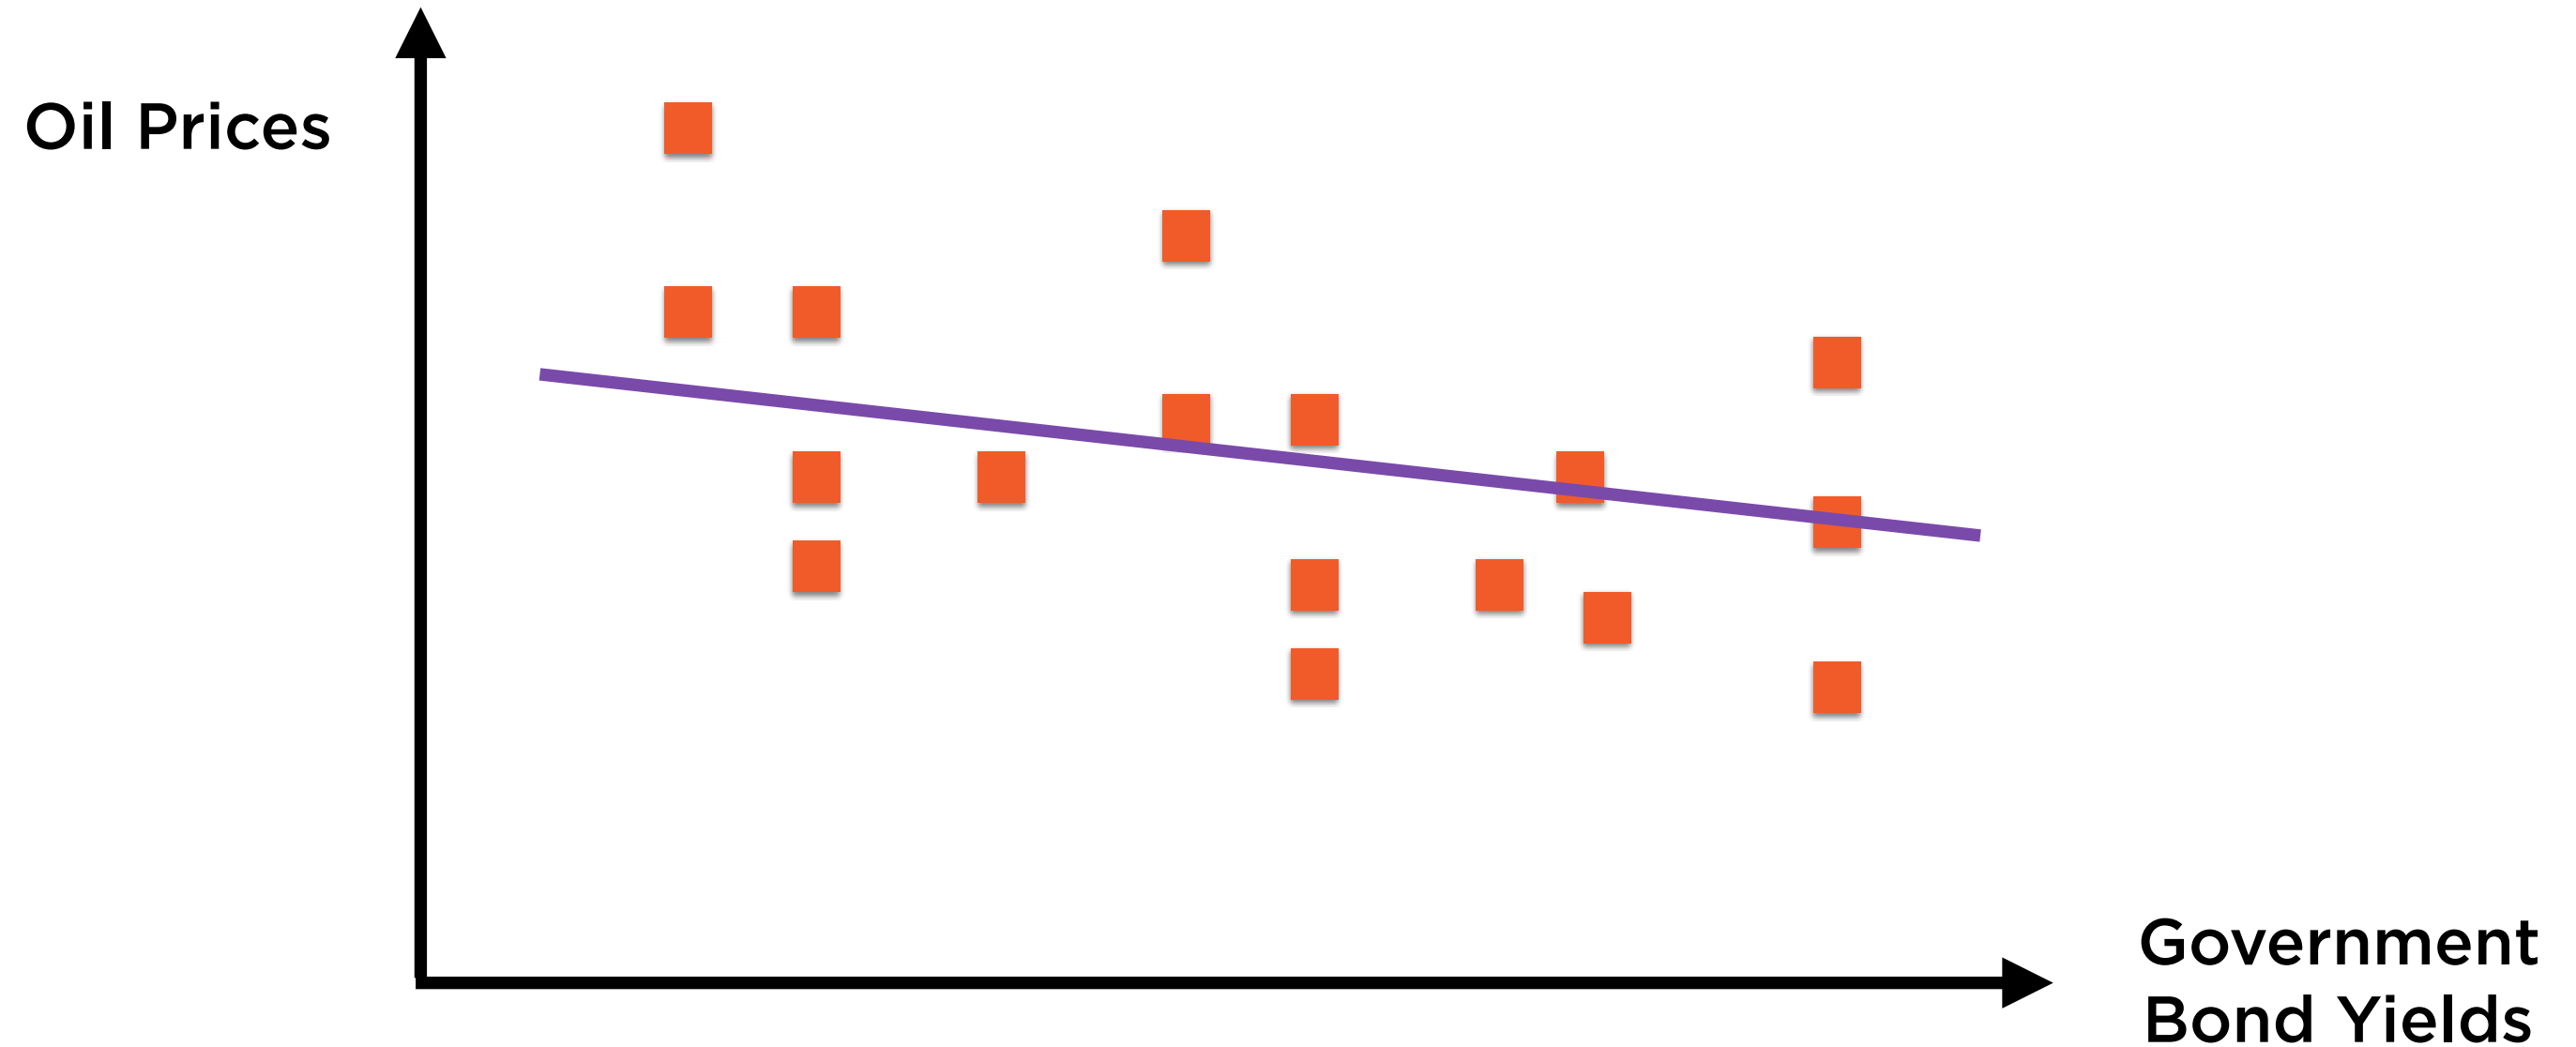
A curve has a “good fit” if the distances of points from the curve are small

# Data in Two Dimensions



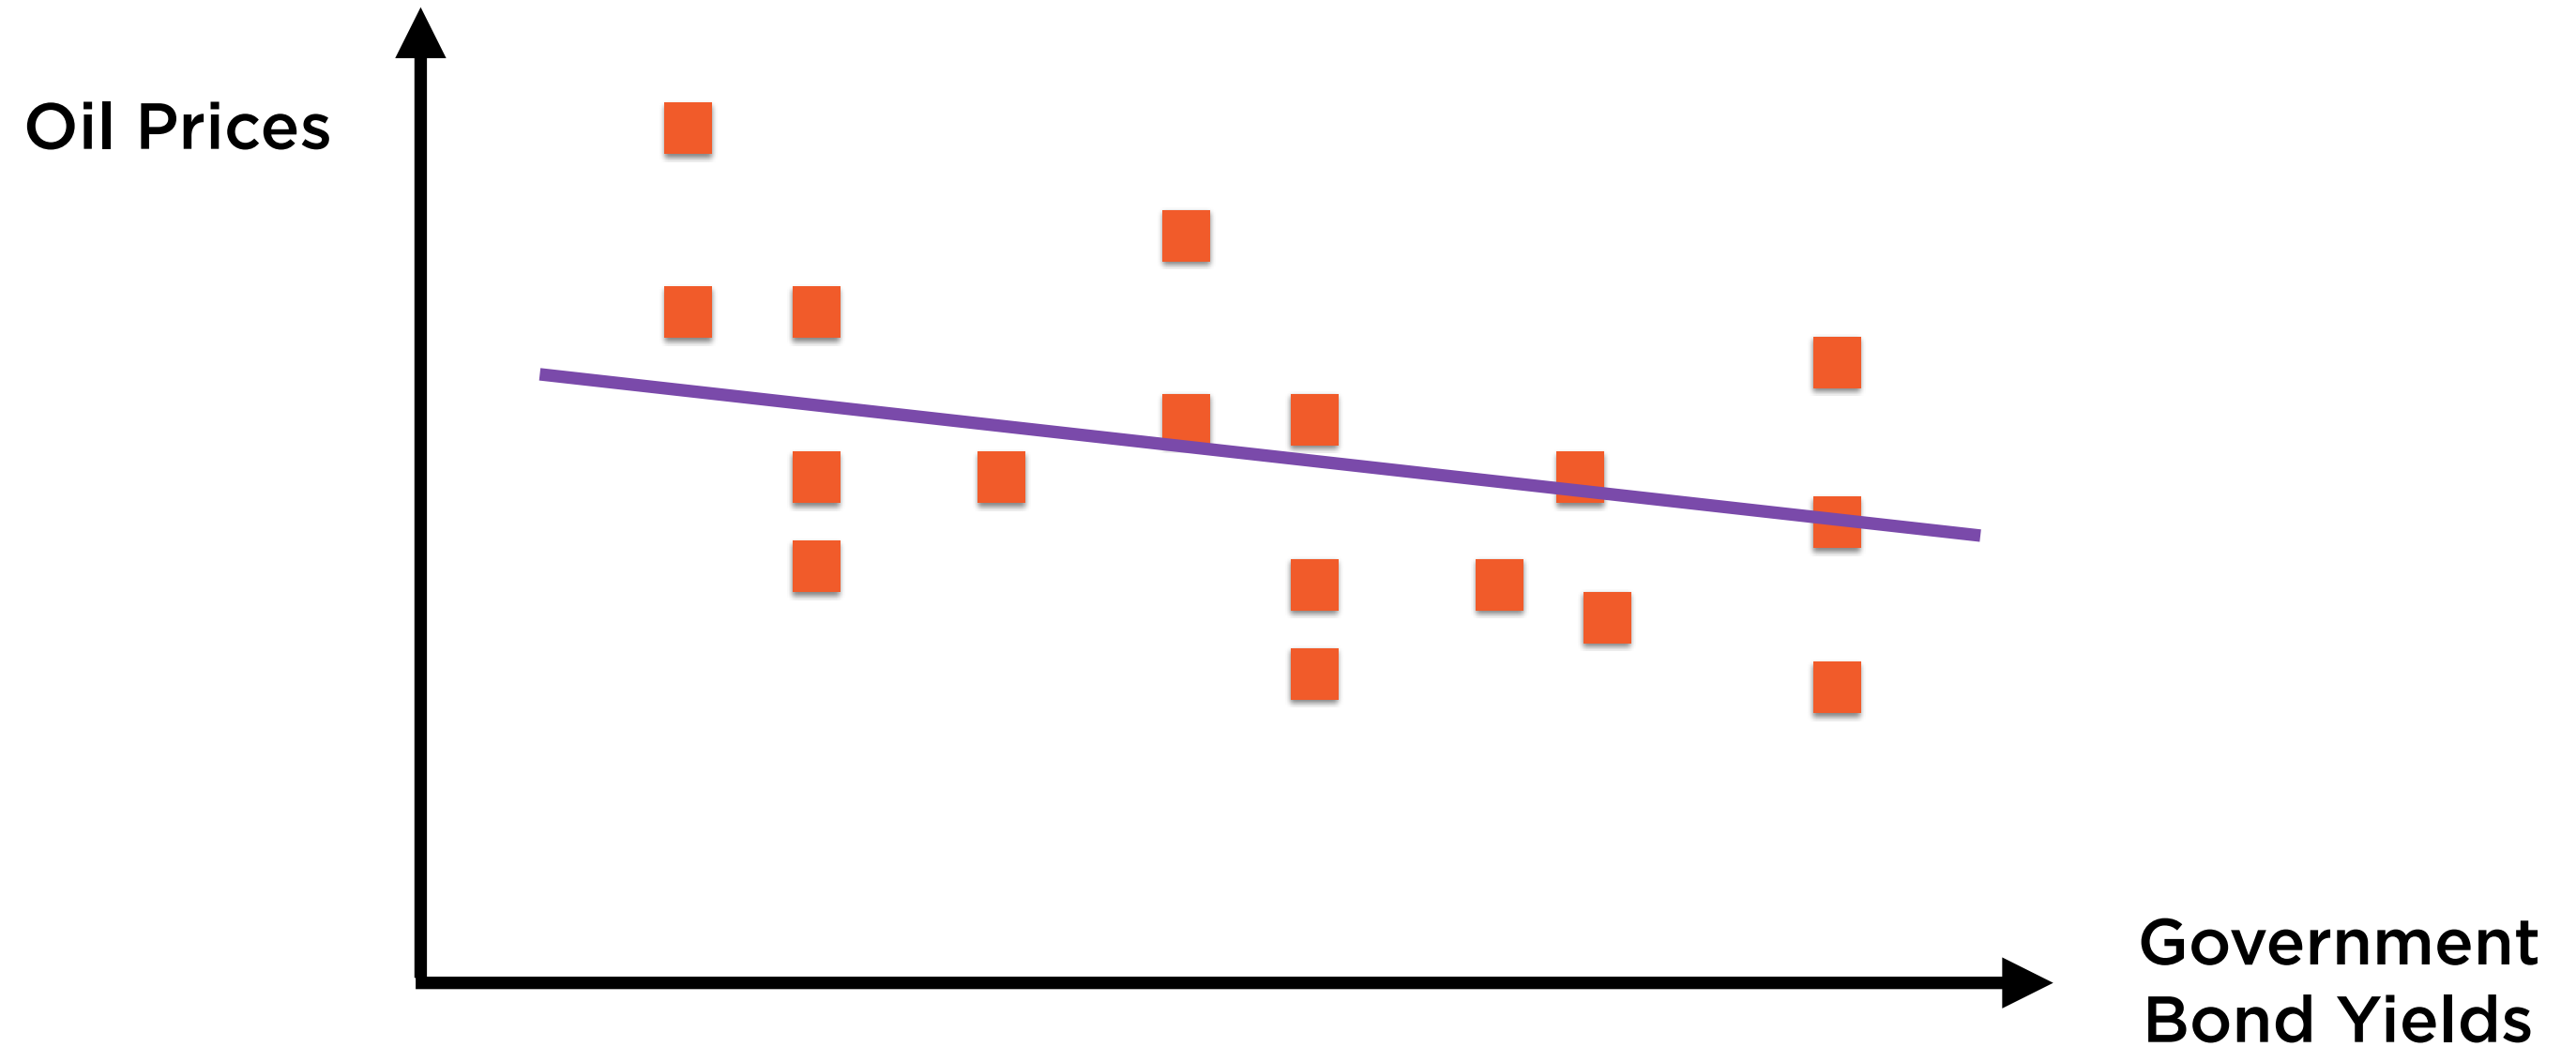
Overfitting by finding a very complicated curve  
often only hurts predictive accuracy

# Data in Two Dimensions



Often, a straight line works just fine

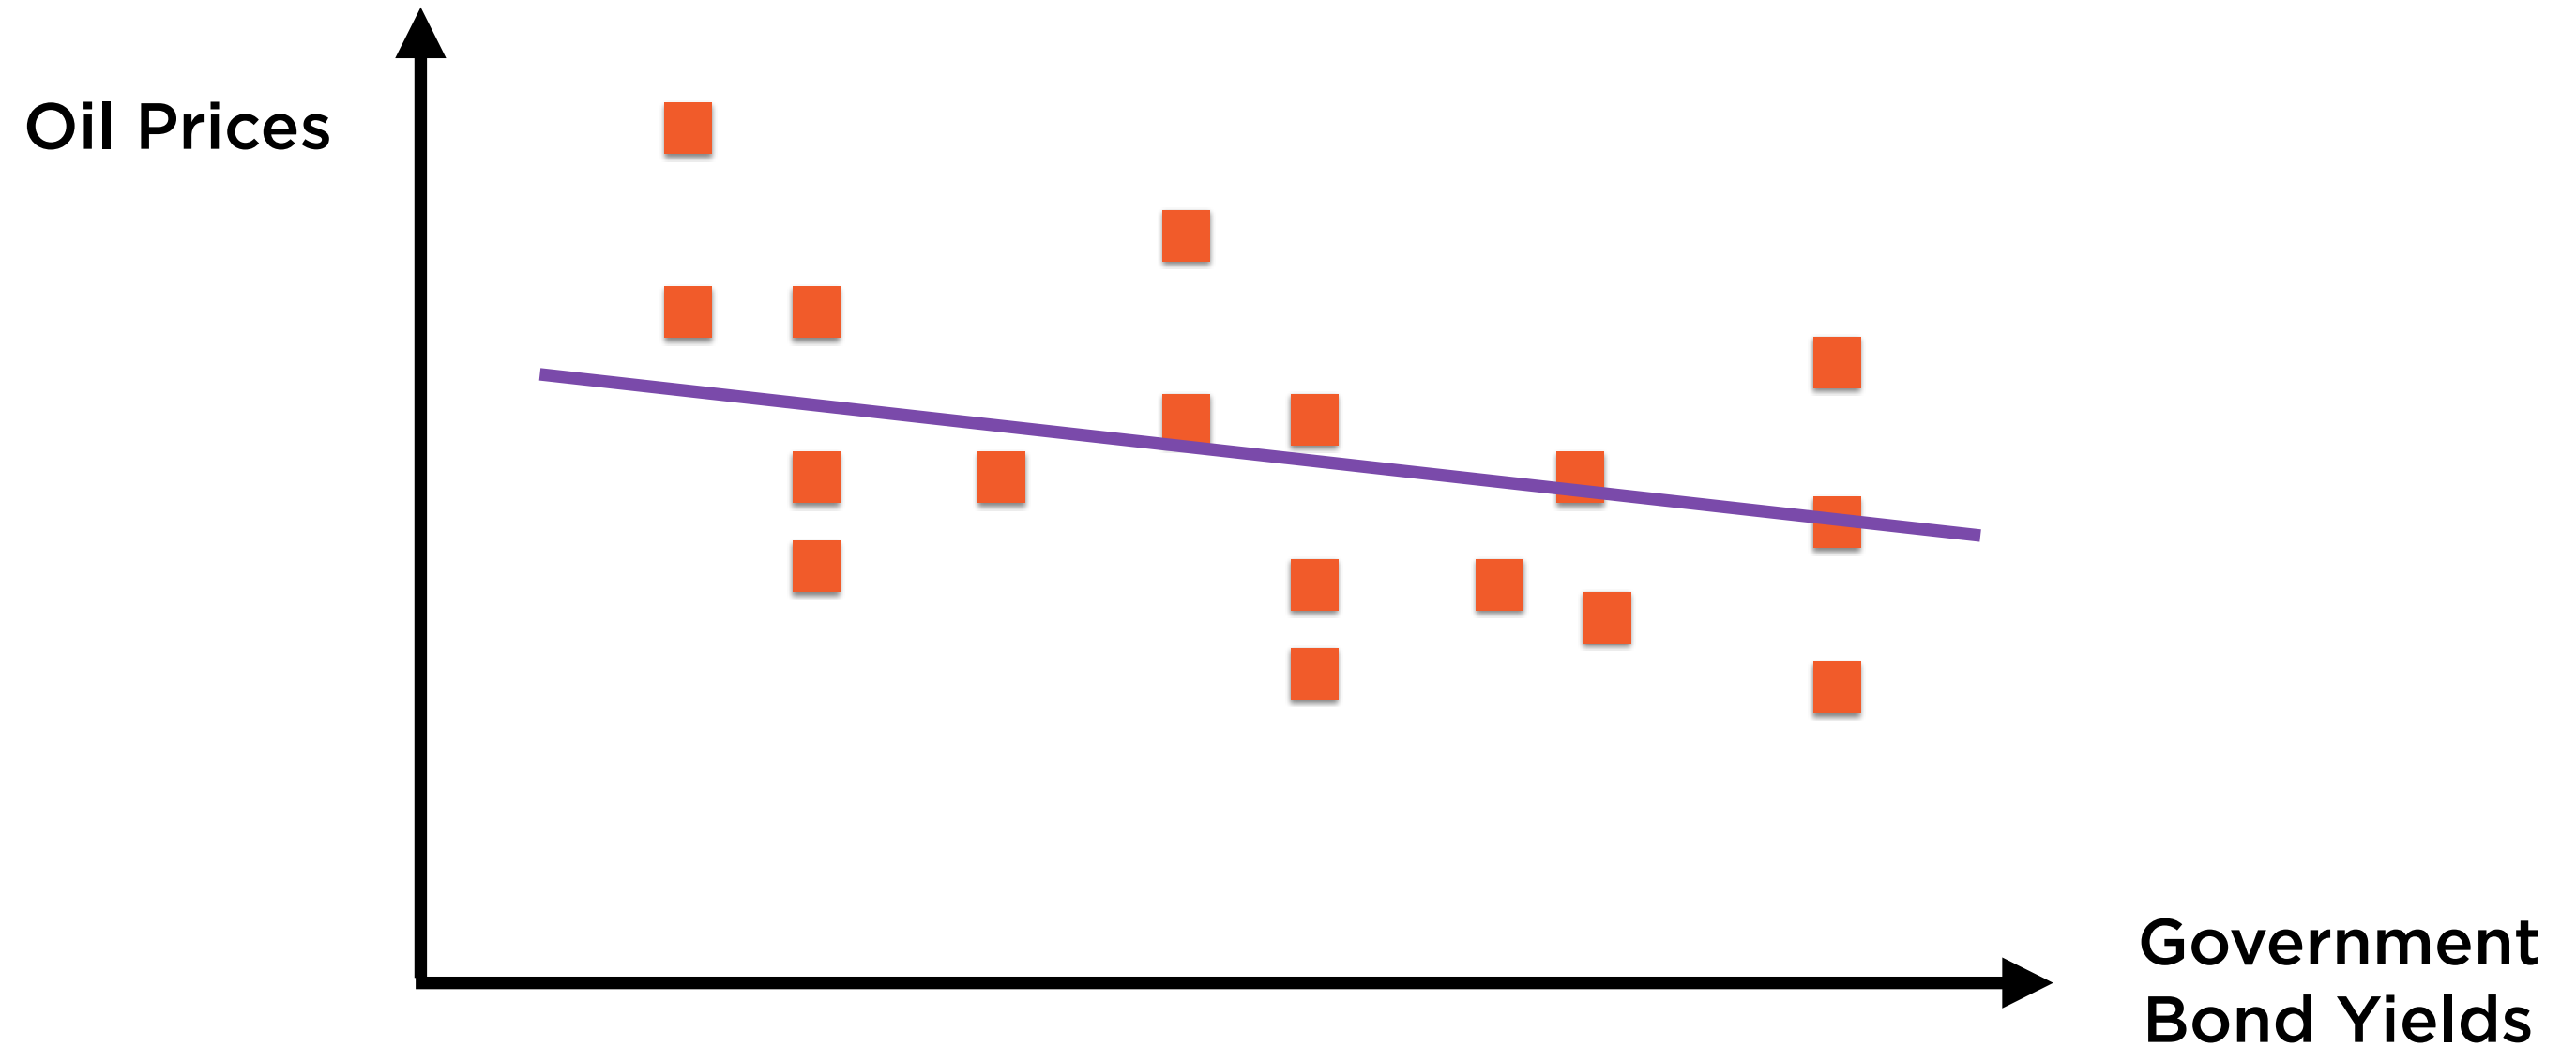
# Data in Two Dimensions



Finding the “best” such straight line is called **Linear Regression**

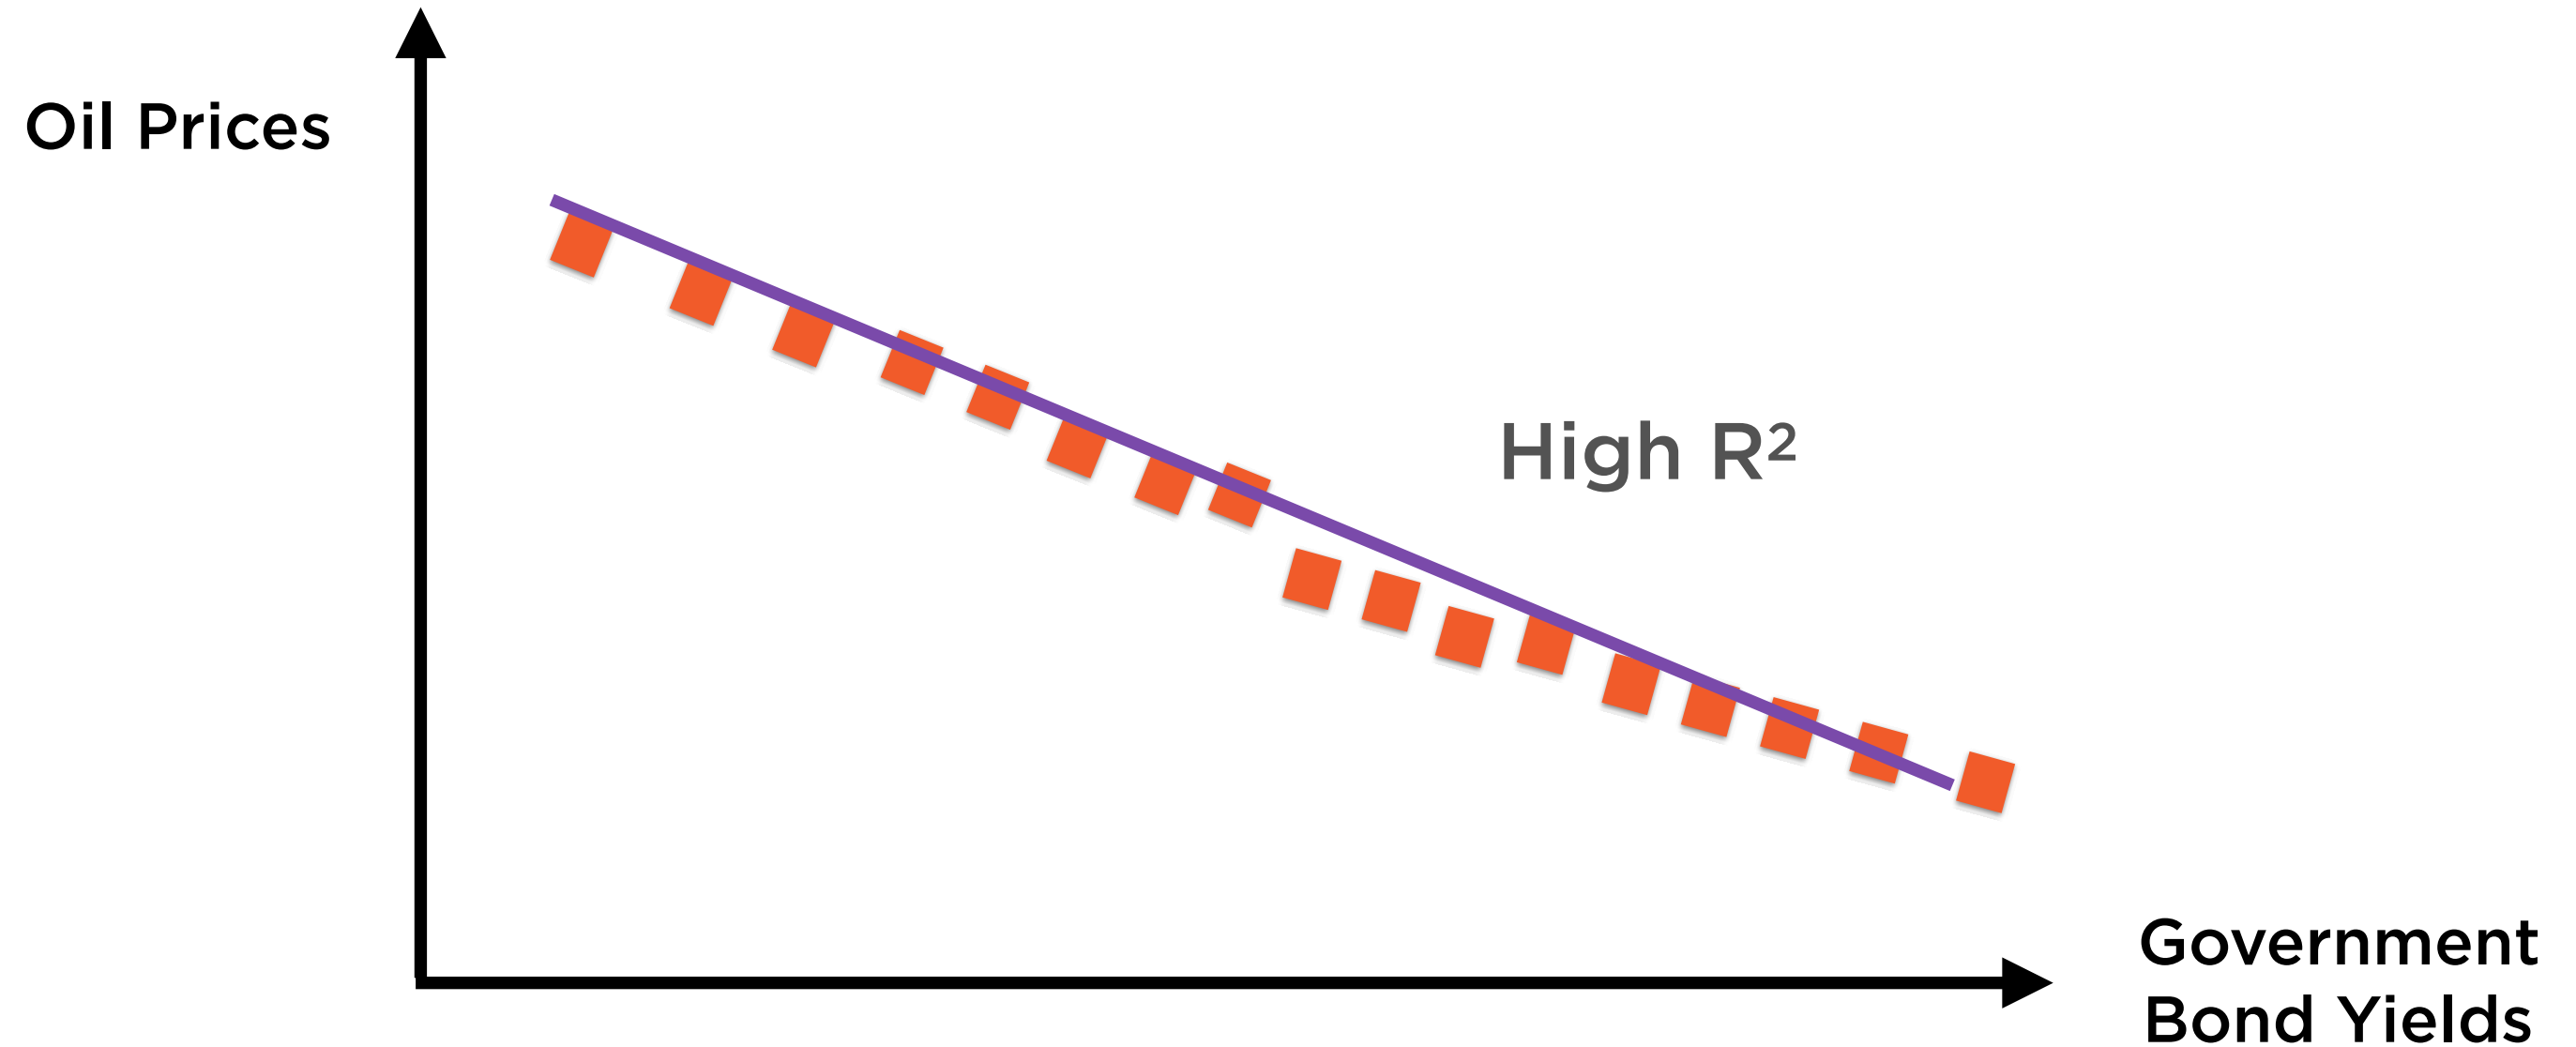


# Data in Two Dimensions



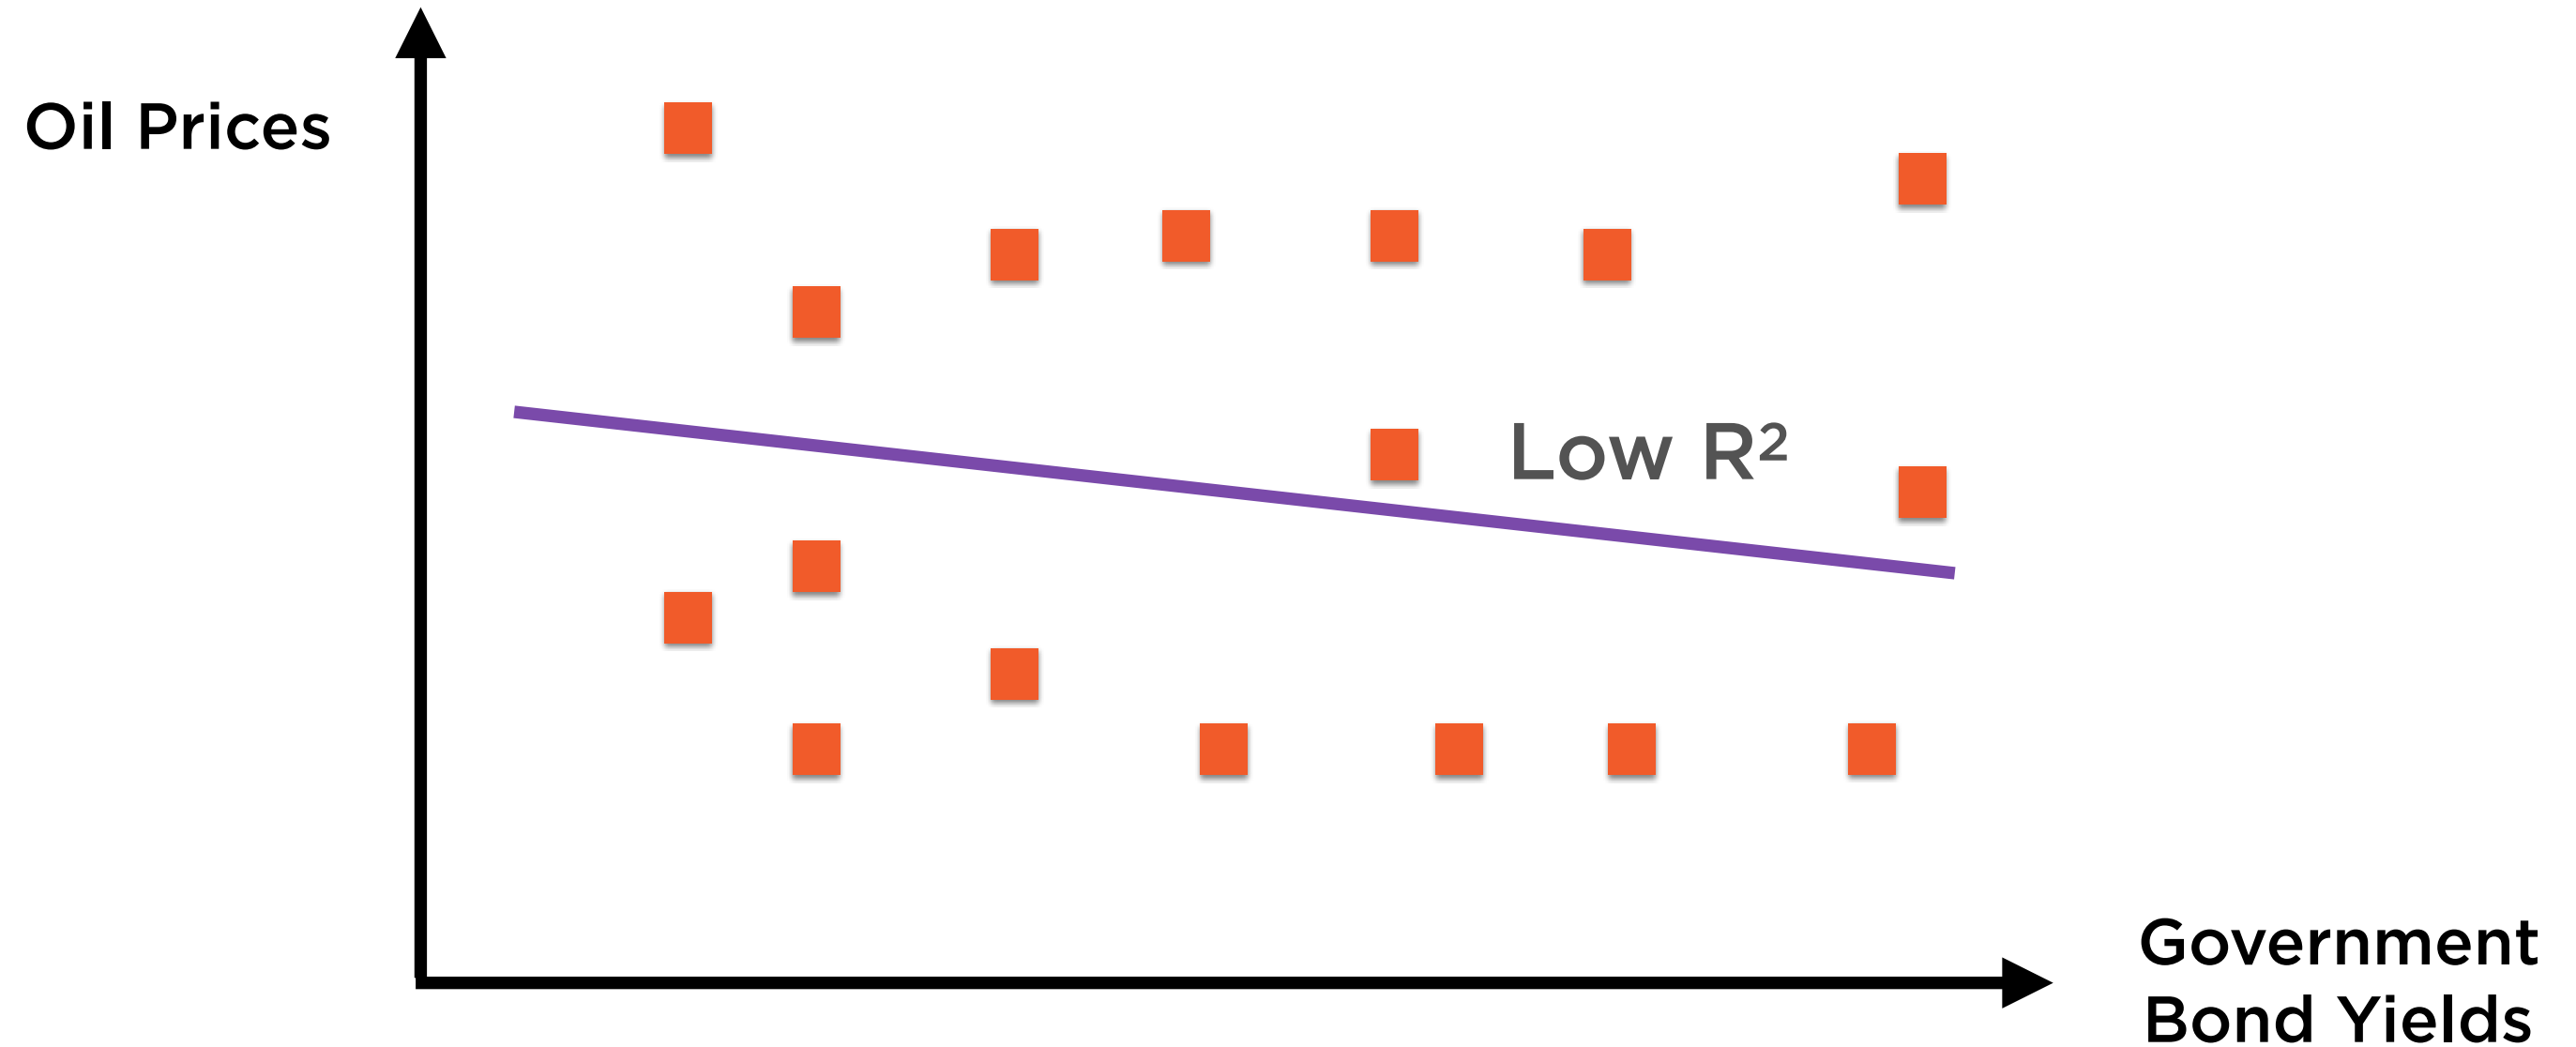
Regression not only gives us the equation of this line, it also signals how reliable the line is

# Data in Two Dimensions



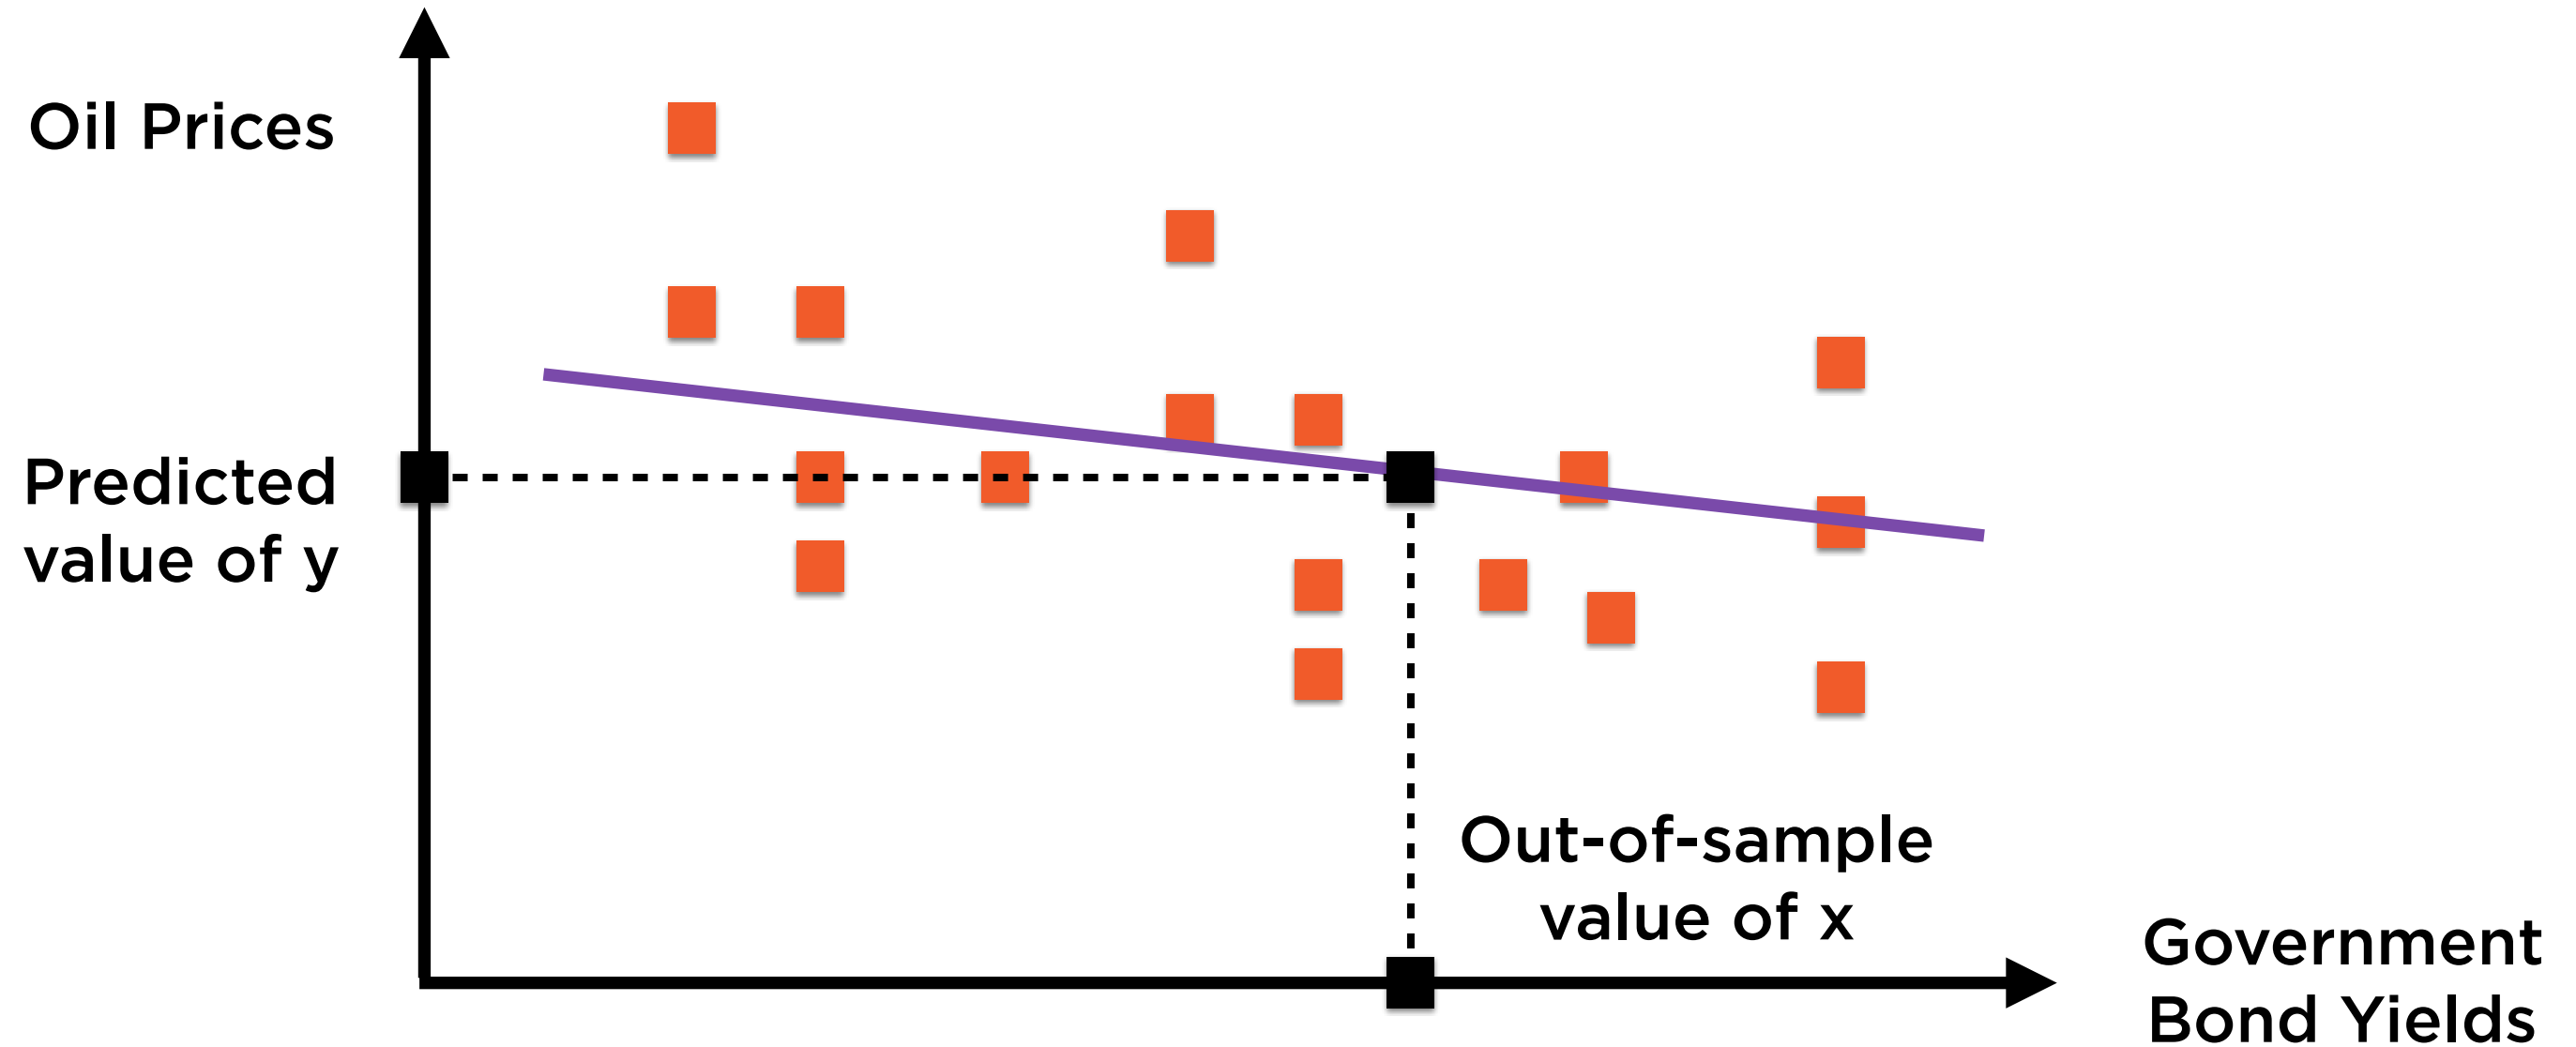
High quality of fit

# Data in Two Dimensions




Low quality of fit

# Prediction Using Regression



Given a new value of  $x$ , use the line to predict the corresponding value of  $y$

# Data in N Dimensions



Linear Regression can easily be extended to n-dimensional data

# Setting Up the Regression Problem

---

X Causes Y



**Cause**

**Independent variable**



**Effect**

**Dependent variable**

X Causes Y



**Cause**

**Explanatory variable**

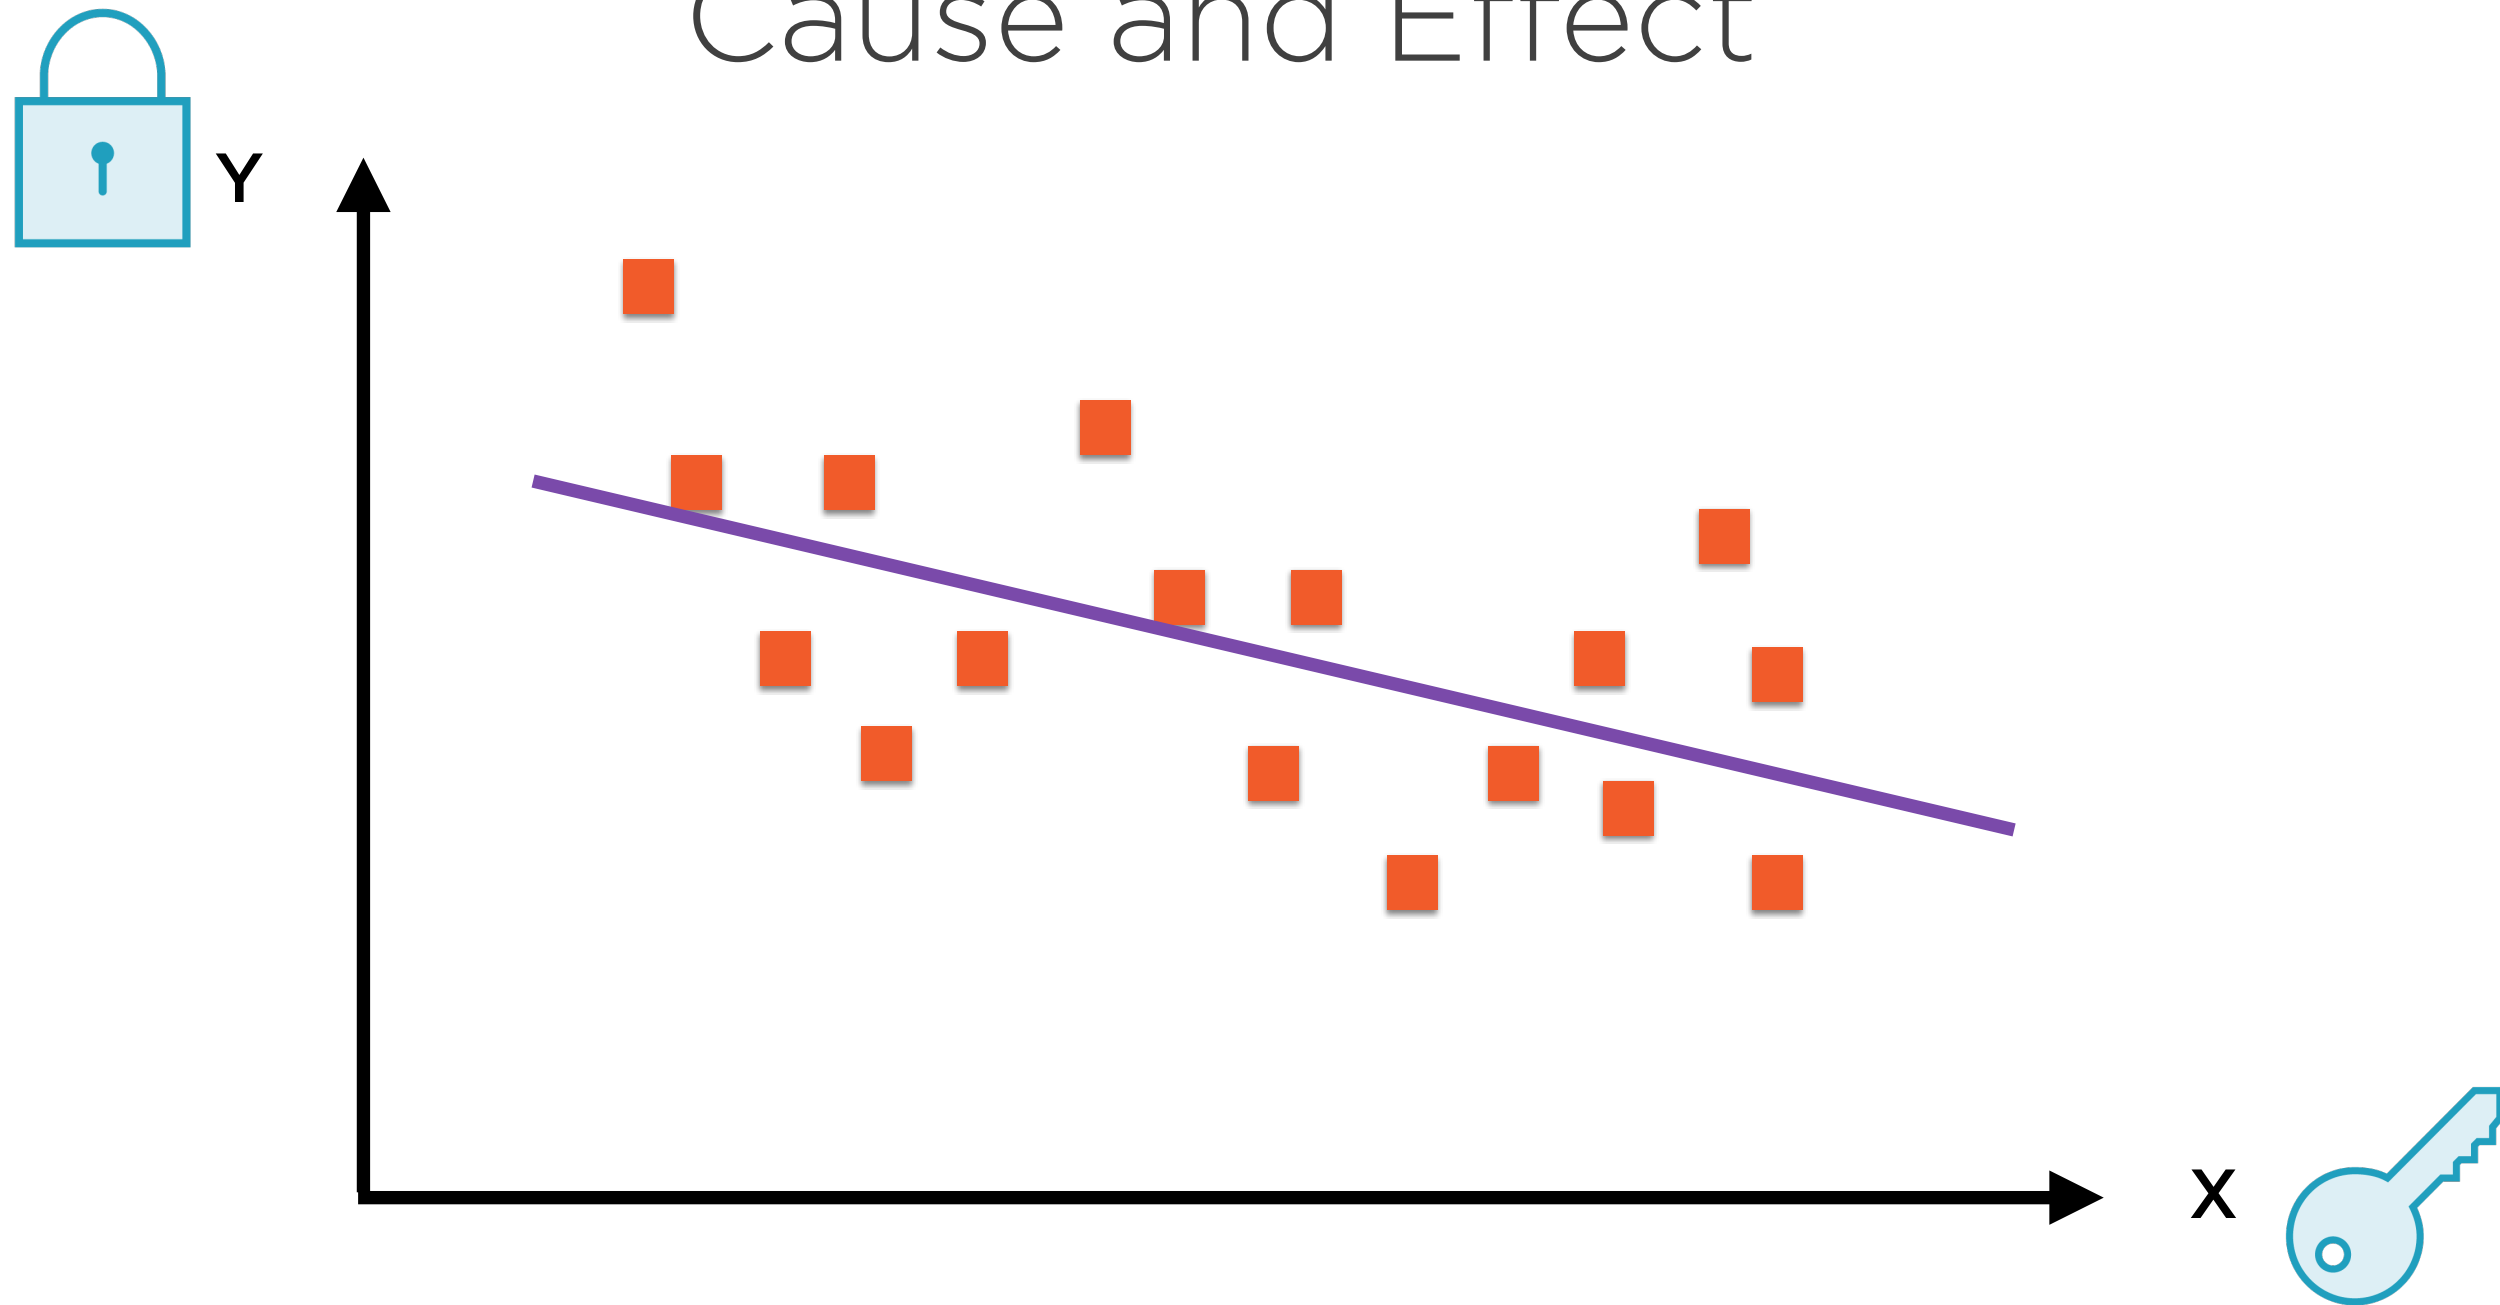


**Effect**

**Dependent variable**



# Cause and Effect

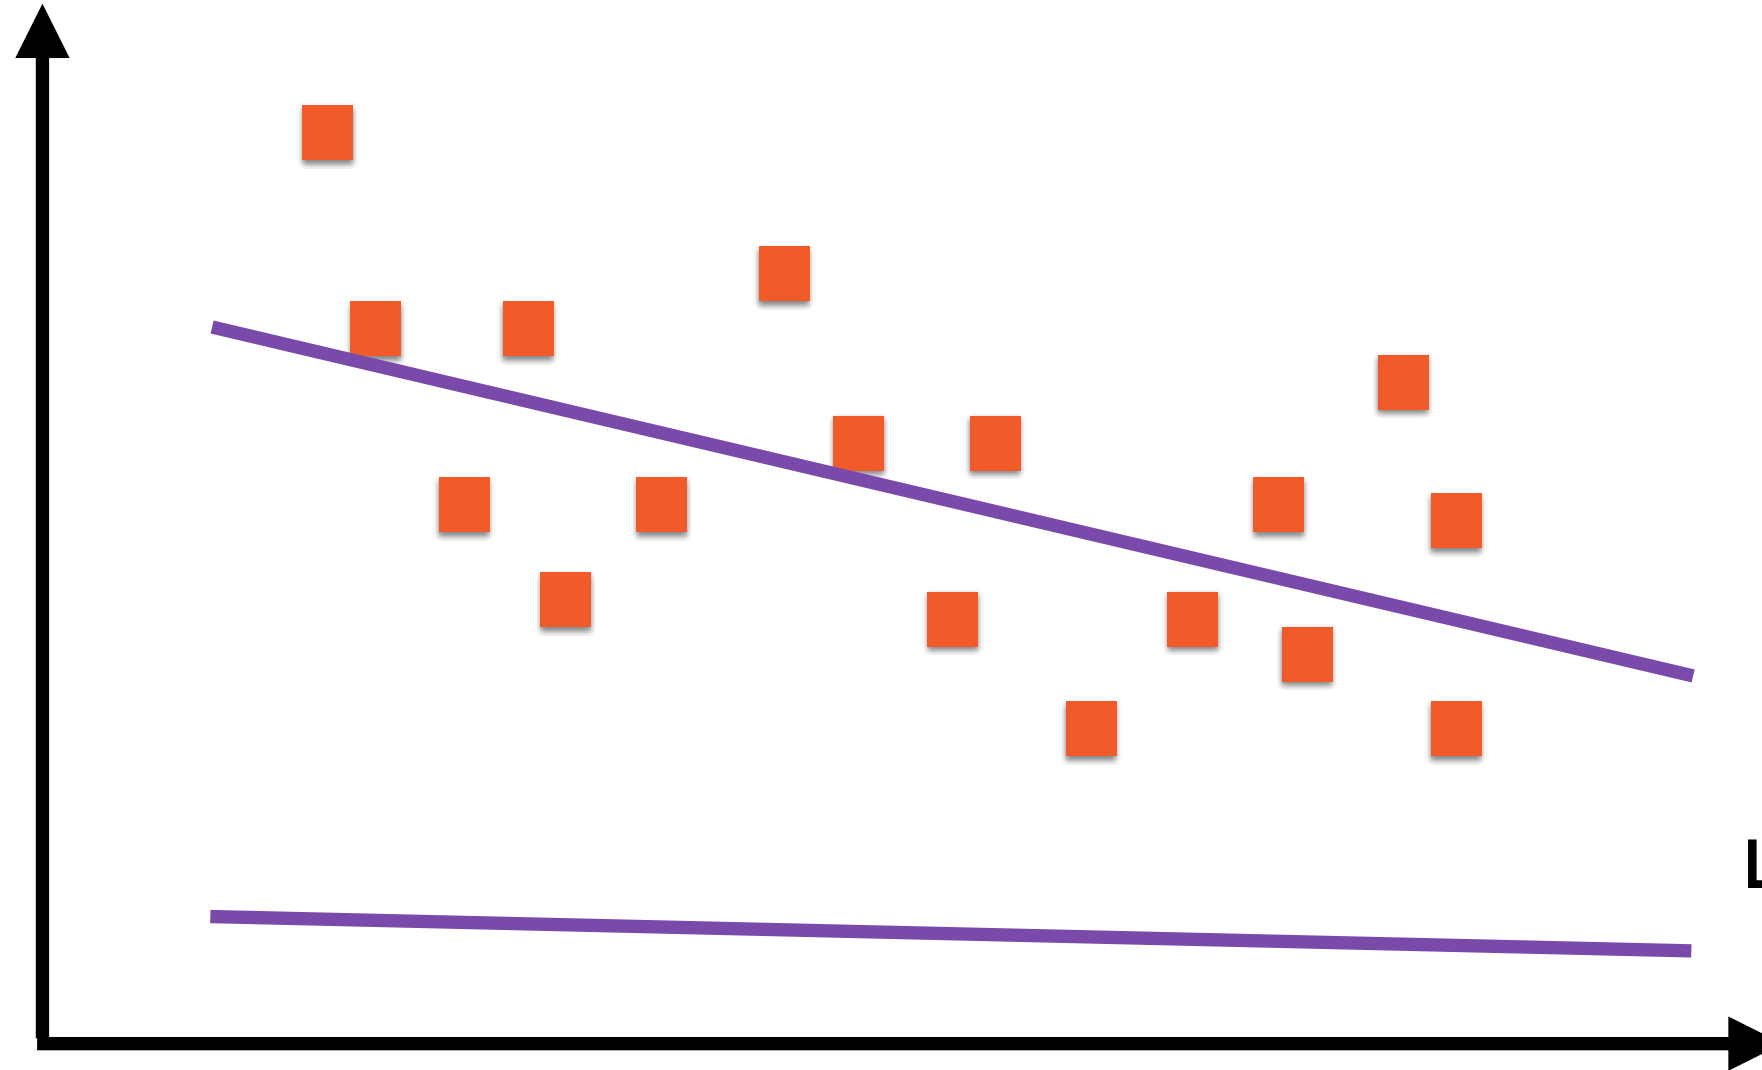


Linear Regression involves finding the “best fit” line

# Cause and Effect



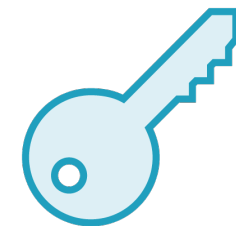
Y



Line 1:  $y = A_1 + B_1x$

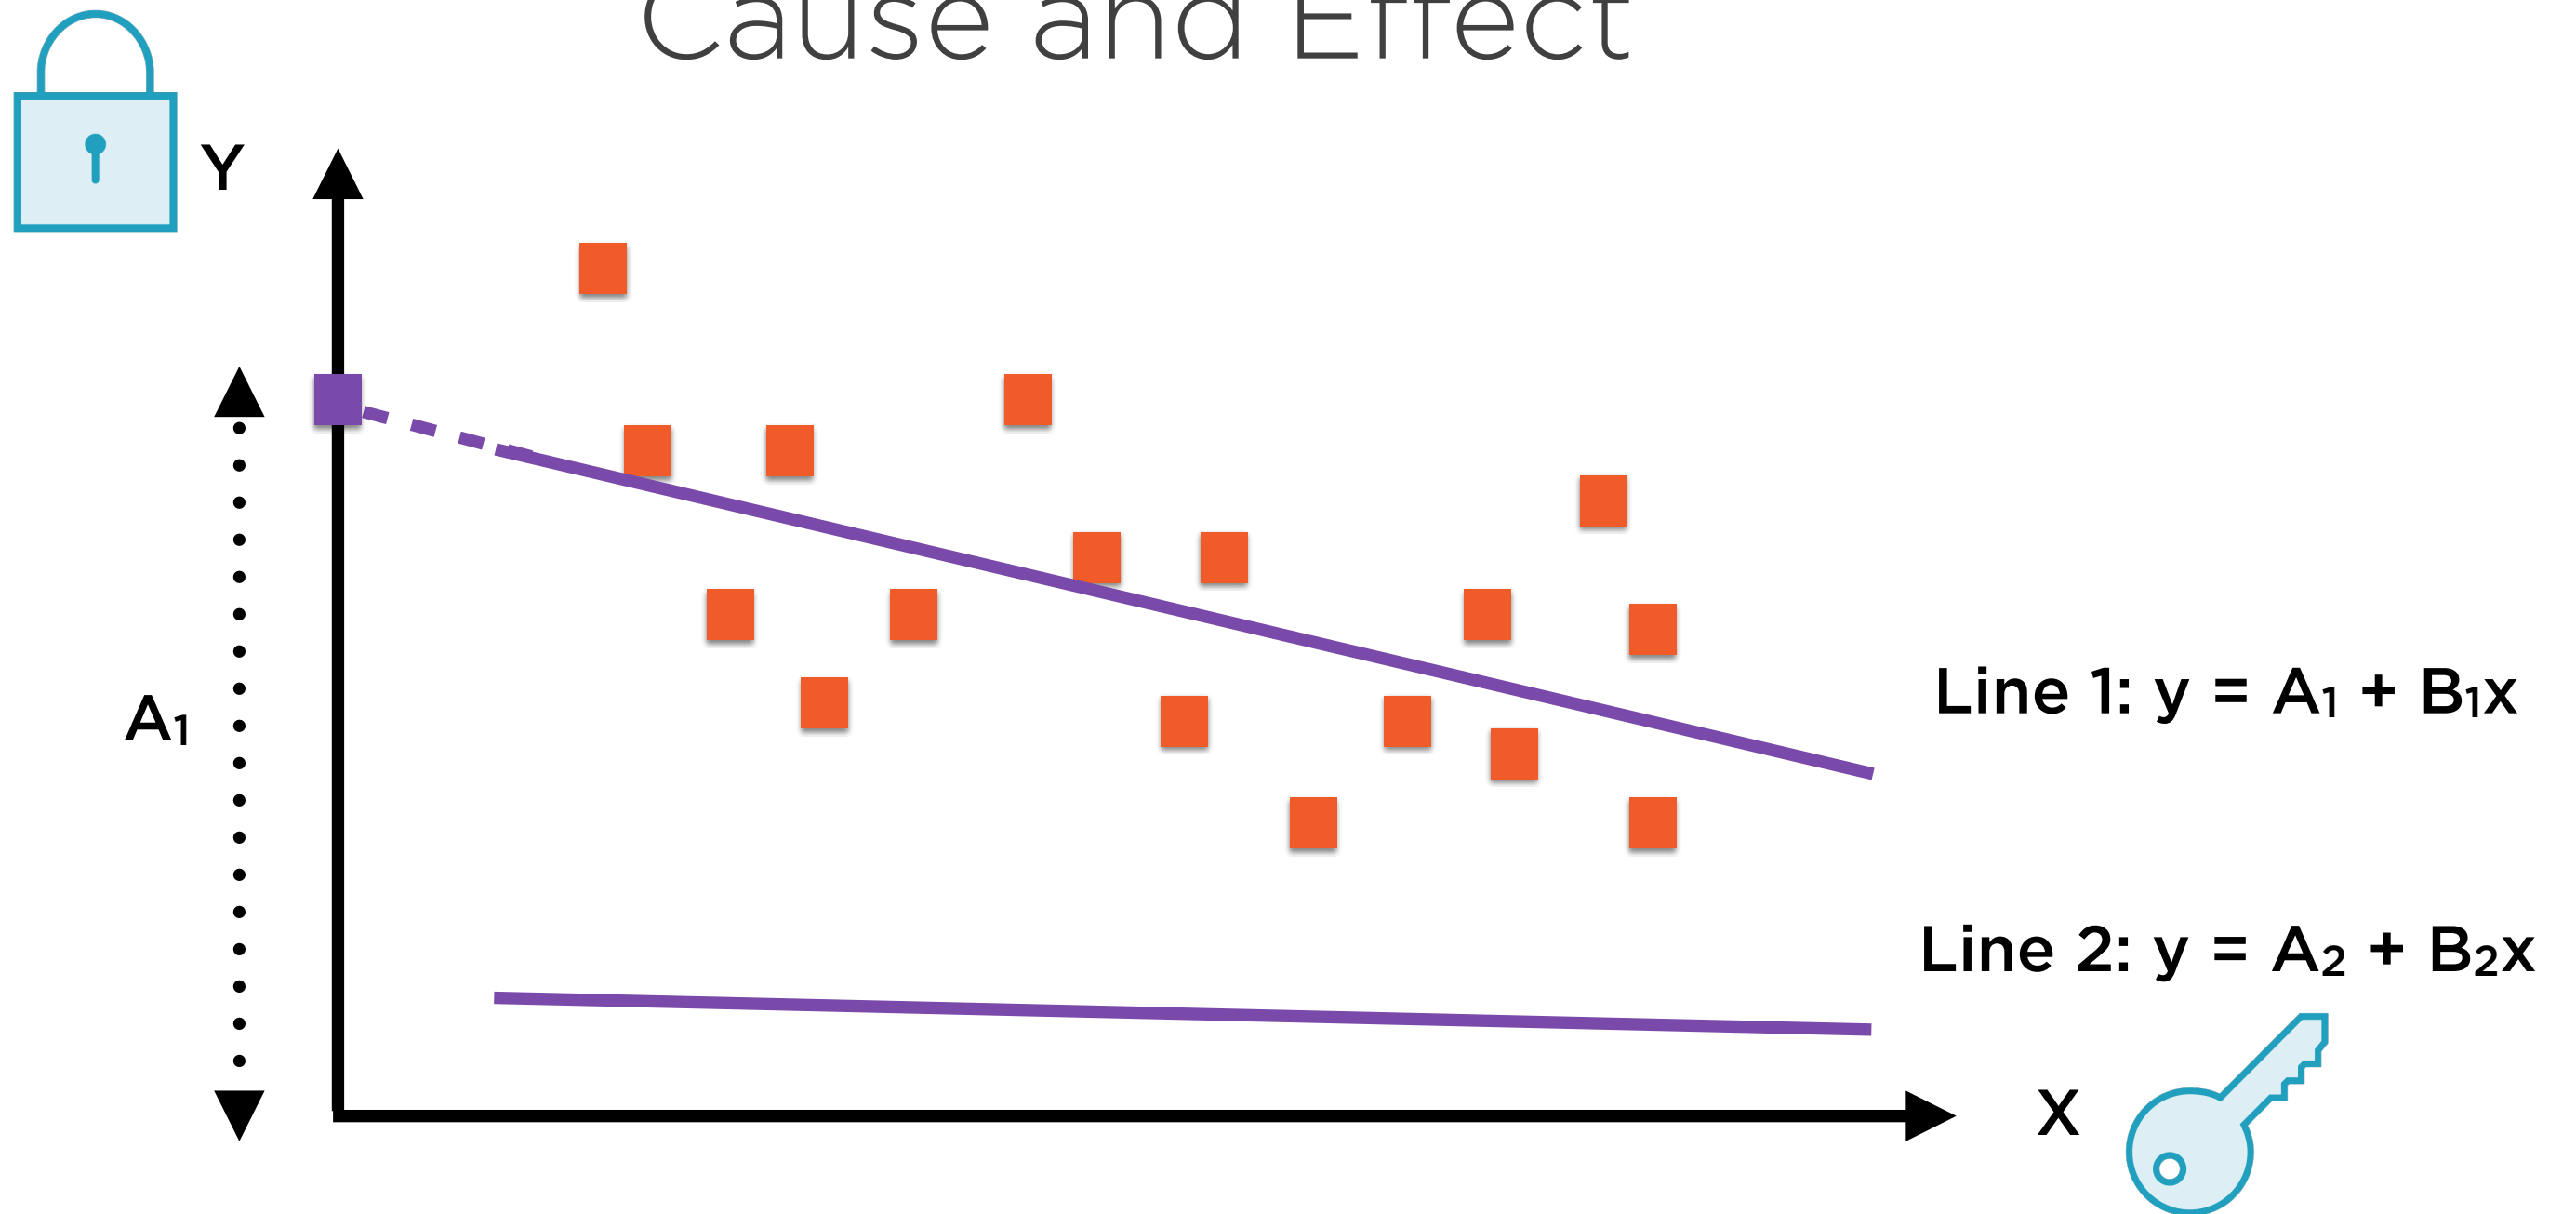
Line 2:  $y = A_2 + B_2x$

X



Let's compare two lines, Line 1 and Line 2

# Cause and Effect




The first line has y-intercept  $A_1$

# Cause and Effect



Y

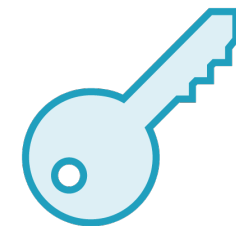
$A_2$



Line 1:  $y = A_1 + B_1x$

Line 2:  $y = A_2 + B_2x$

X

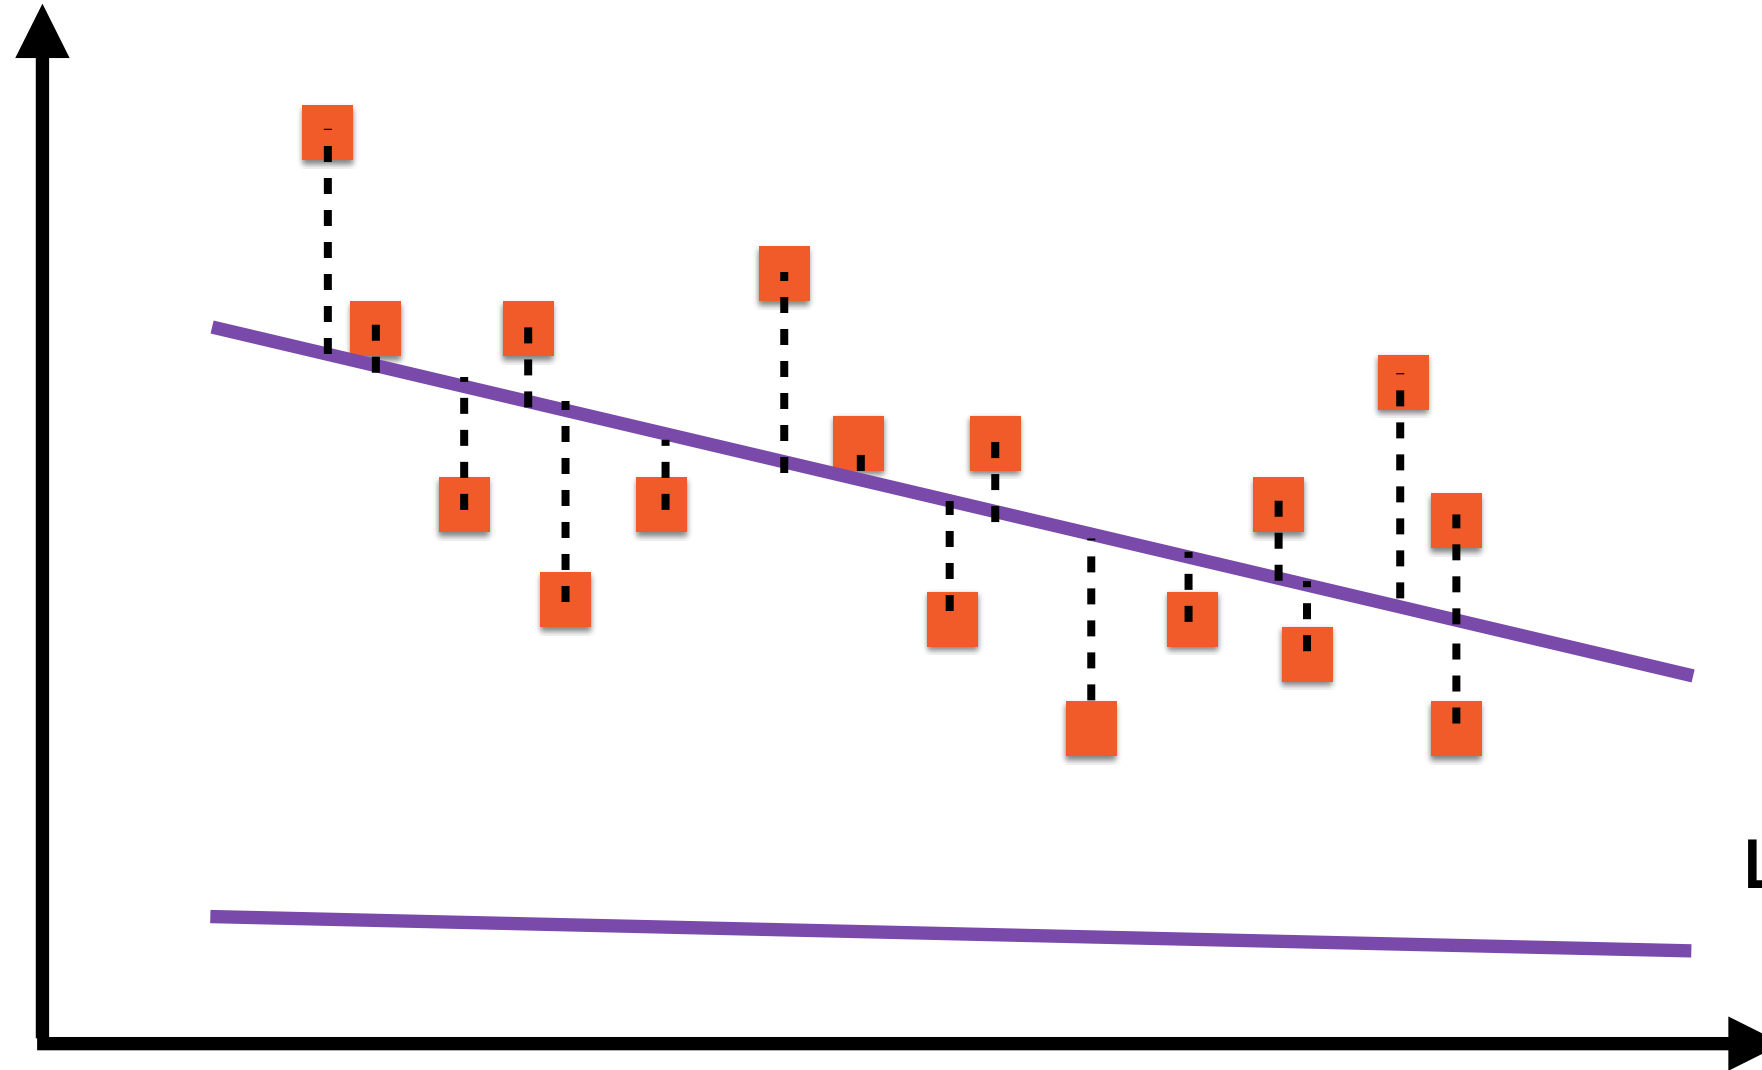


The second line has y-intercept  $A_2$

# Minimizing Least Square Error




Y



Line 1:  $y = A_1 + B_1x$

Line 2:  $y = A_2 + B_2x$


X



# Minimizing Least Square Error




Y



Line 1:  $y = A_1 + B_1x$

Line 2:  $y = A_2 + B_2x$

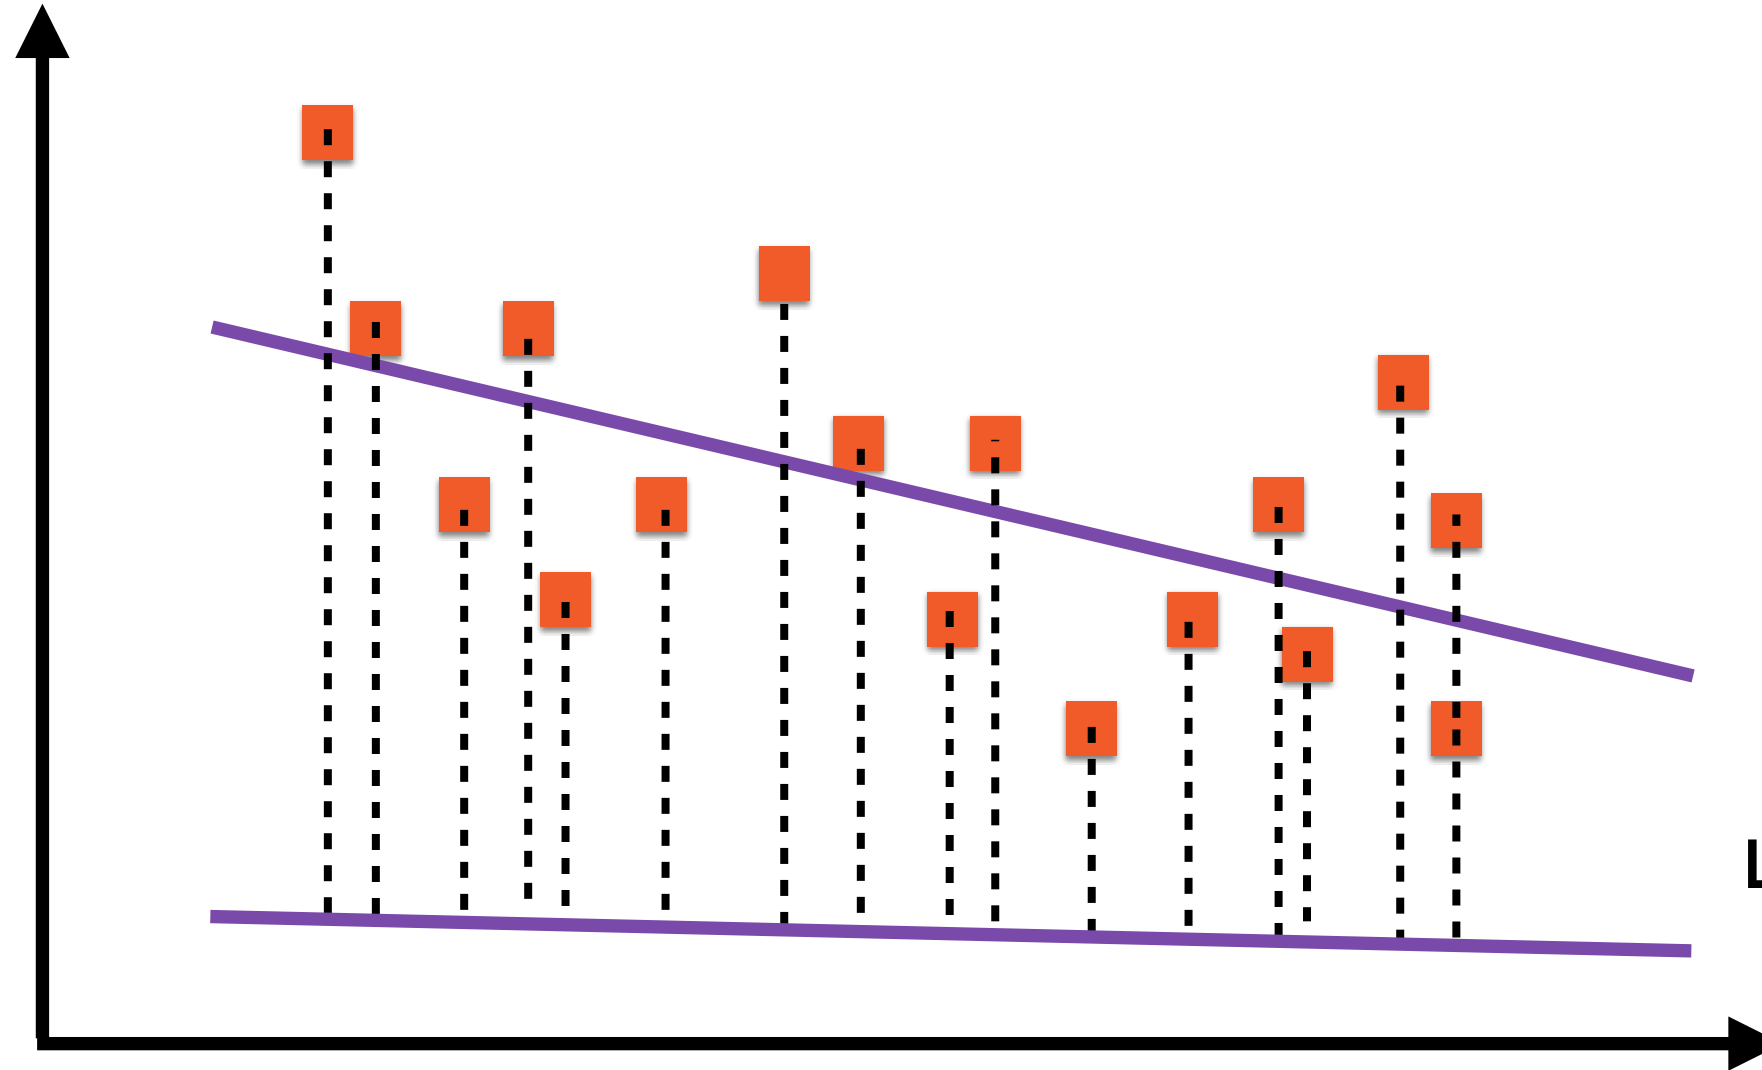
X



# Minimizing Least Square Error



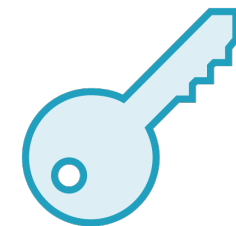
Y



Line 1:  $y = A_1 + B_1x$

Line 2:  $y = A_2 + B_2x$

X

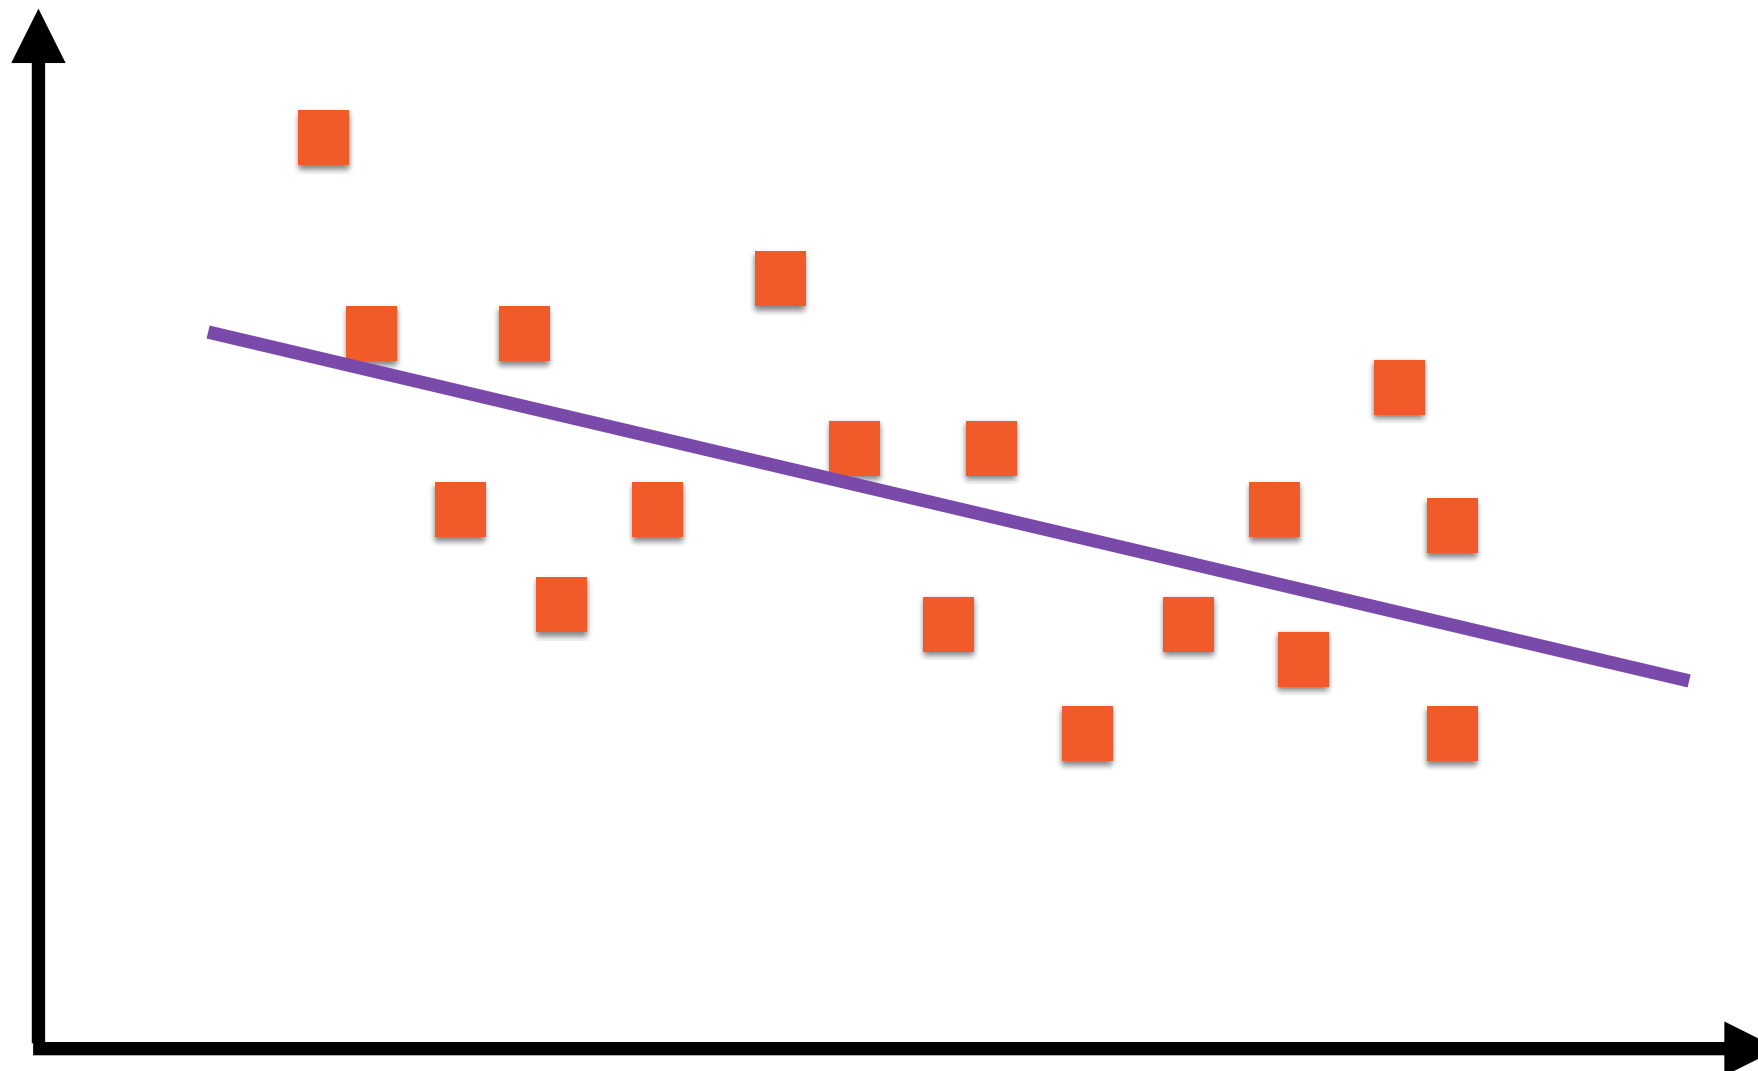


The “best fit” line is the one where the sum of the squares of the lengths of the errors is minimum

# Best-fit Line

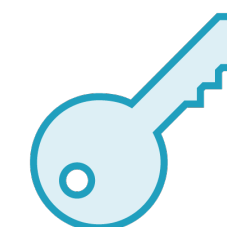


Y



Line 1:  $y = A_1 + B_1x$

X



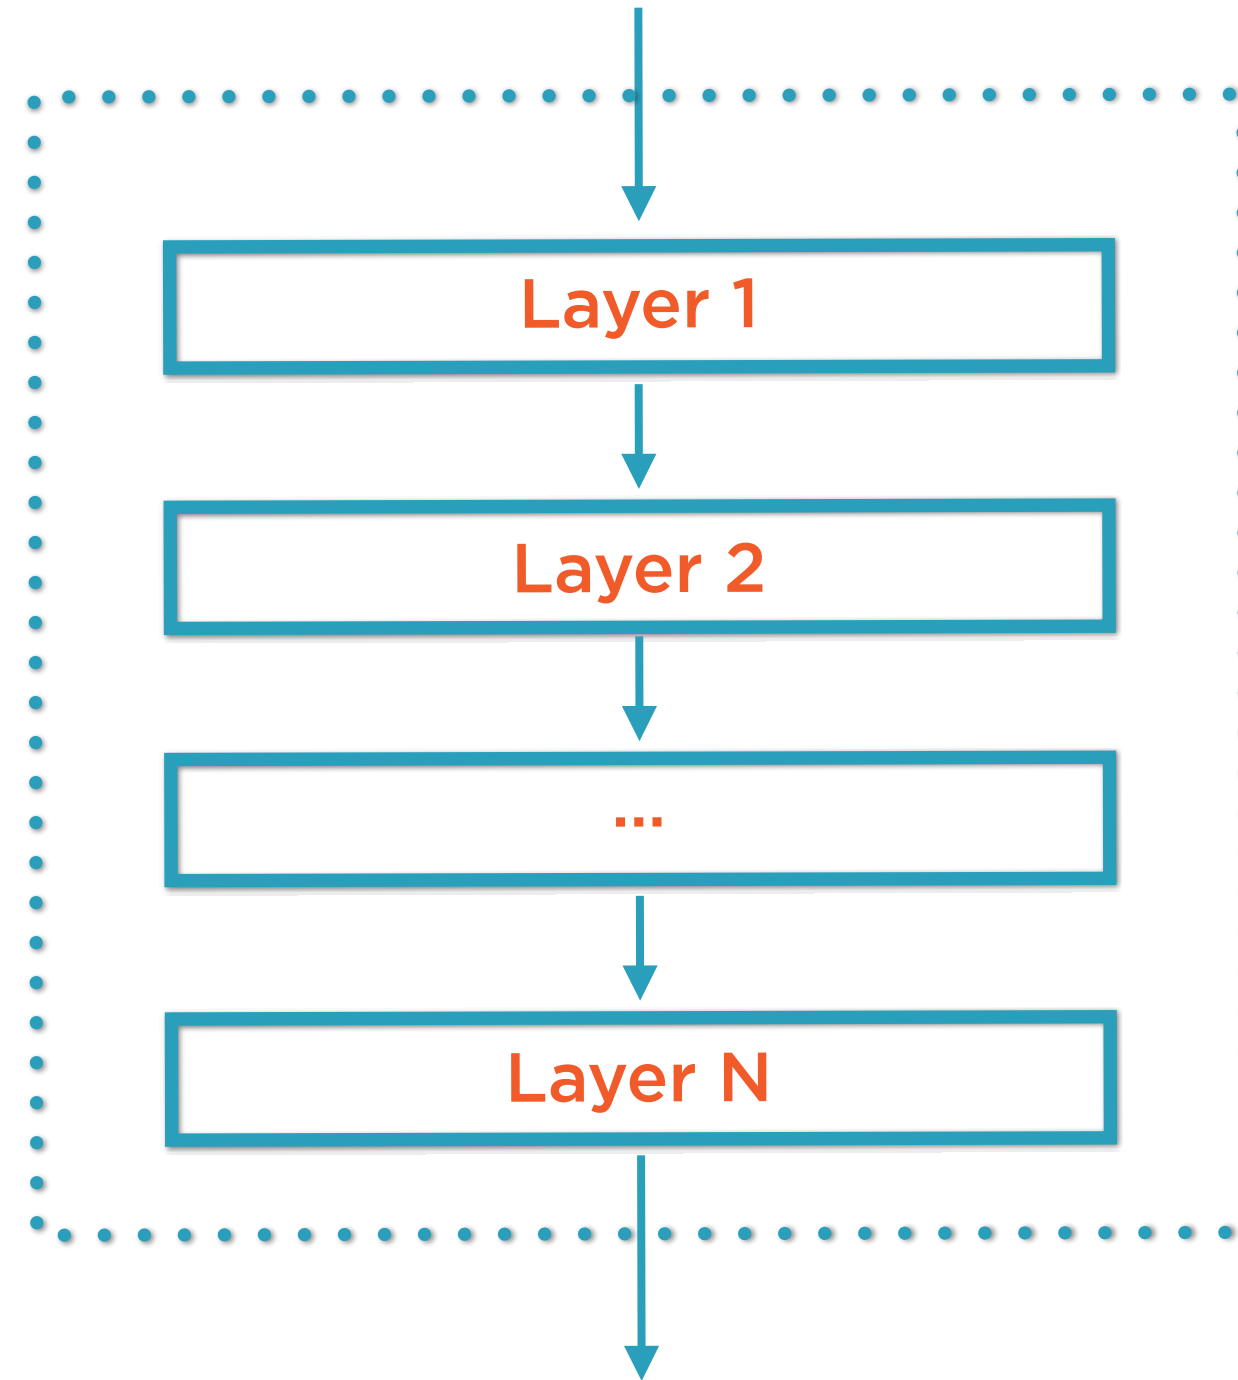
The “best fit” line is the one where the sum of the squares of the lengths of these dotted lines is minimum



# Gradient Descent

---

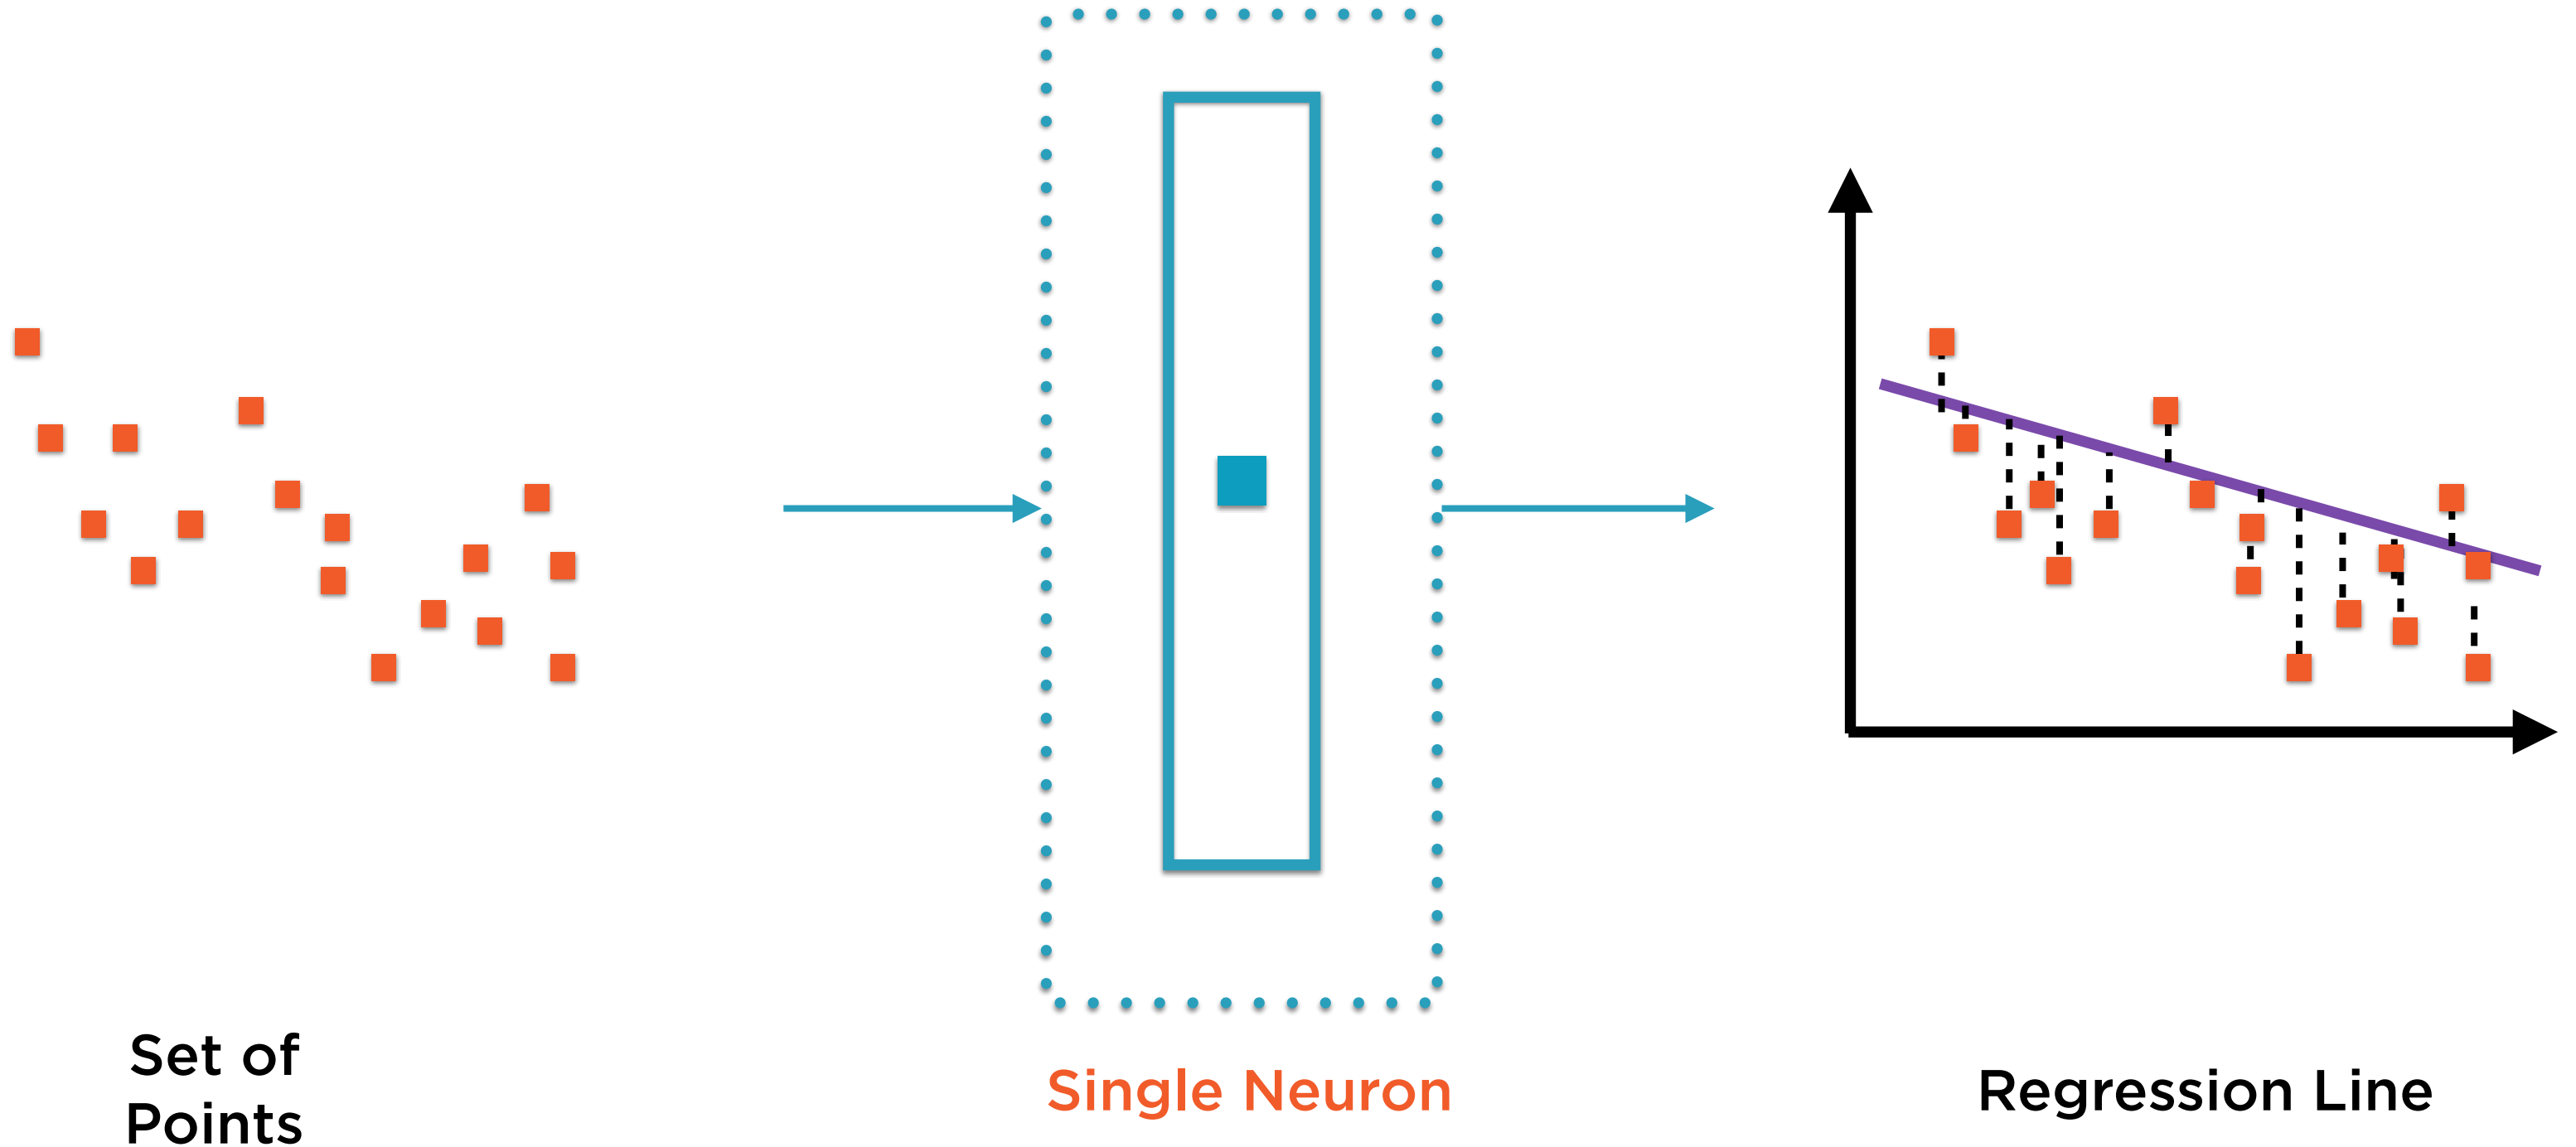
# Neural Network Model



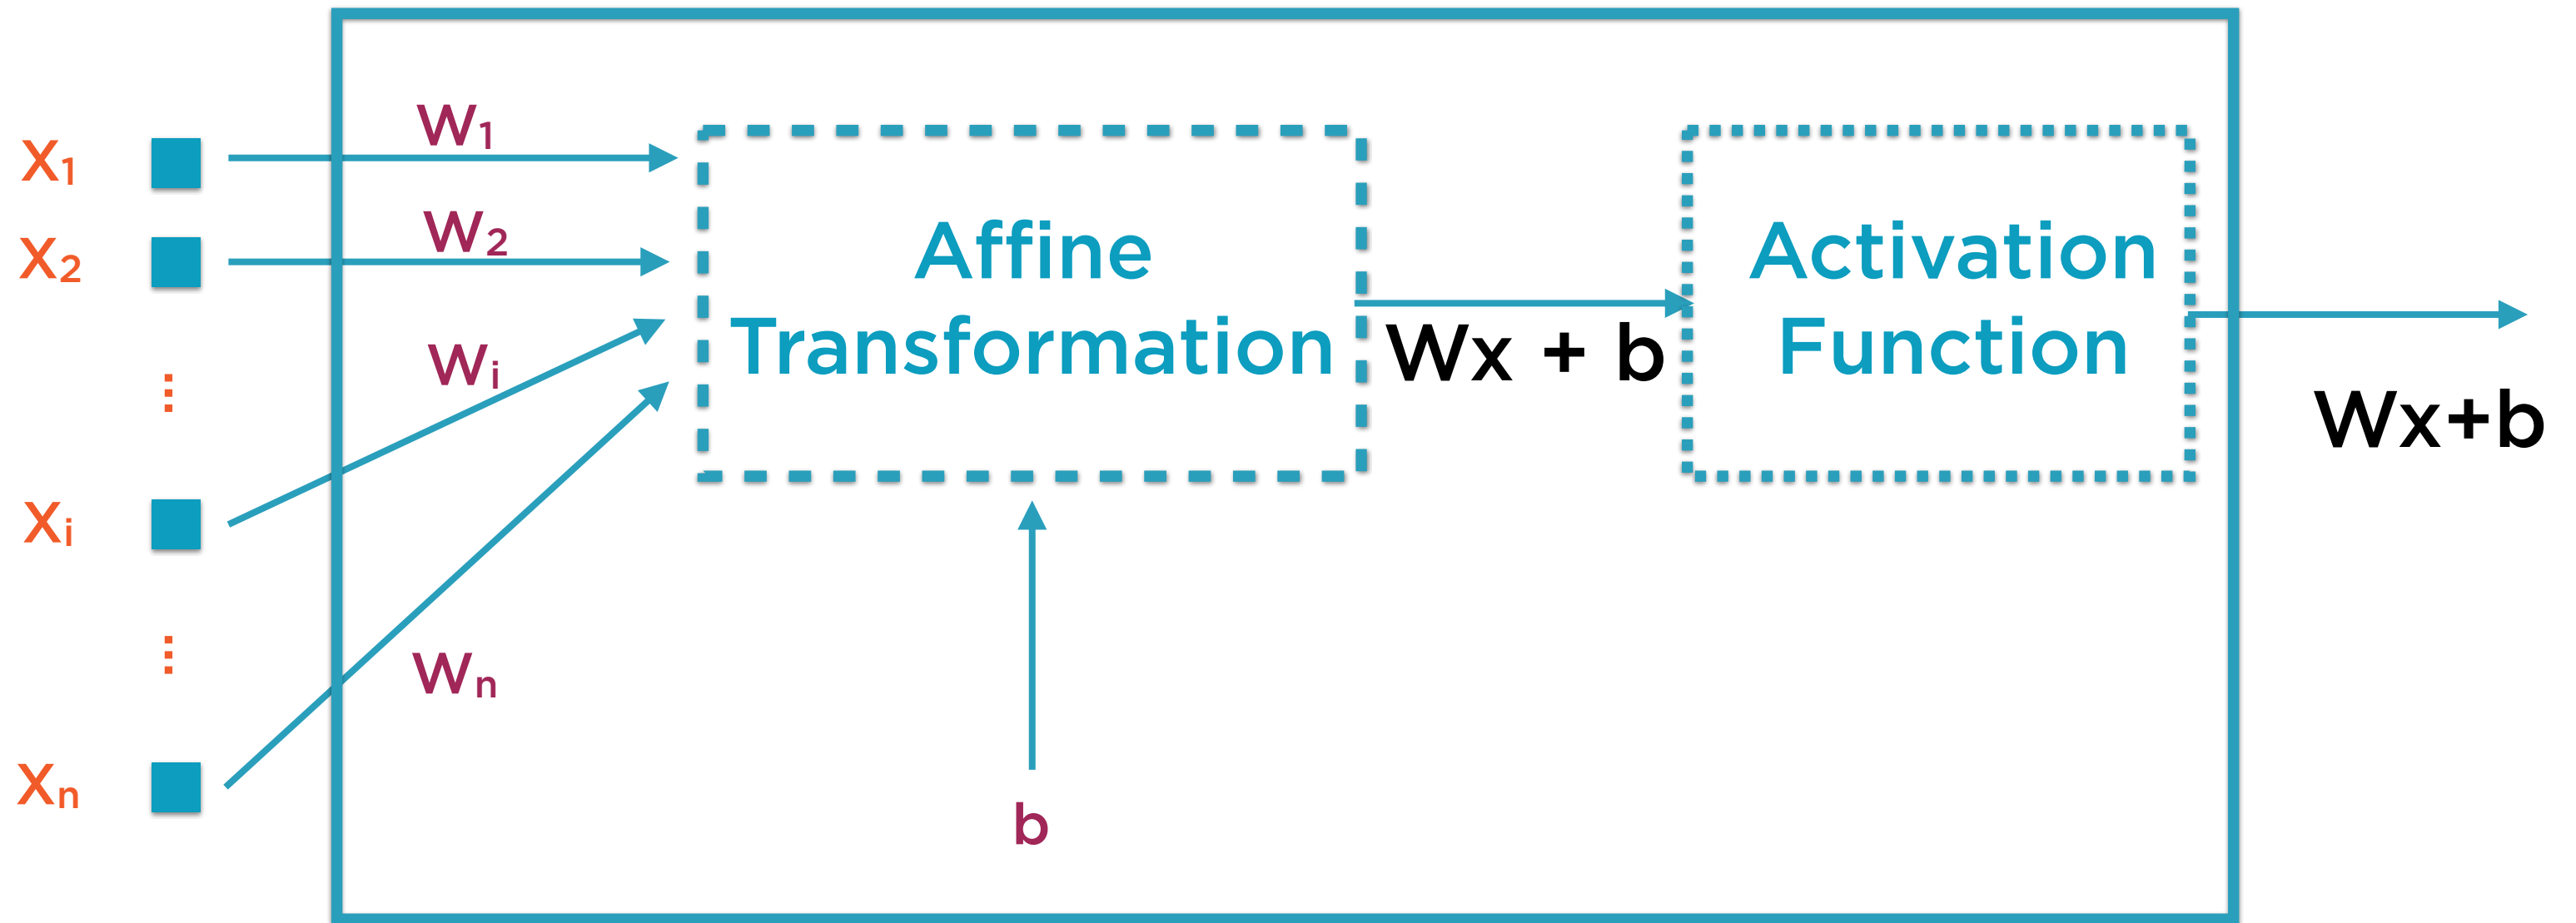
**Network of interconnected layers**

The **weights** and **biases** of individual neurons are determined during the **training** process

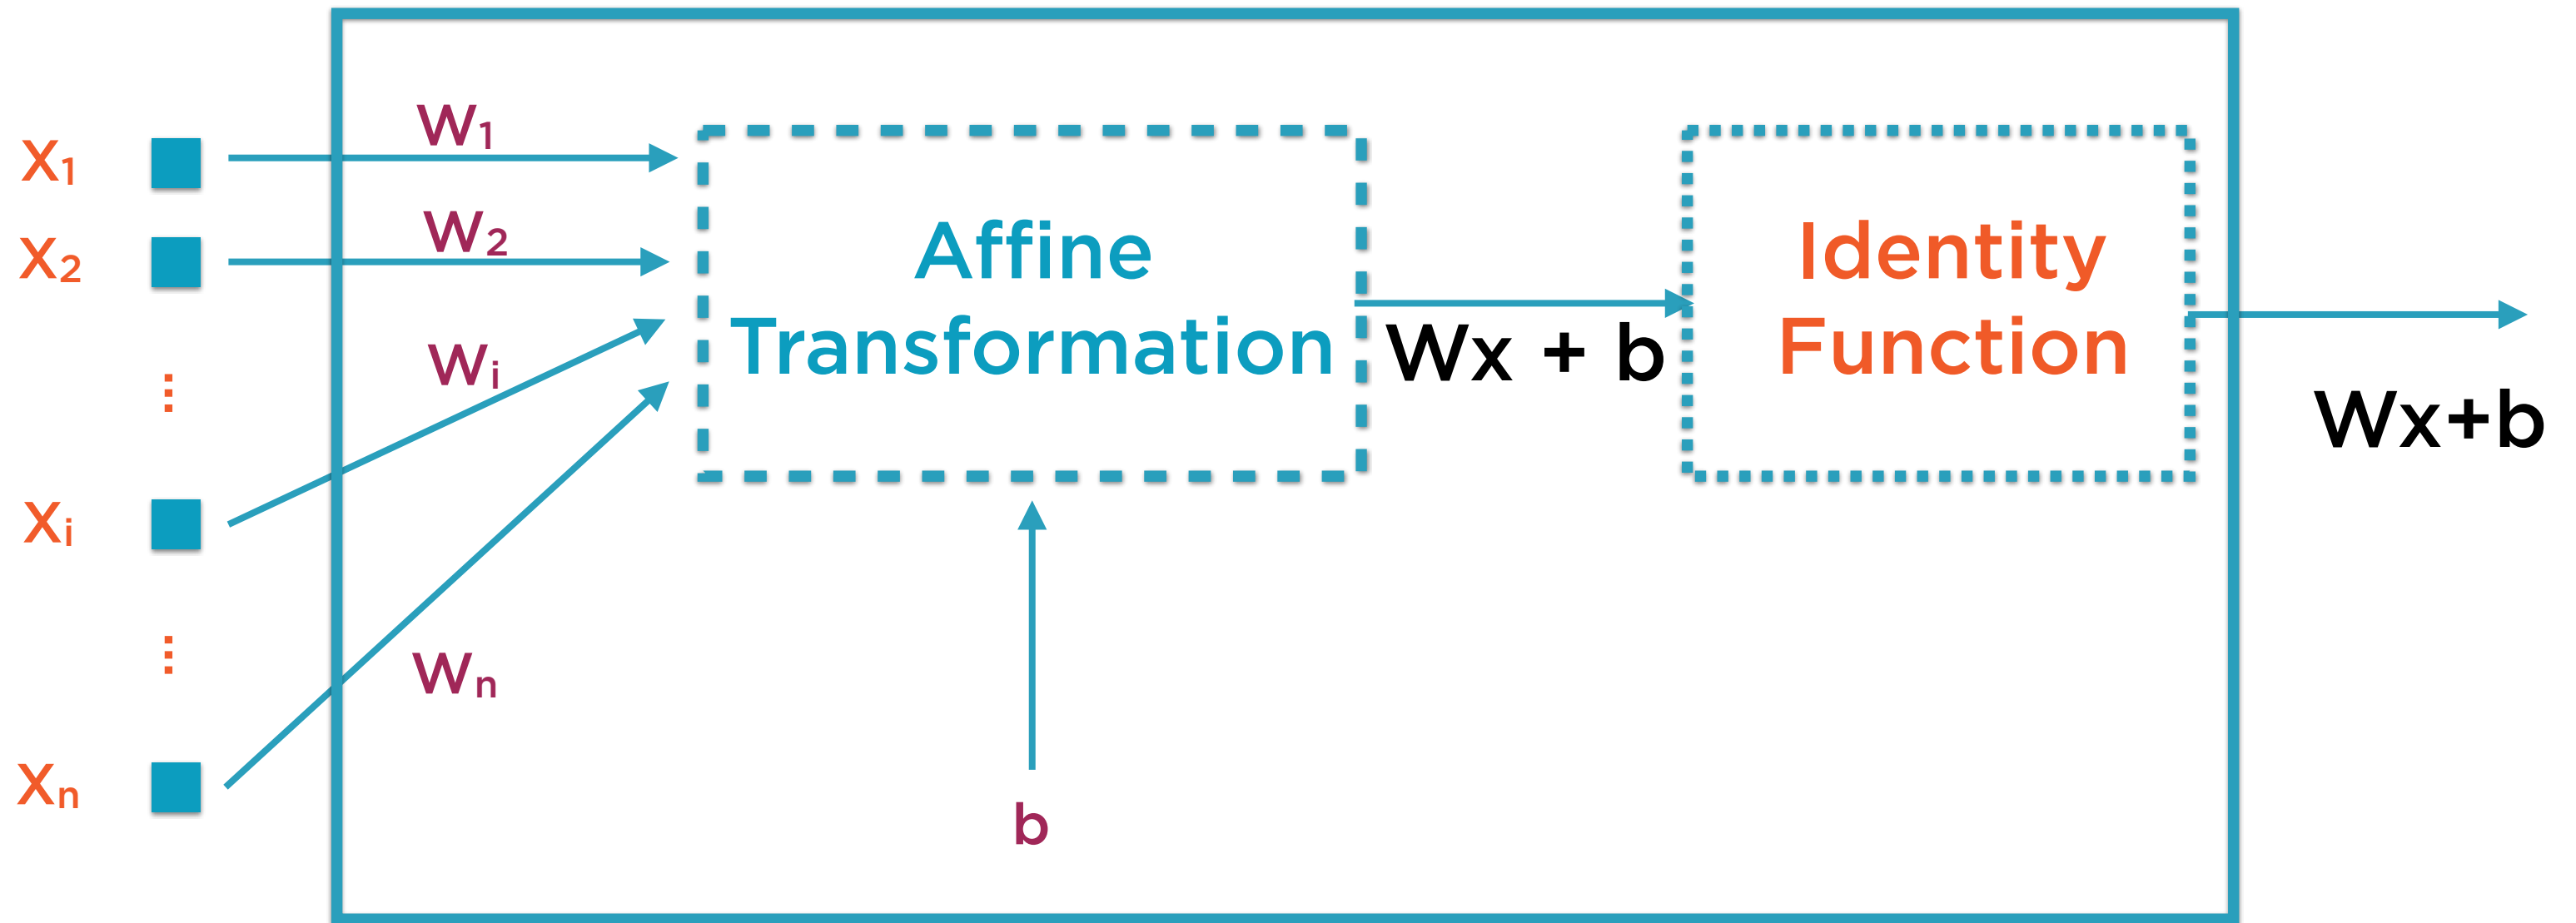
# Regression: The Simplest Neural Network




# Regression: The Simplest Neural Network



# Regression: The Simplest Neural Network



# Best-fit Line



The “best fit” line is the one where the sum of the squares of the lengths of these dotted lines is minimum

The actual training of a neural network happens via Gradient Descent Optimization



# Linear Regression as an Optimization Problem



## **Objective Function**

**Minimize variance of  
the residuals (MSE)**

# Linear Regression as an Optimization Problem



## Objective Function

Minimize variance of  
the residuals (MSE)



## Constraints

Express relationship as  
a straight line

$$y = Wx + b$$

# Linear Regression as an Optimization Problem



## Objective Function


Minimize variance of  
the residuals (MSE)



## Constraints

Express relationship as  
a straight line


$$y = Wx + b$$



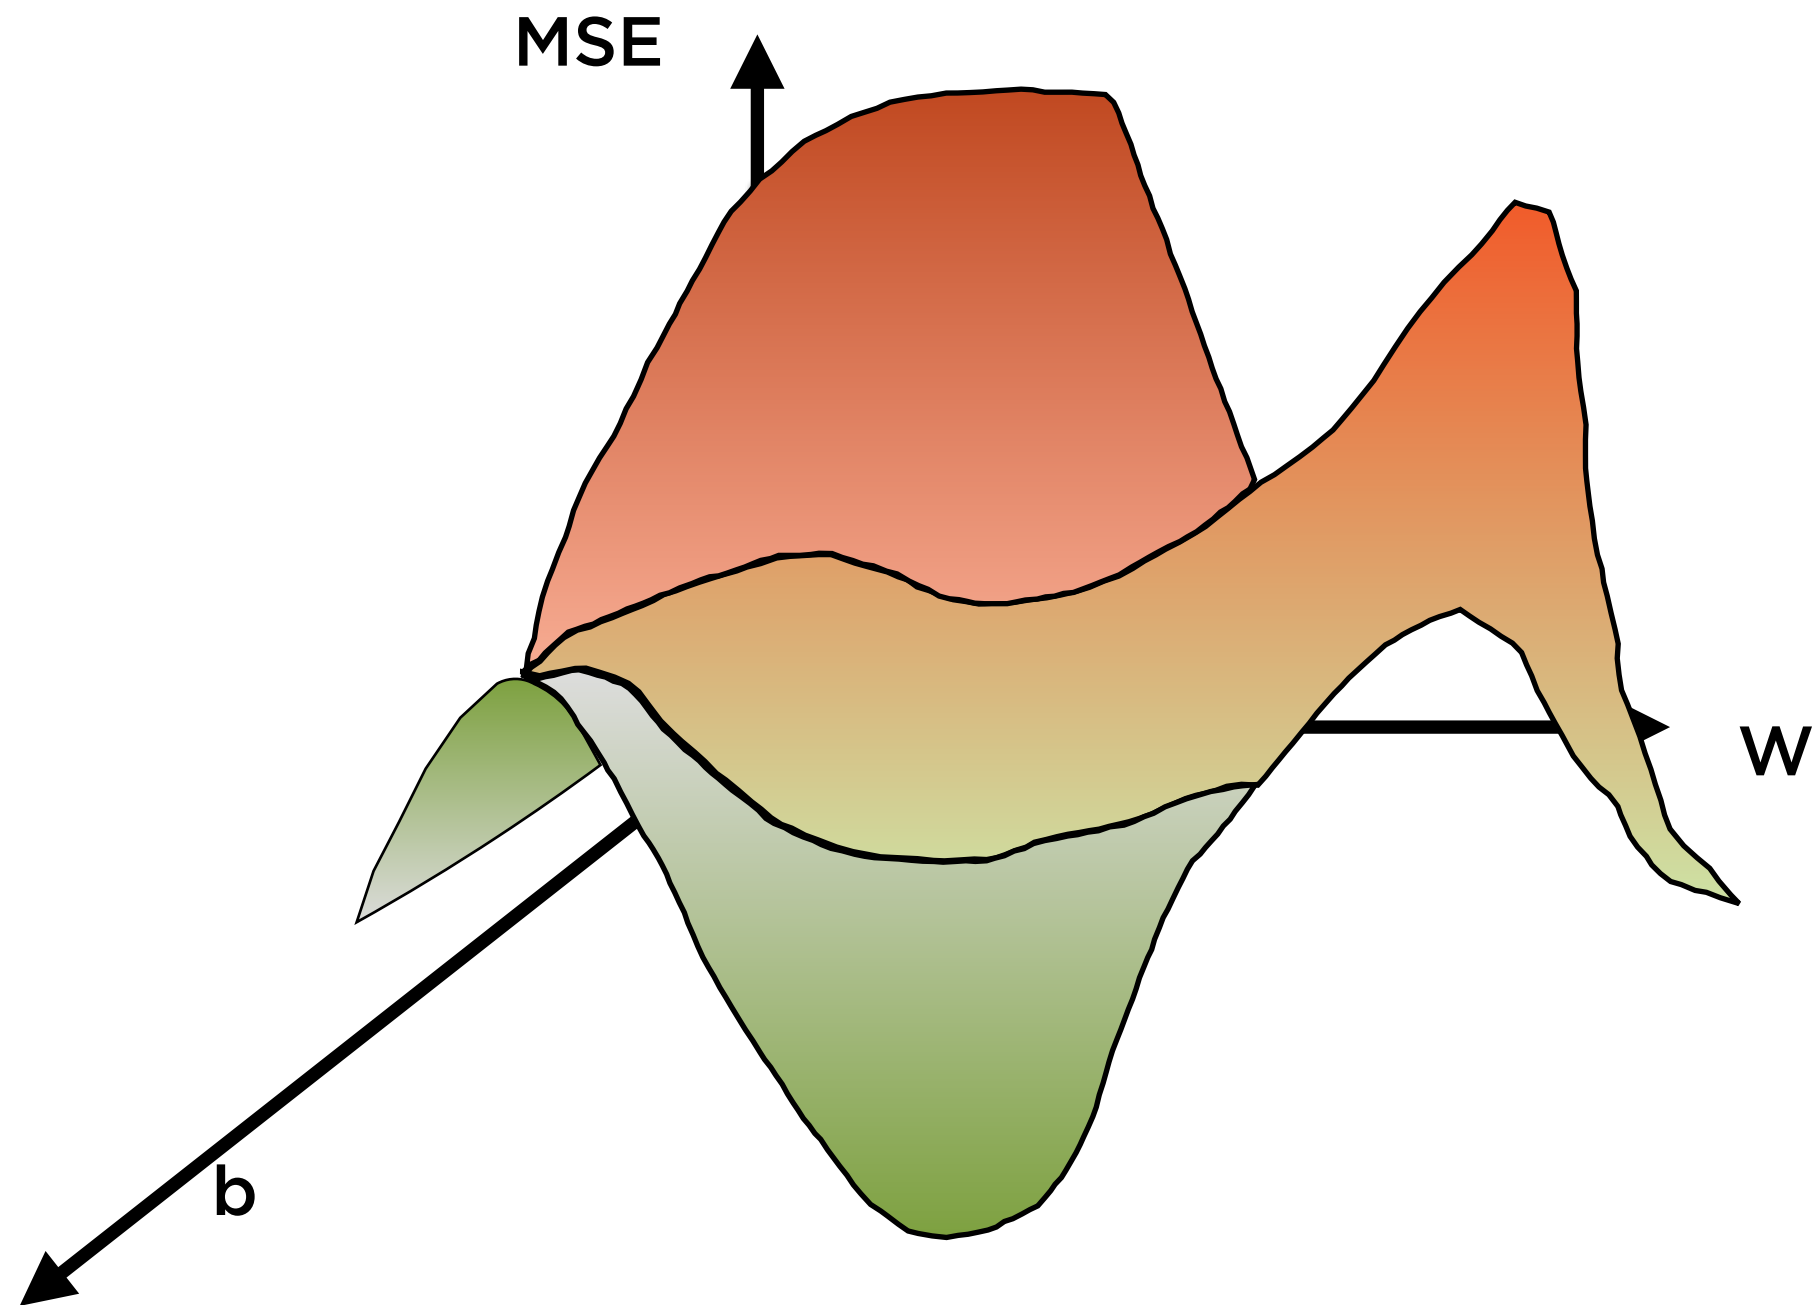
## Decision Variables

Values of  $W$  and  $b$


# Minimizing MSE



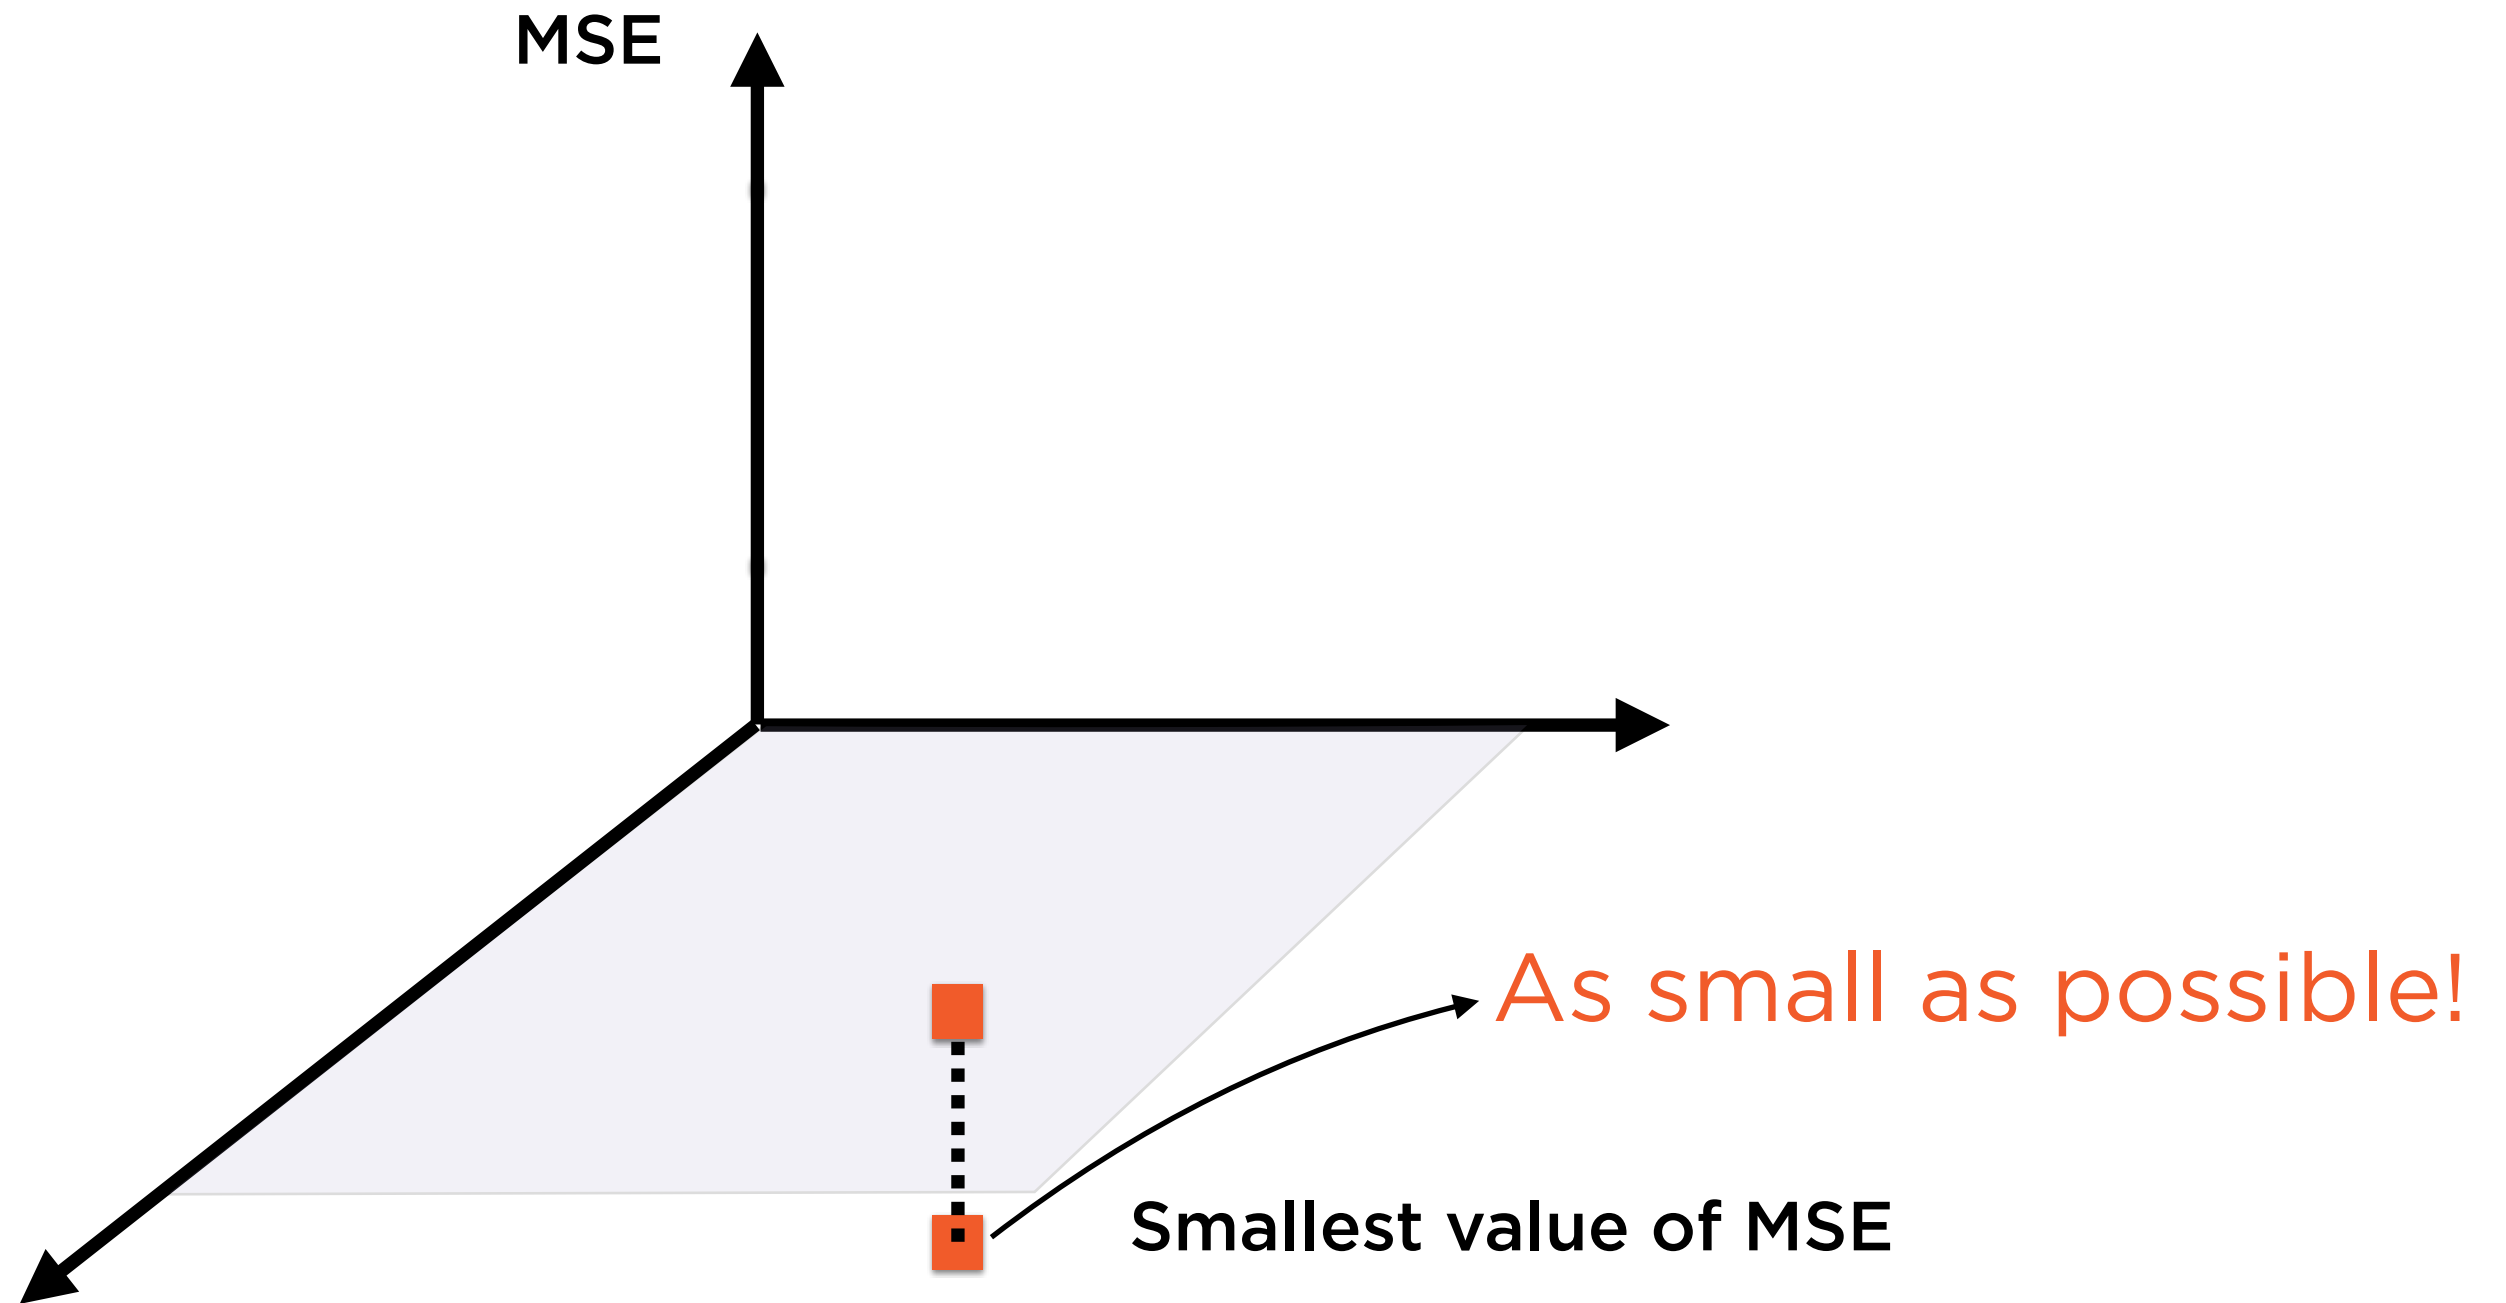
# Minimizing MSE



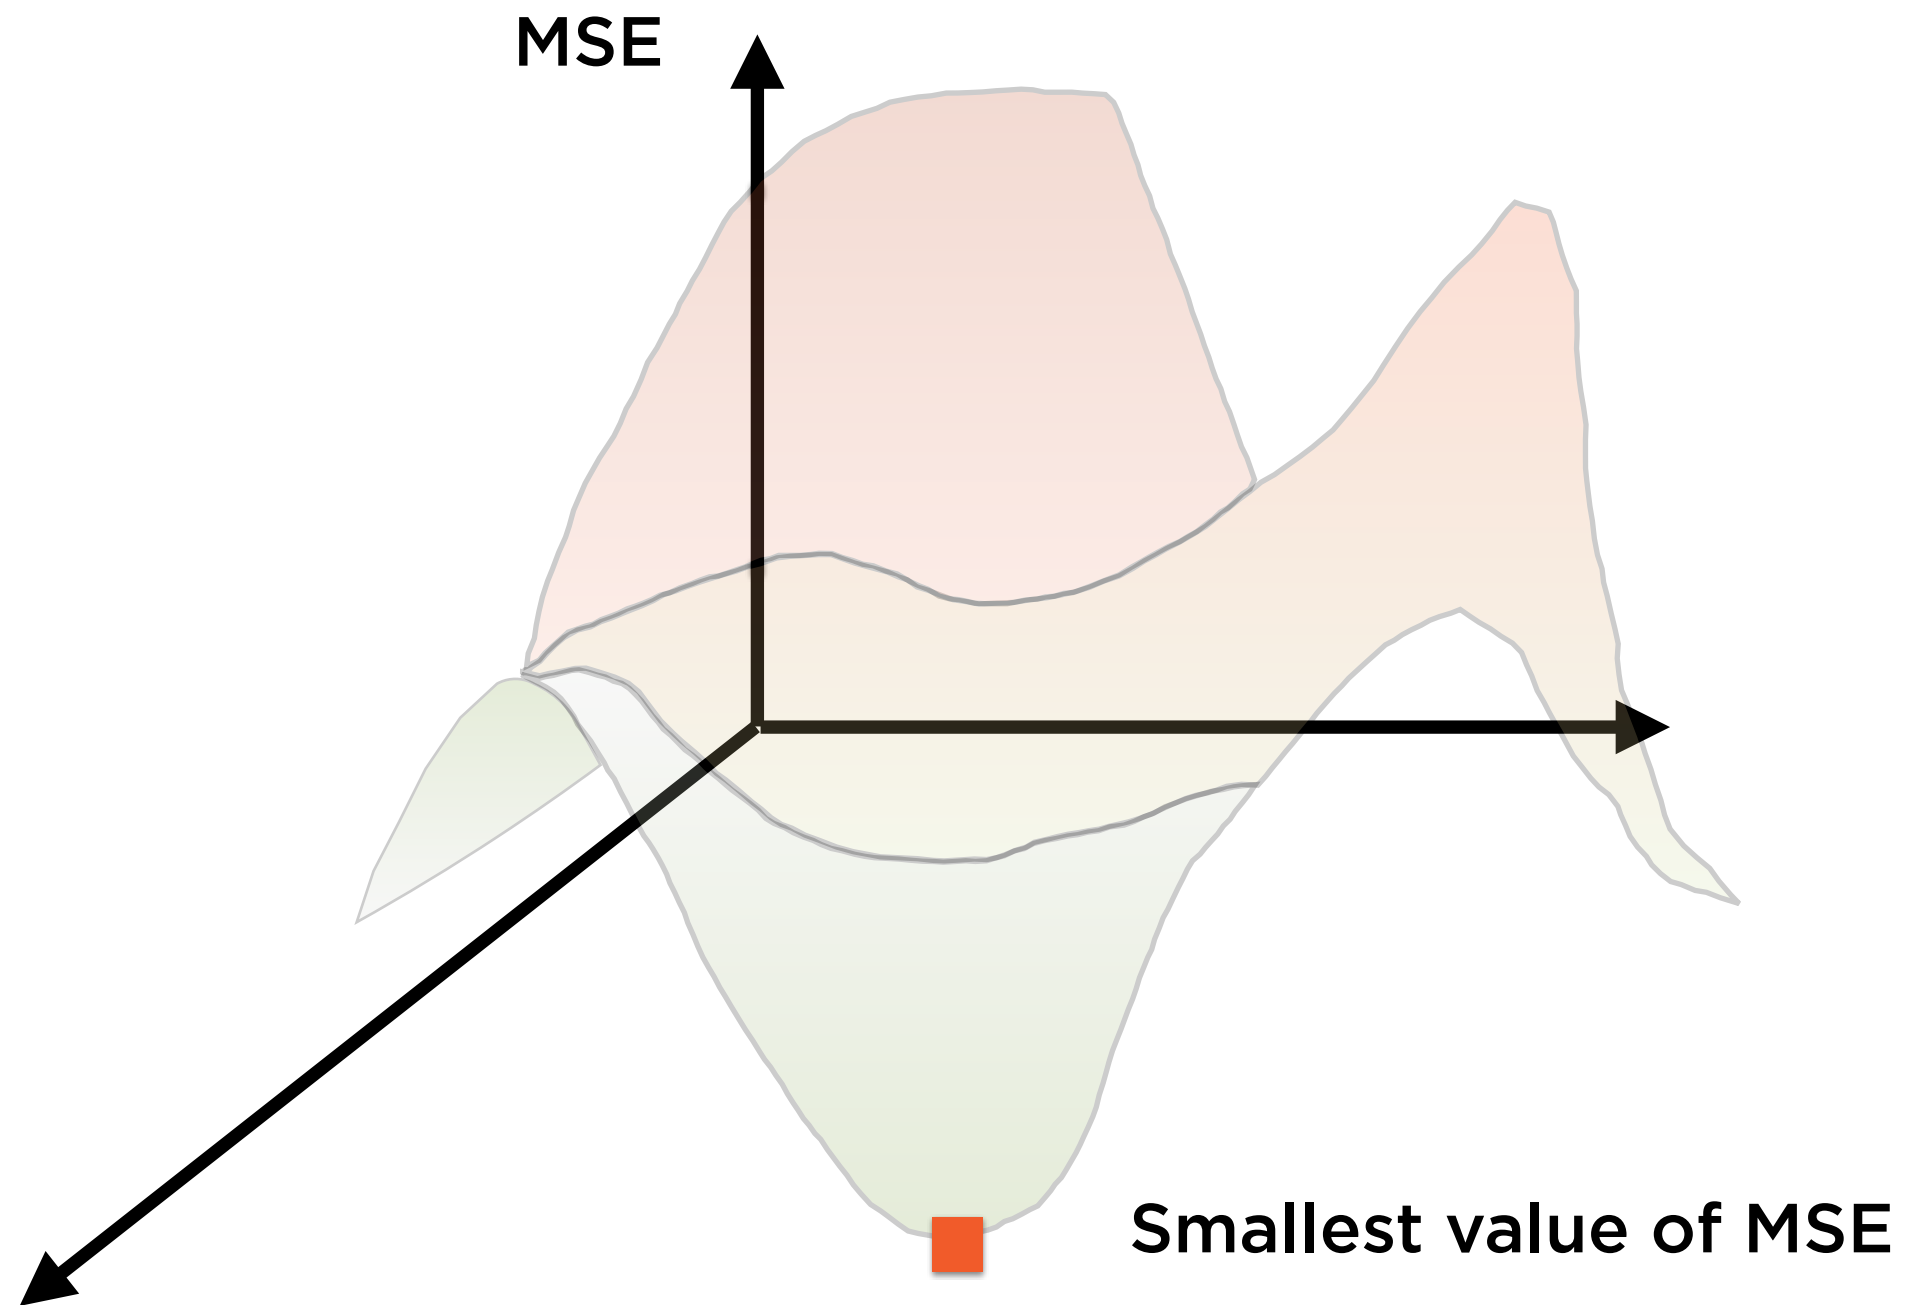
# Minimizing MSE



# Minimizing MSE

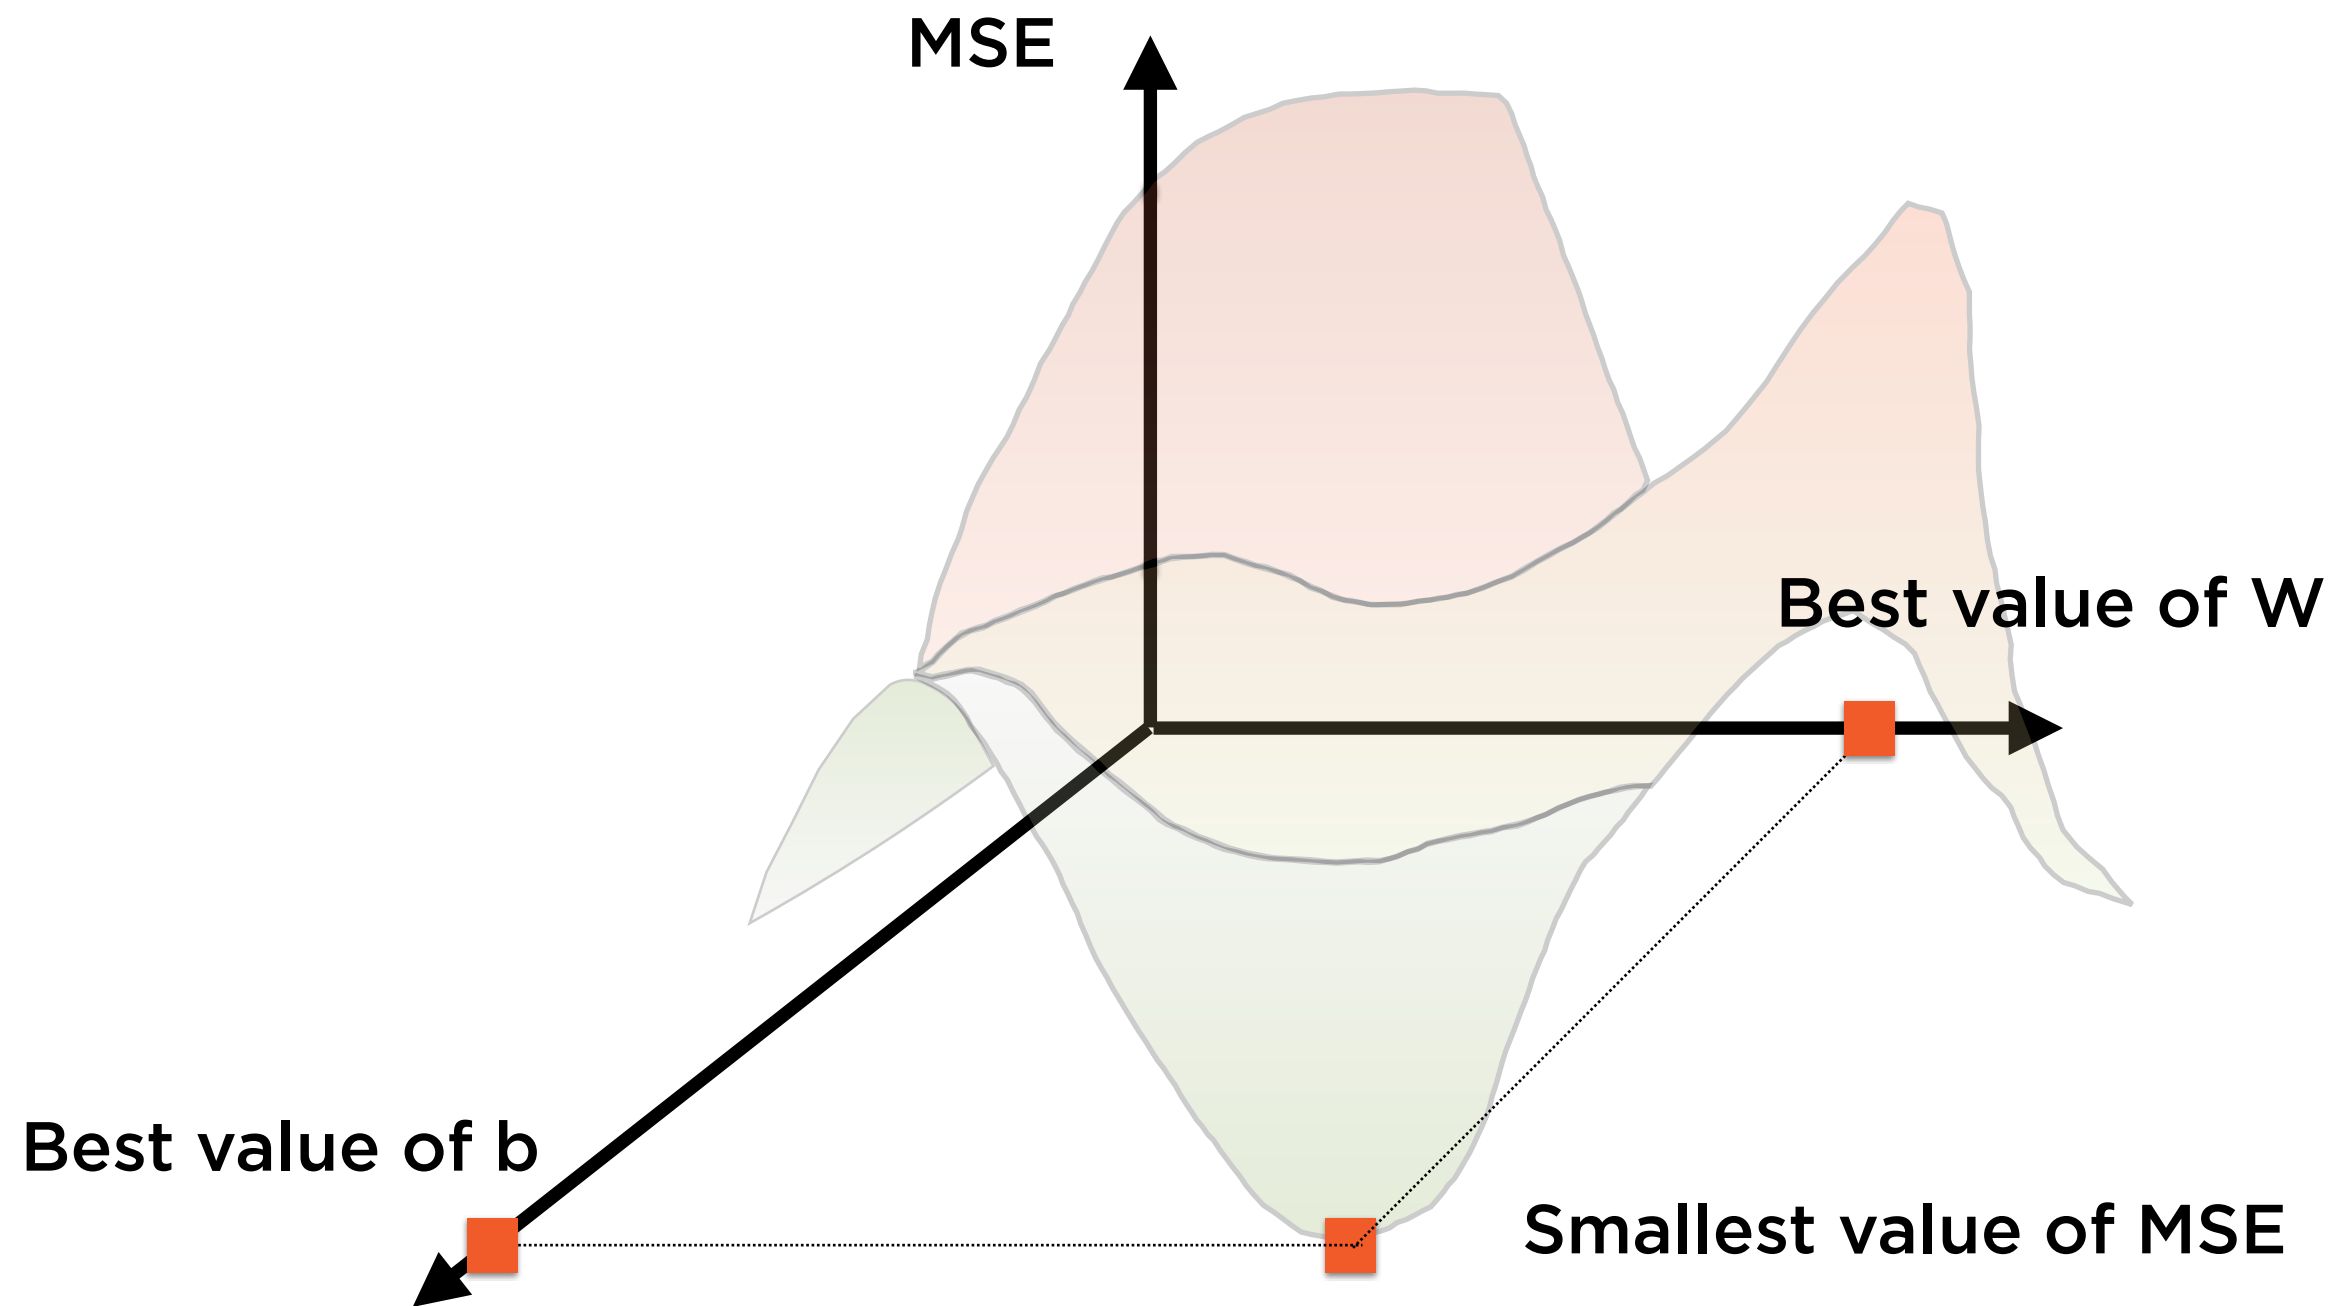


# Minimizing MSE



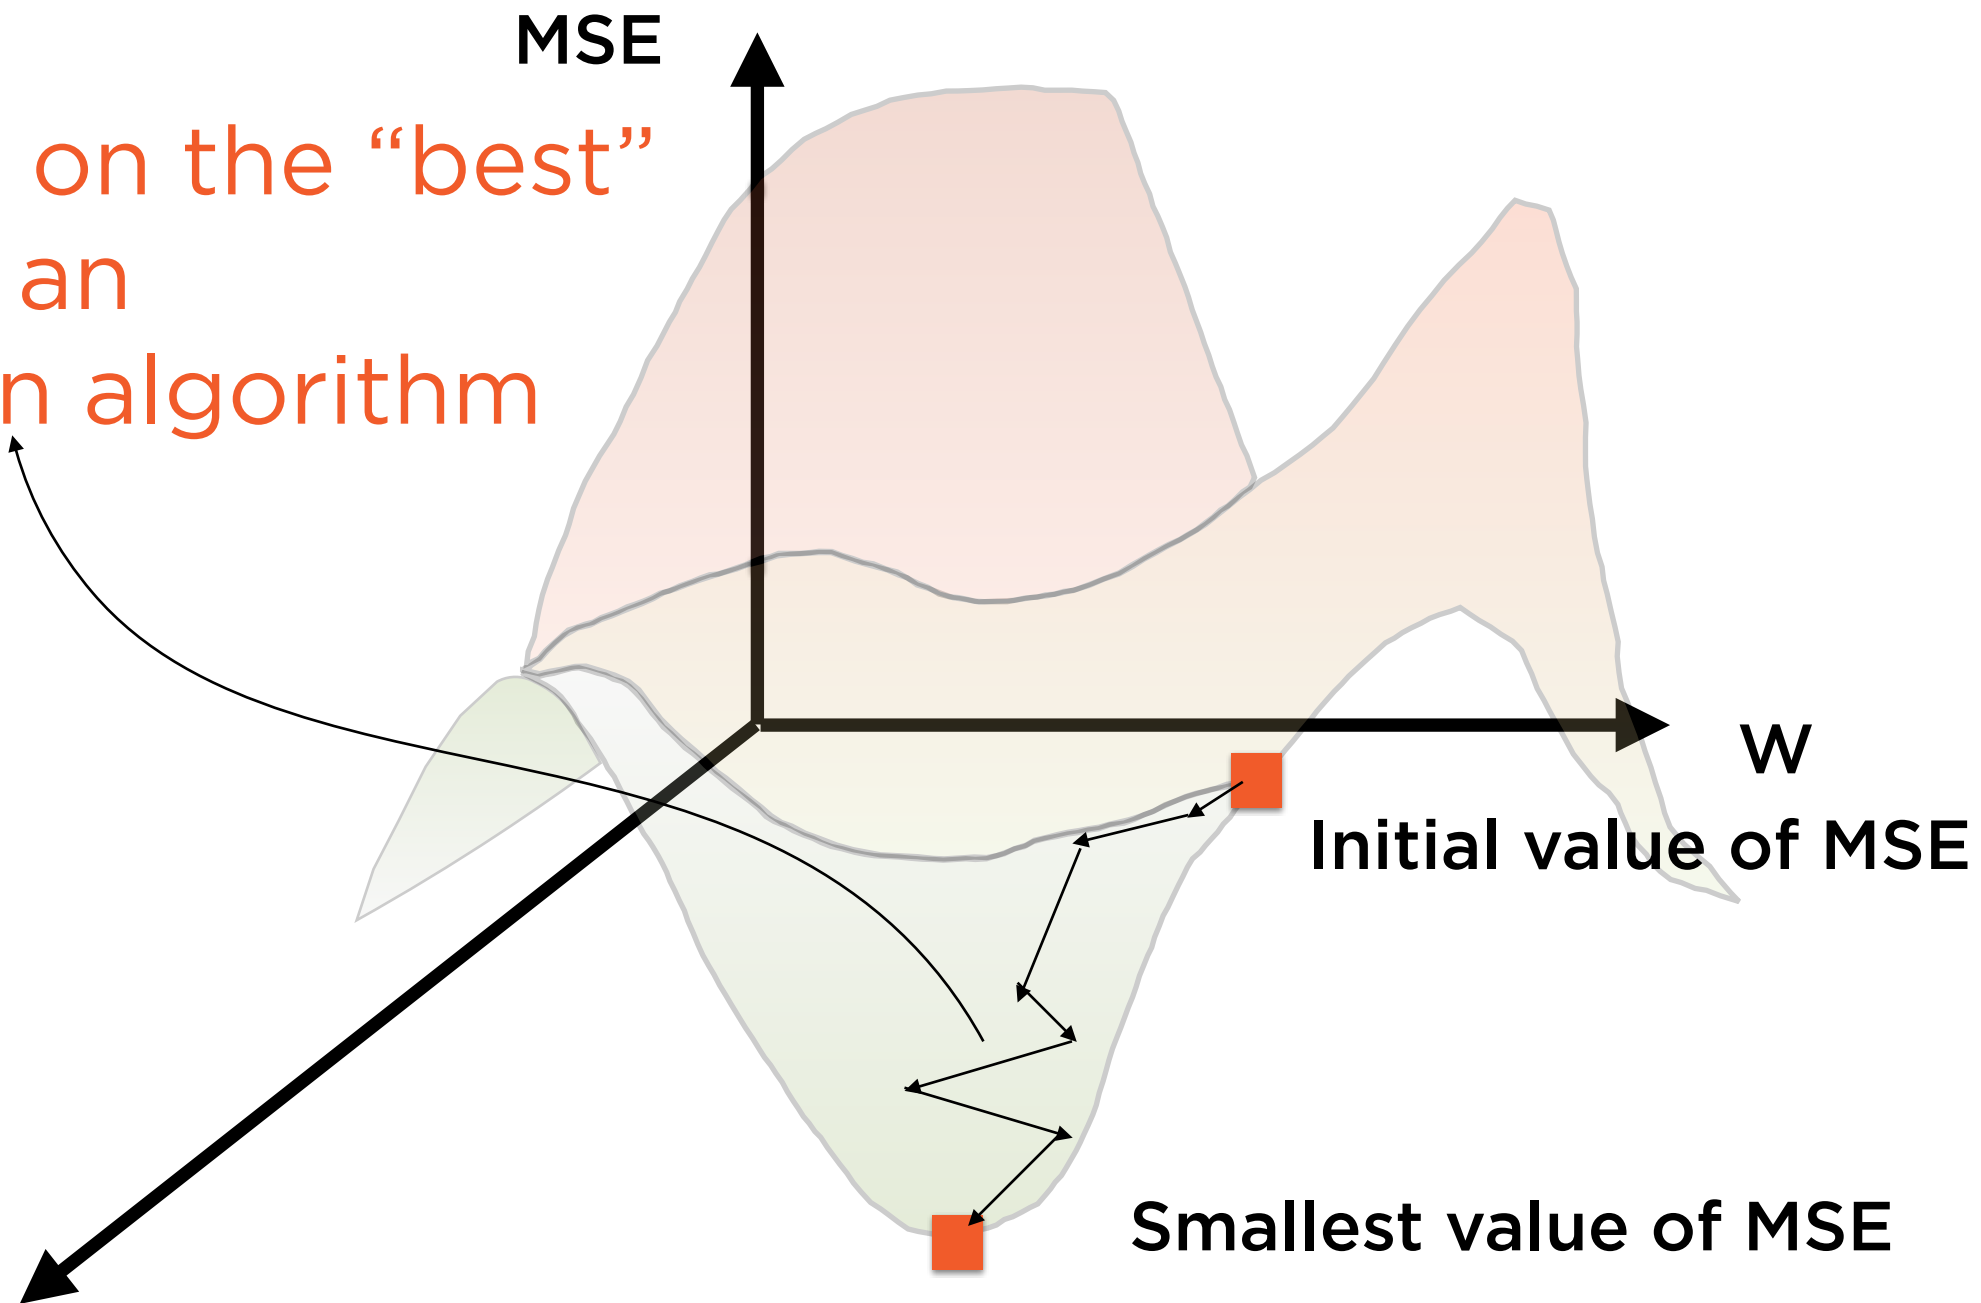


# Minimizing MSE




# “Gradient Descent”

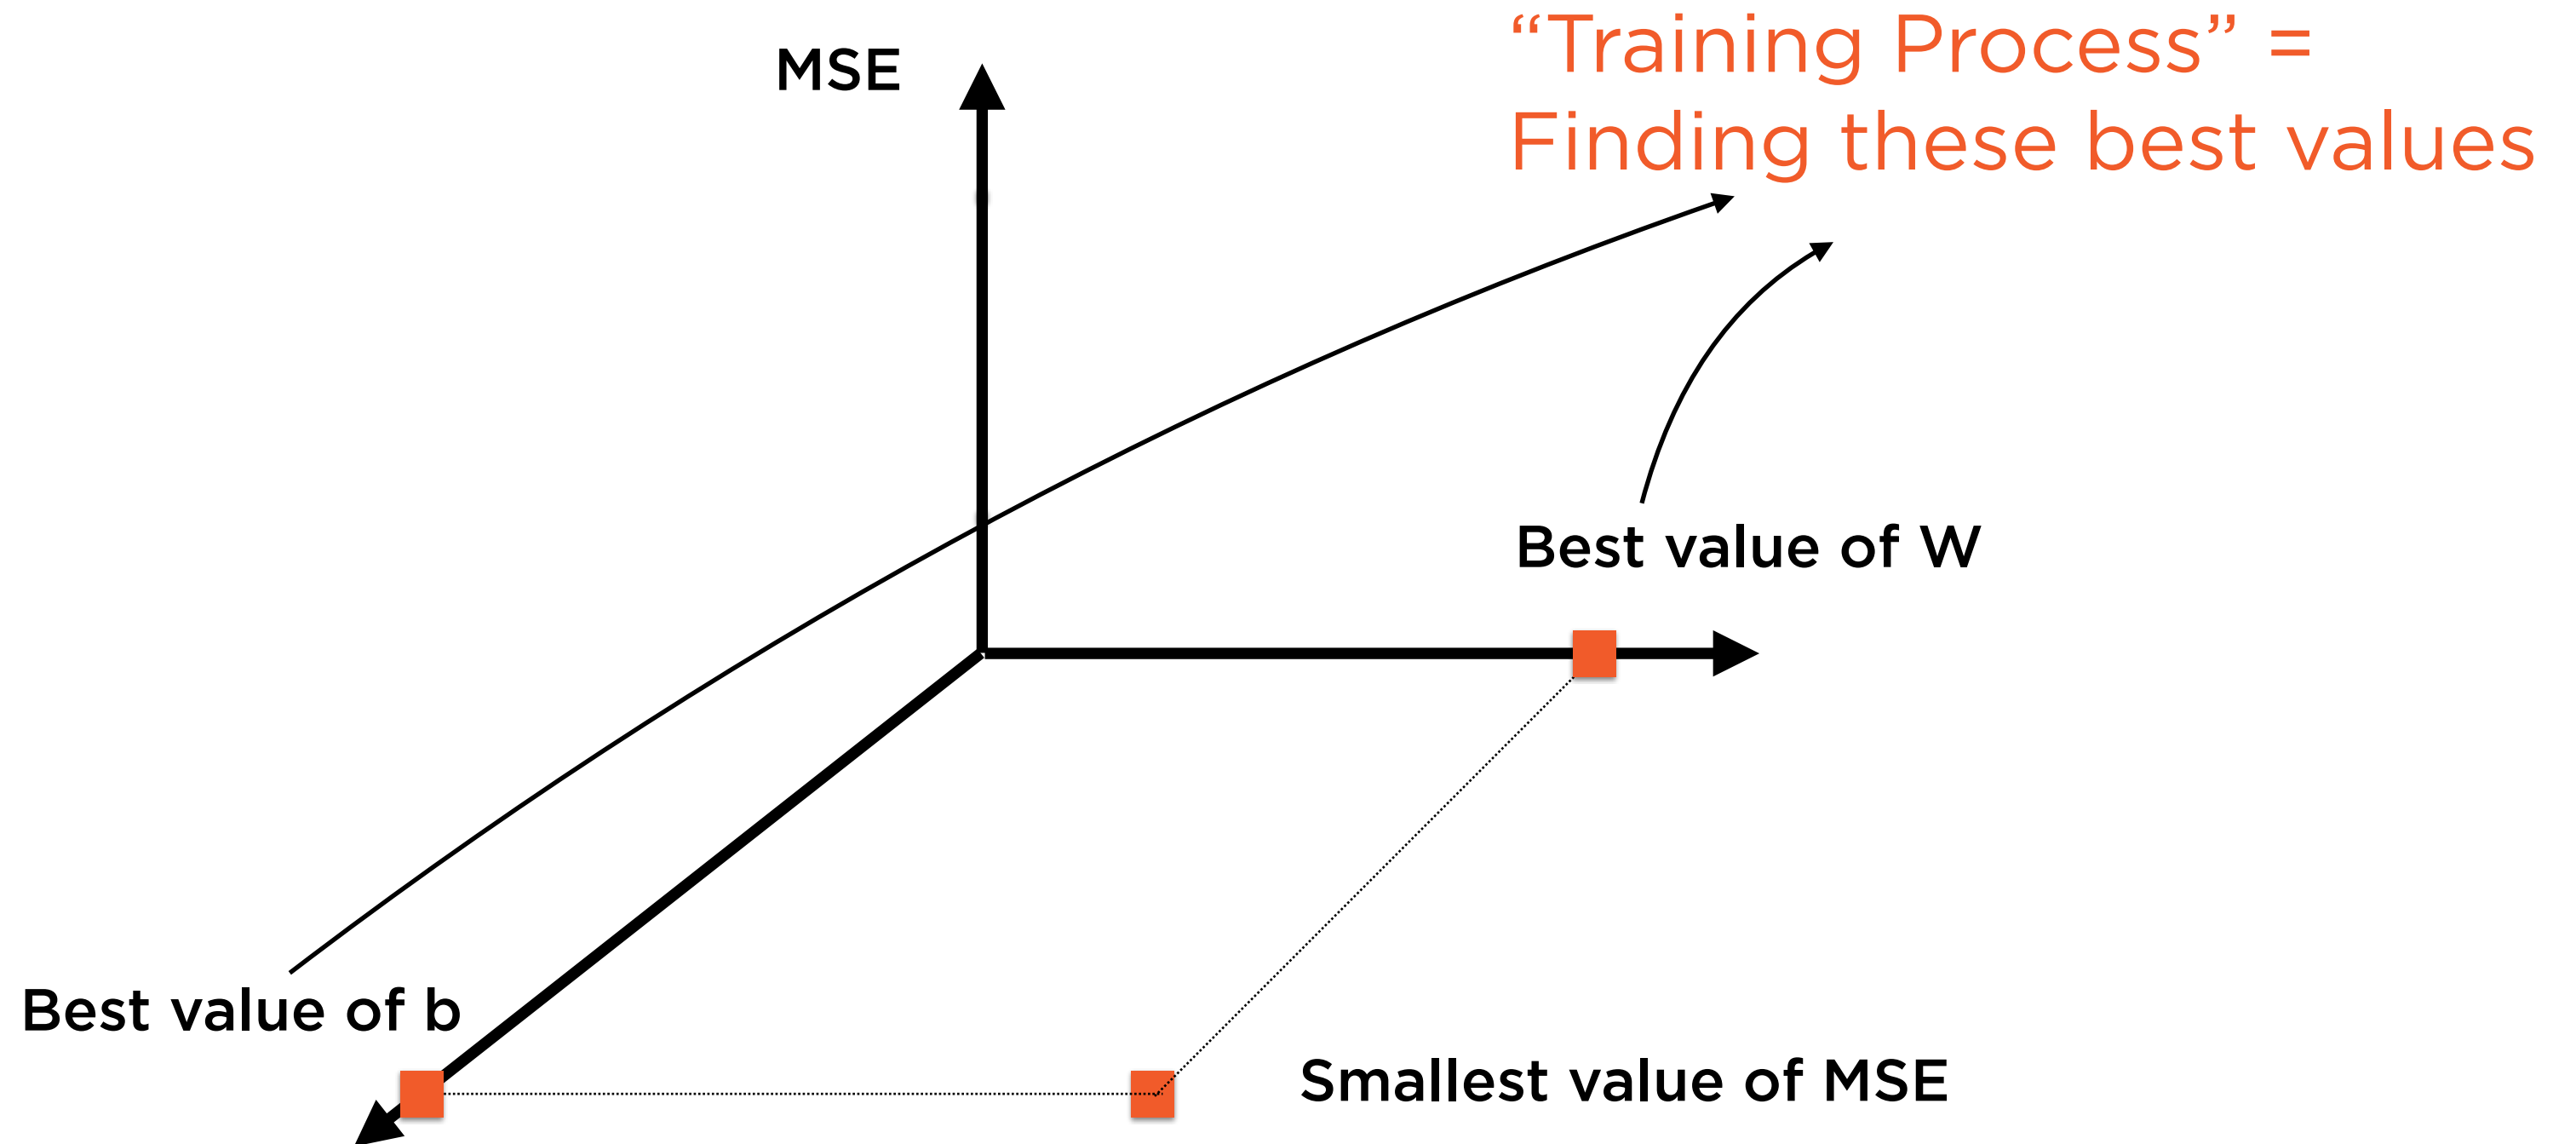
Converging on the “best”  
value using an  
optimization algorithm




# Minimizing MSE



# “Training” the Algorithm

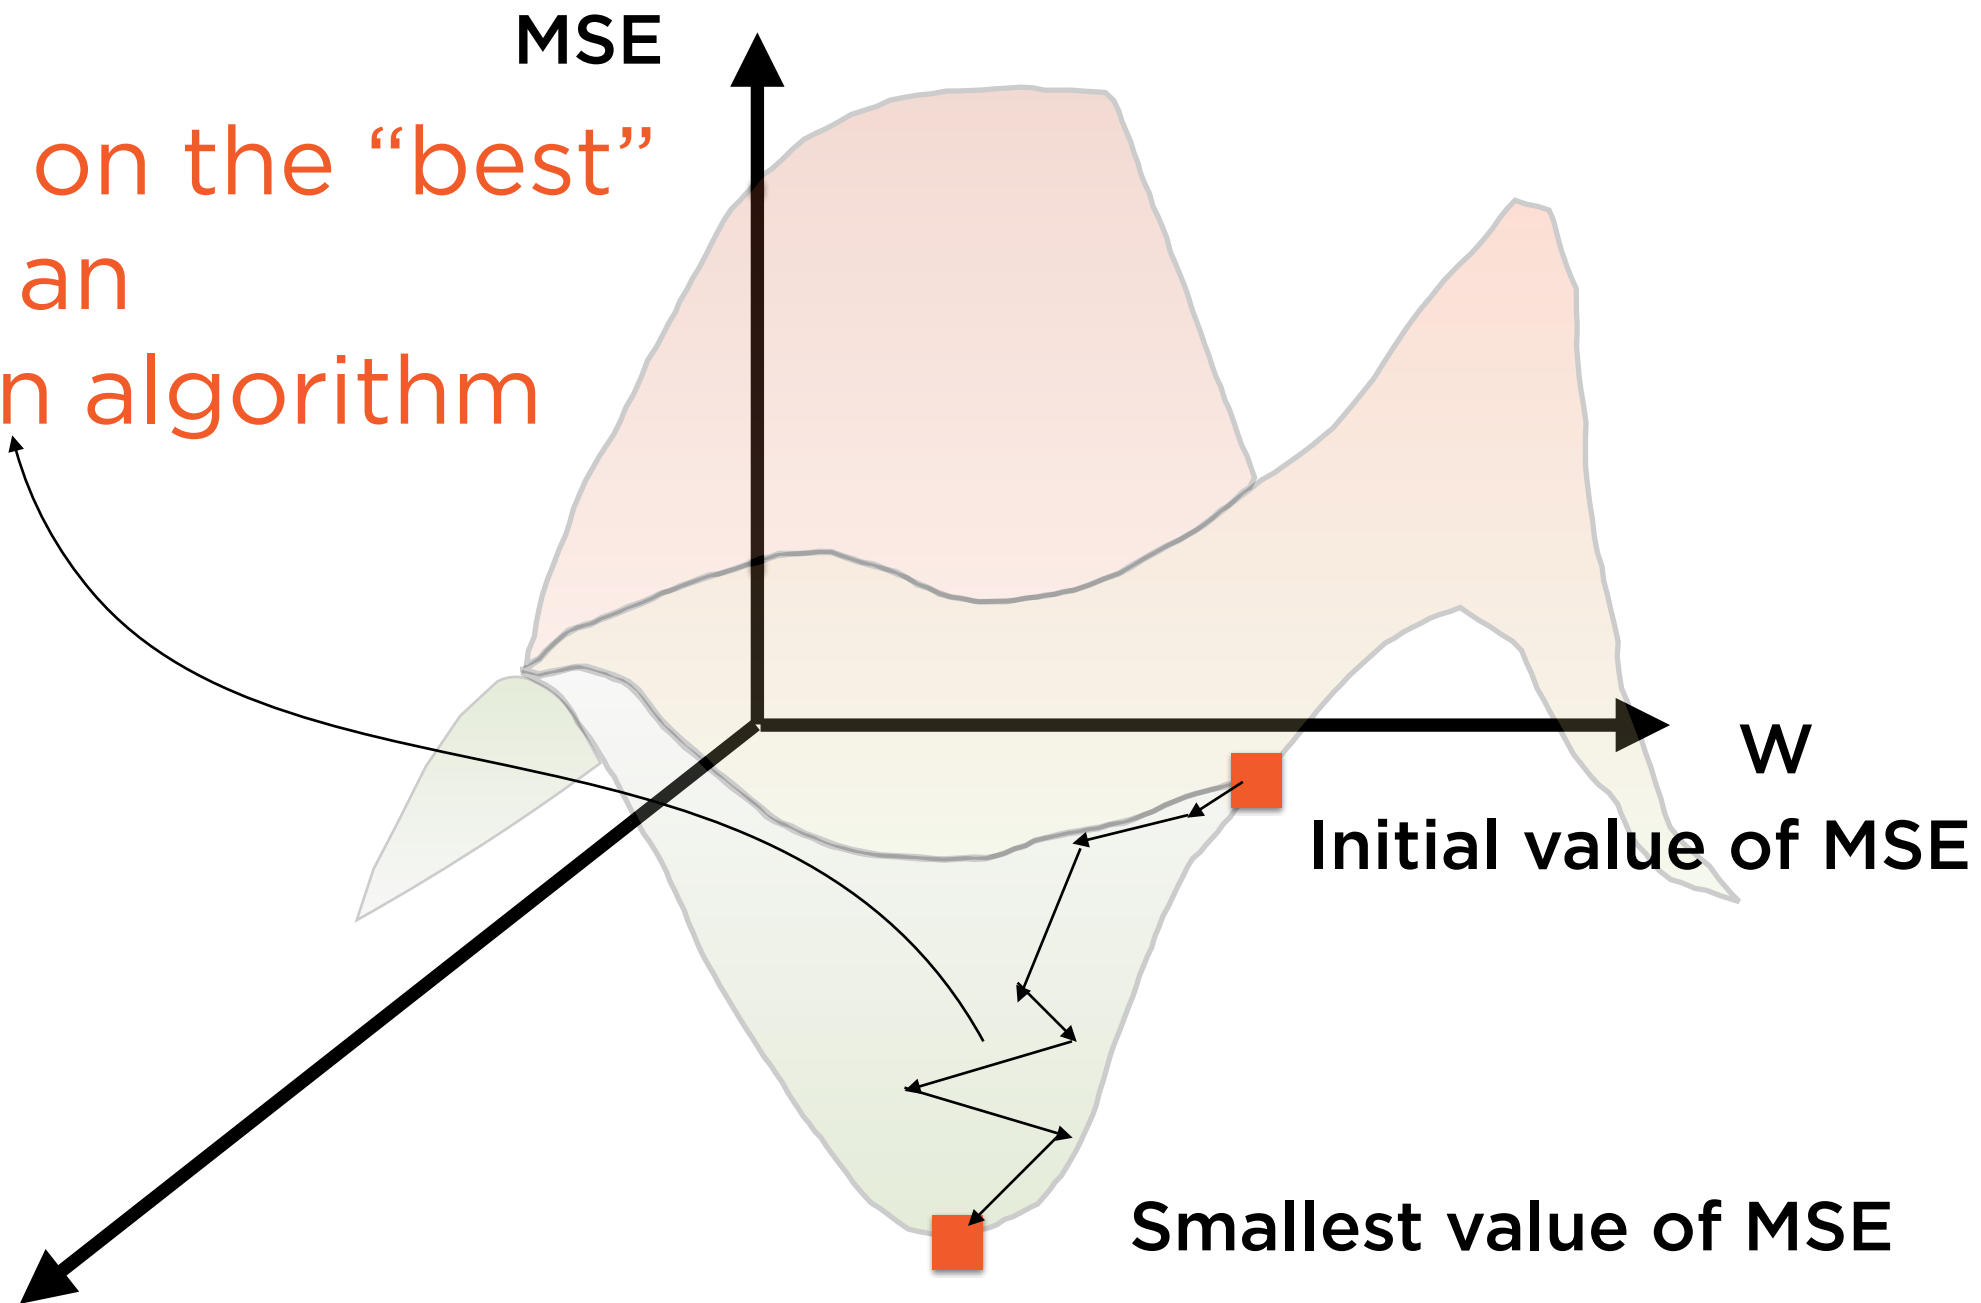


# Start Somewhere



# “Gradient Descent”

Converging on the “best”  
value using an  
optimization algorithm



Demo

**Simple Regression Using Weights  
Biases and Autograd**

# Overfitted Models

---

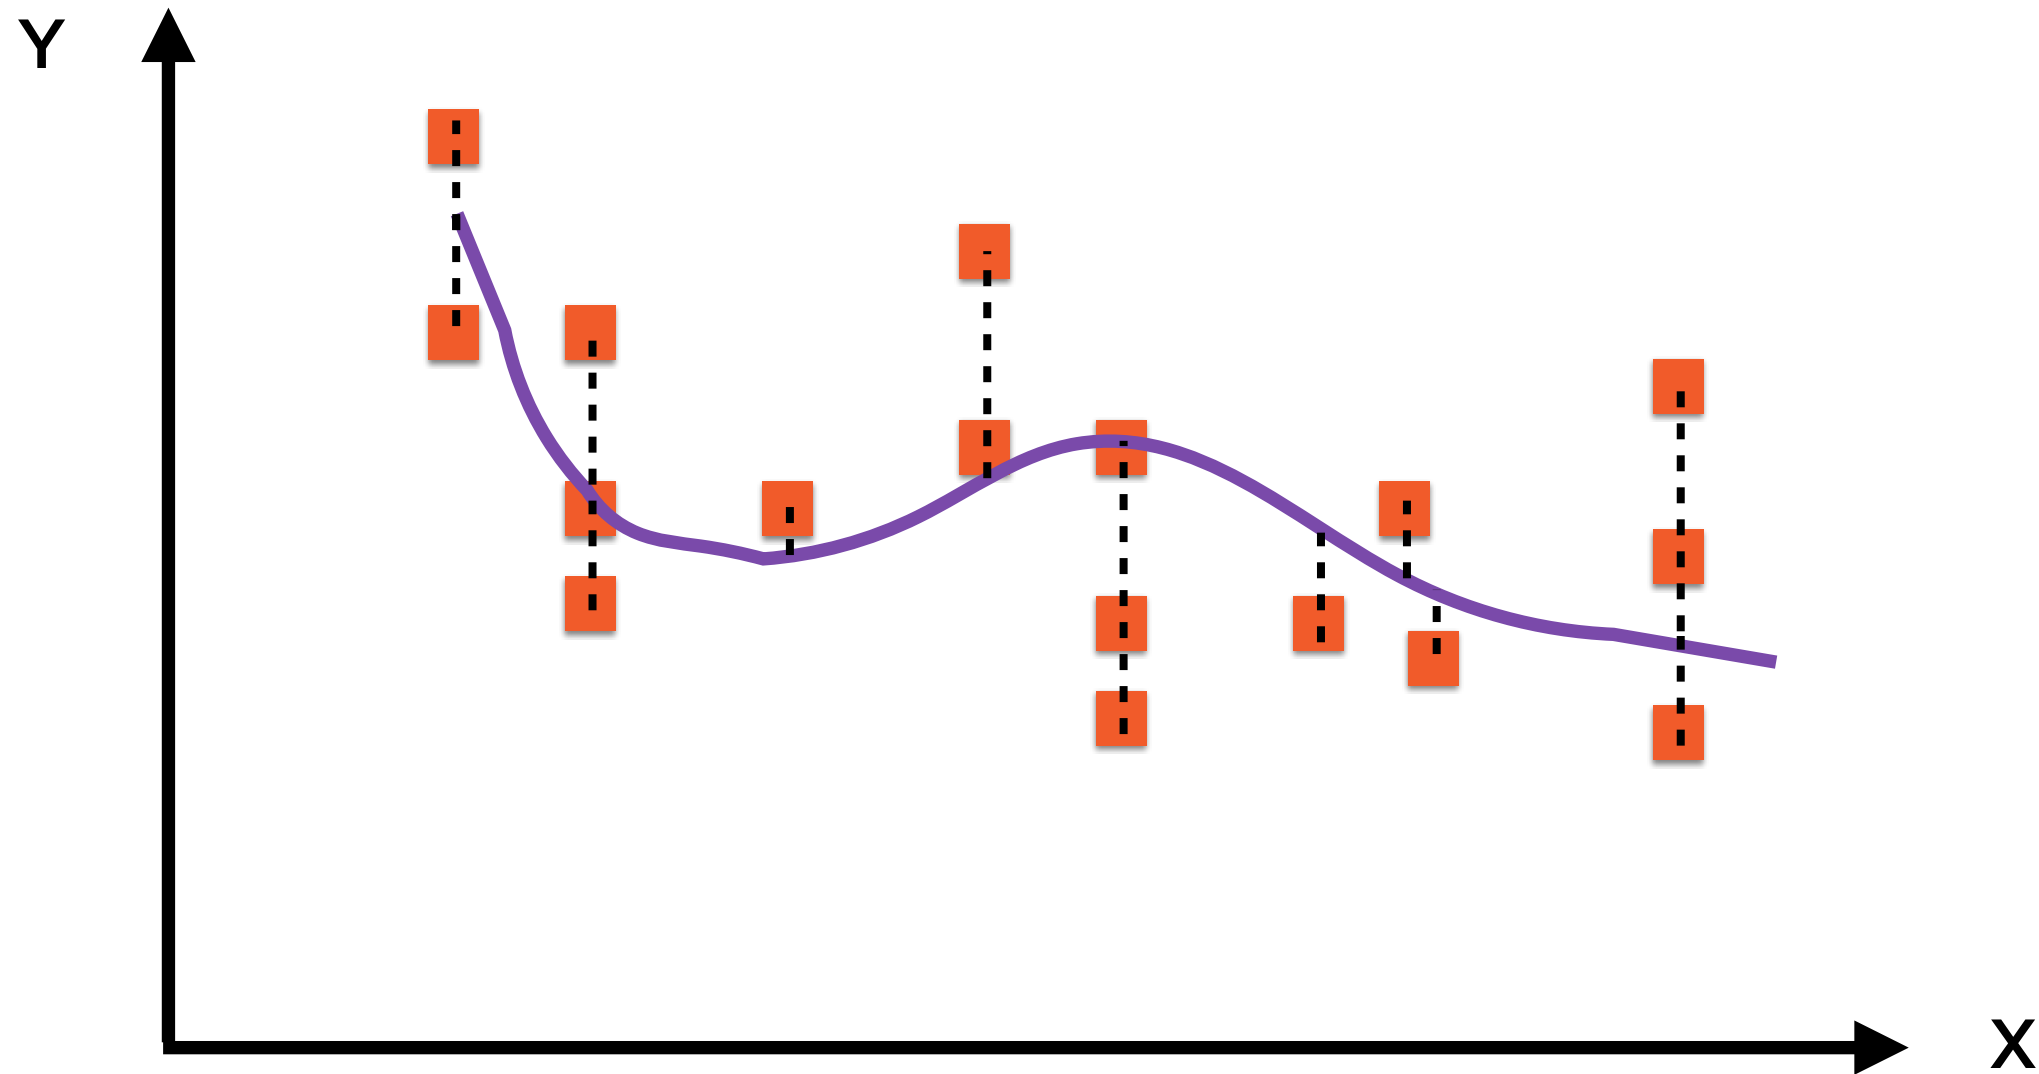


# Connecting the Dots



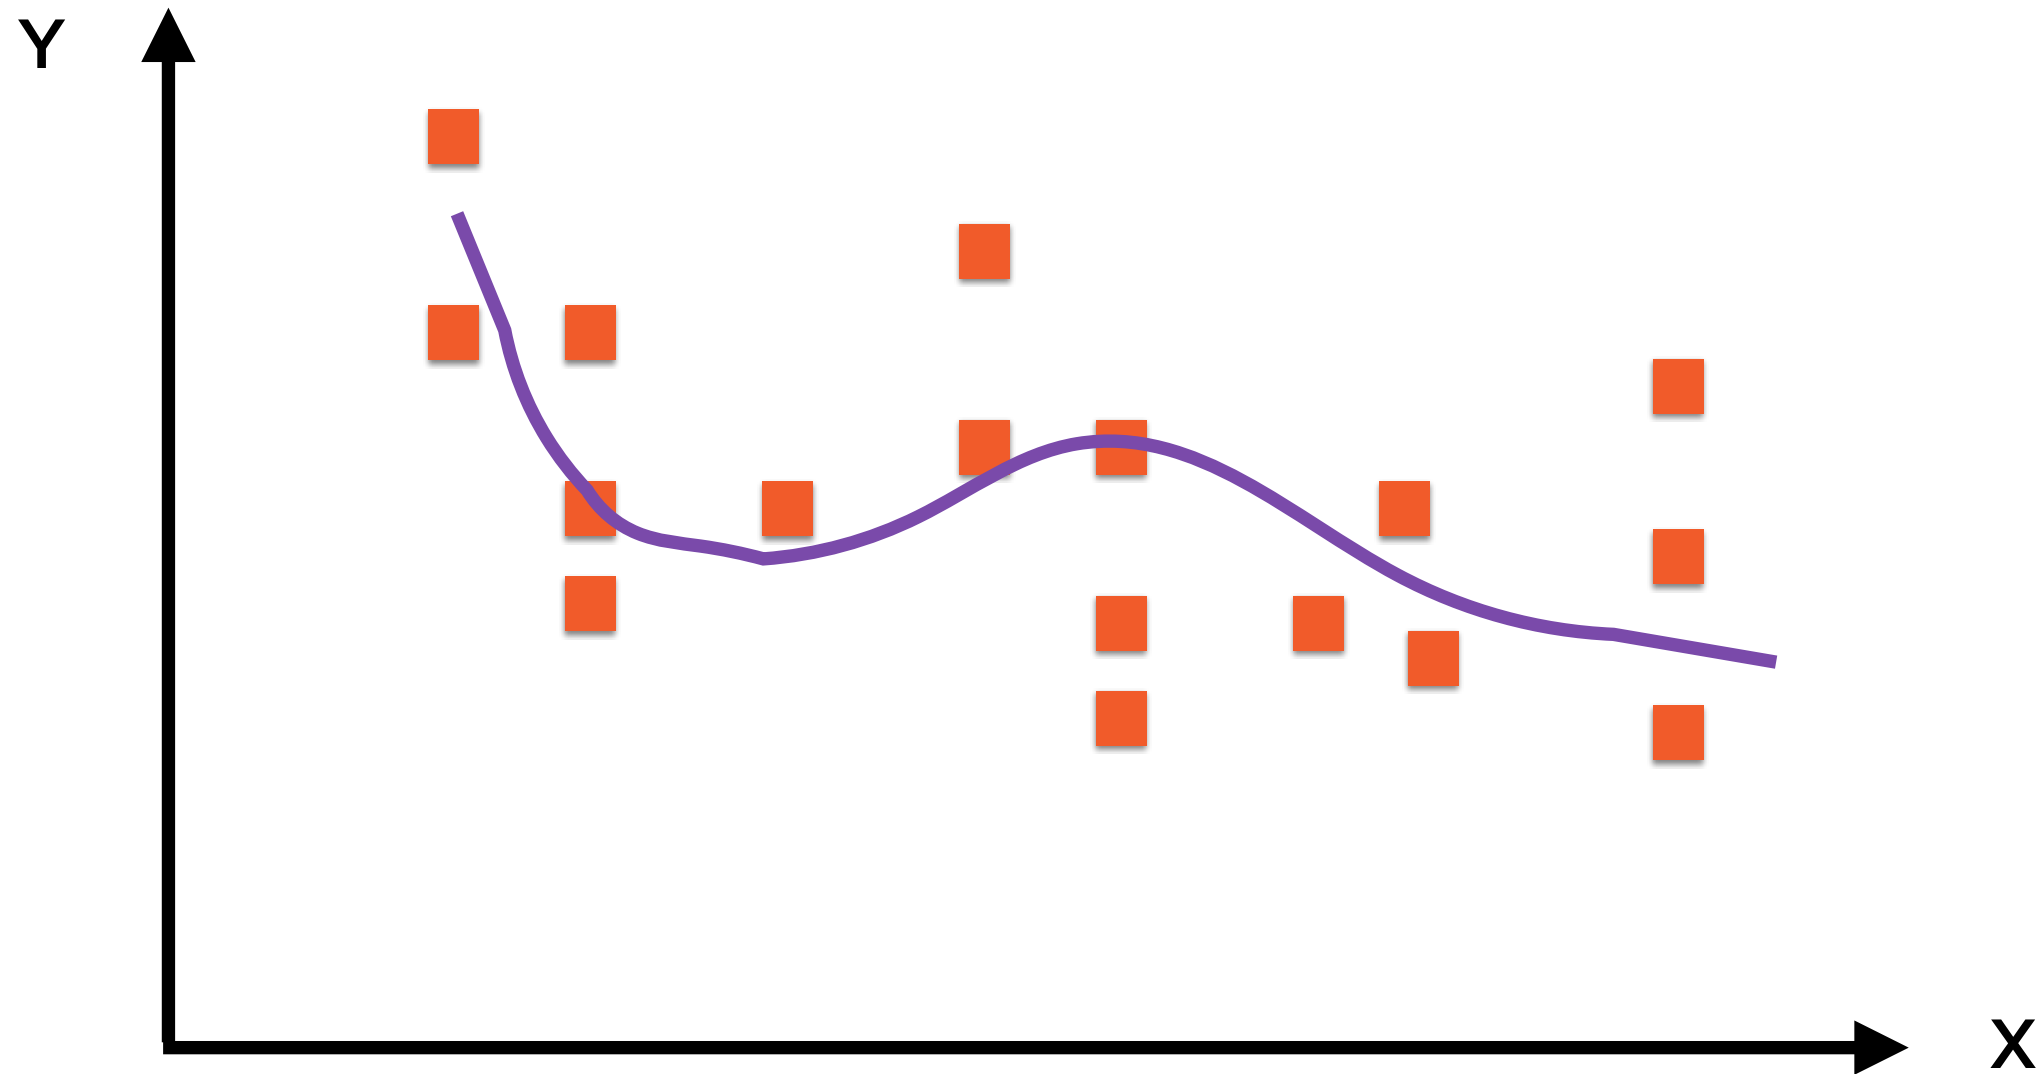
Challenge: Fit the “best” curve through these points

Good Fit?



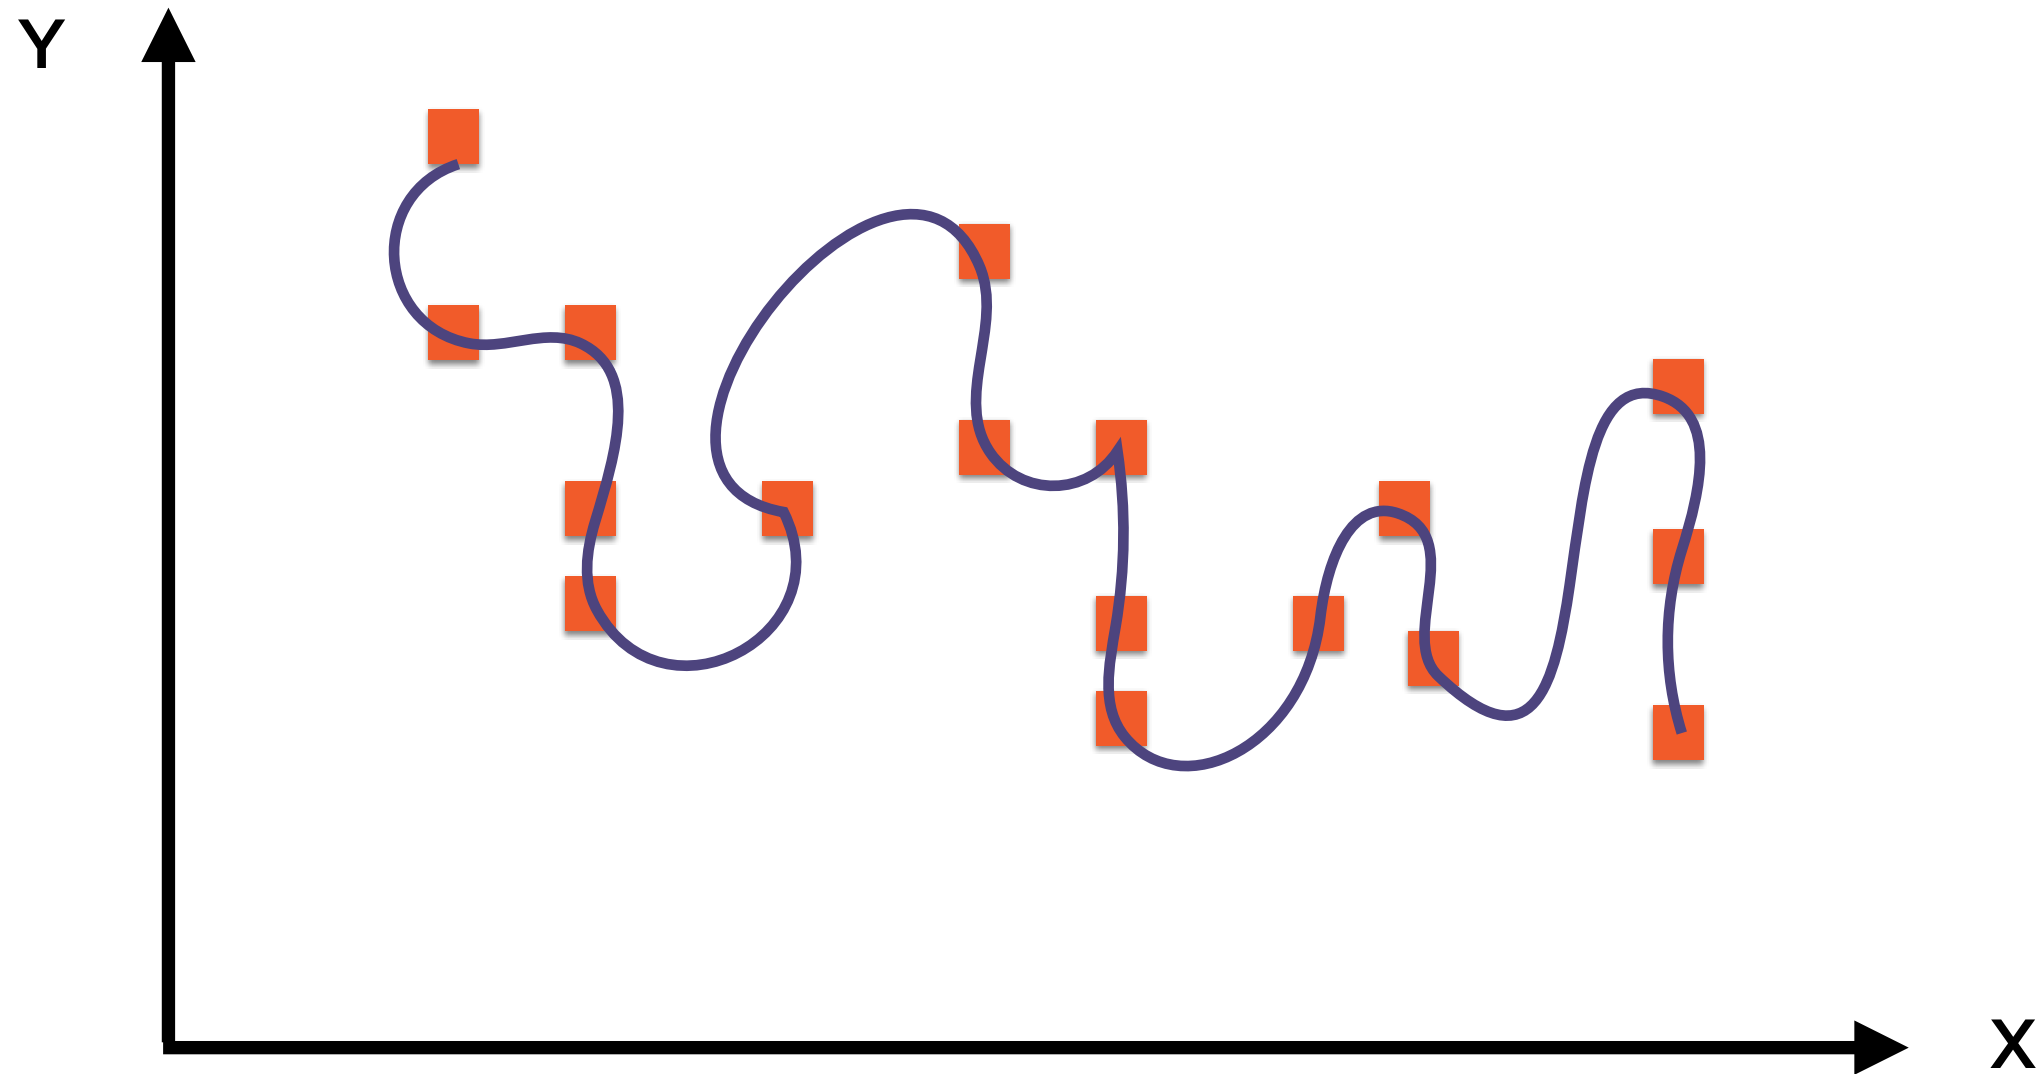
A curve has a “good fit” if the distances of points from the curve are small

# Connecting the Dots



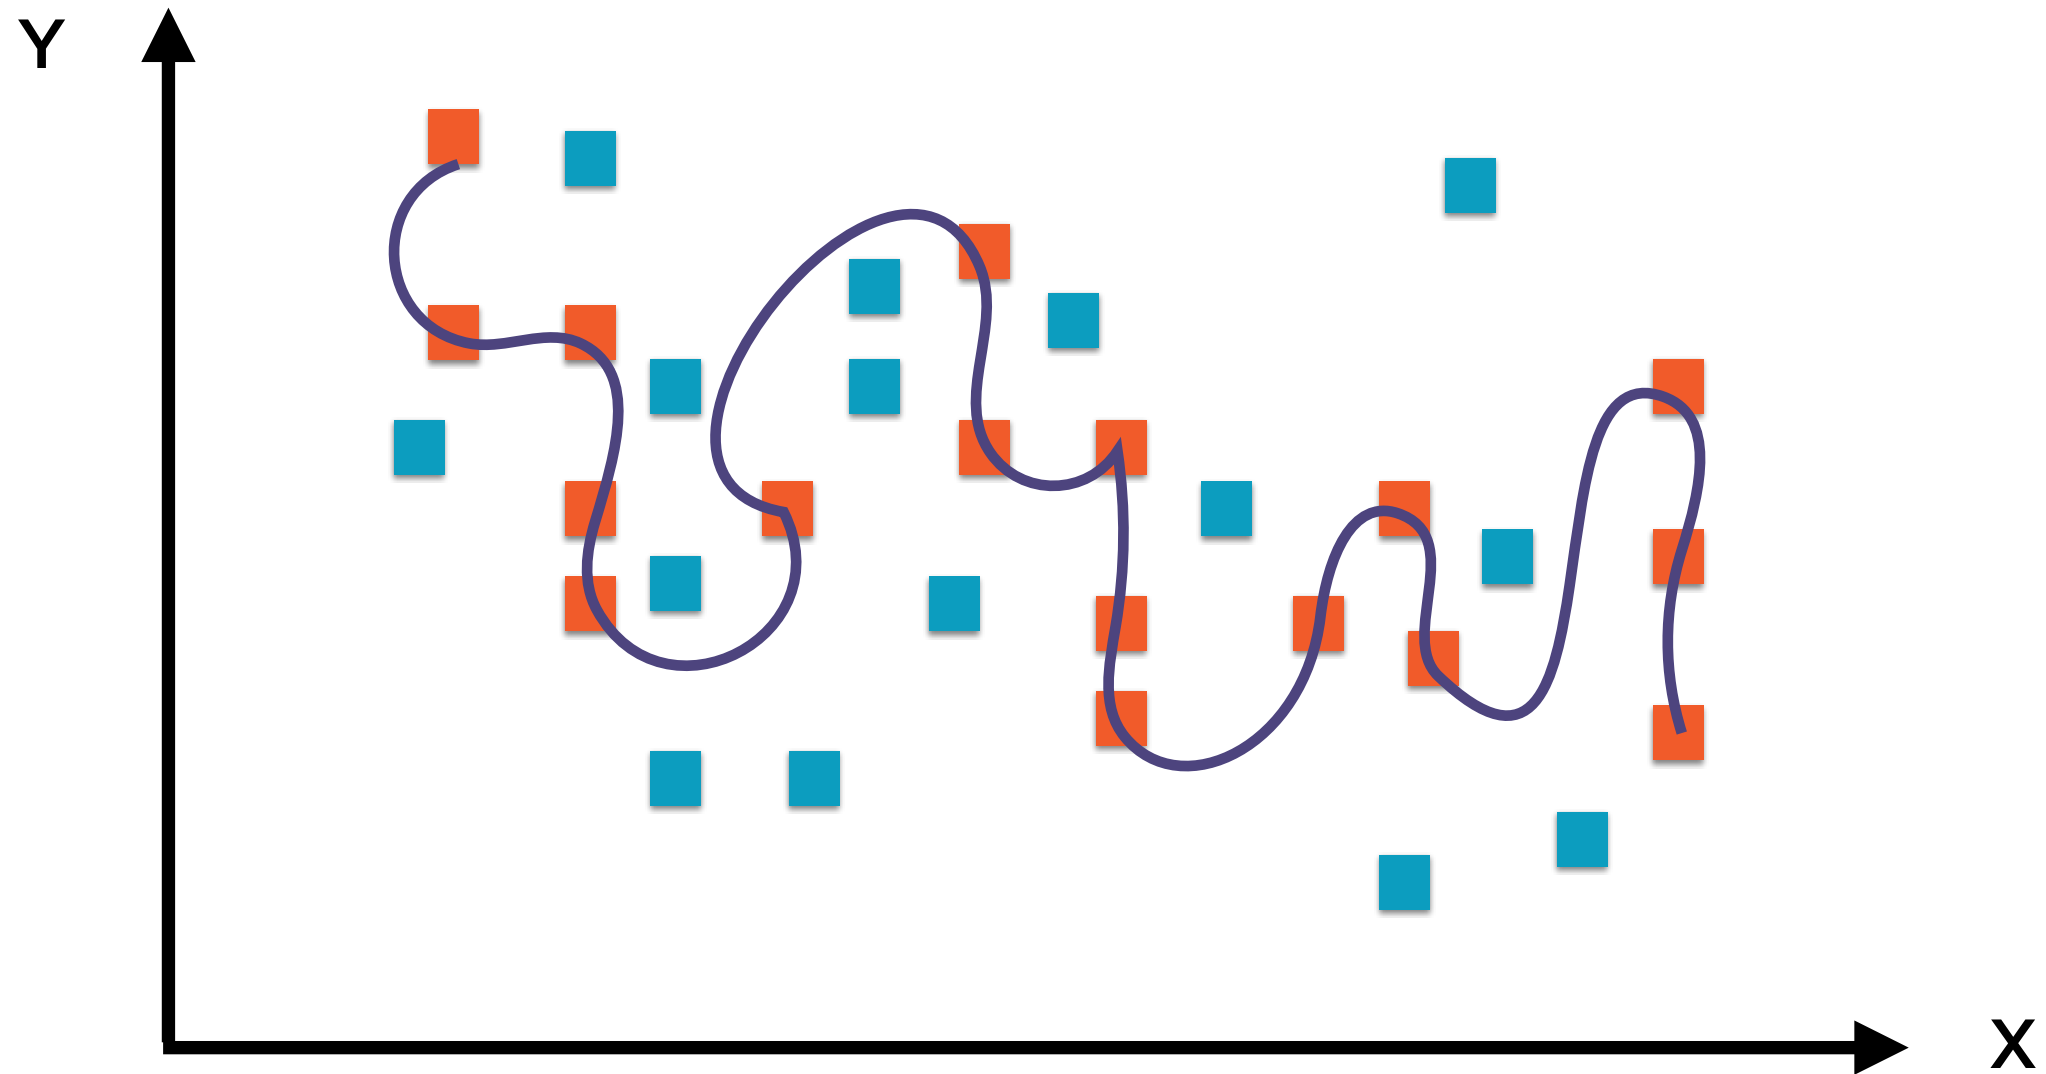
We could draw a pretty complex curve

# Connecting the Dots



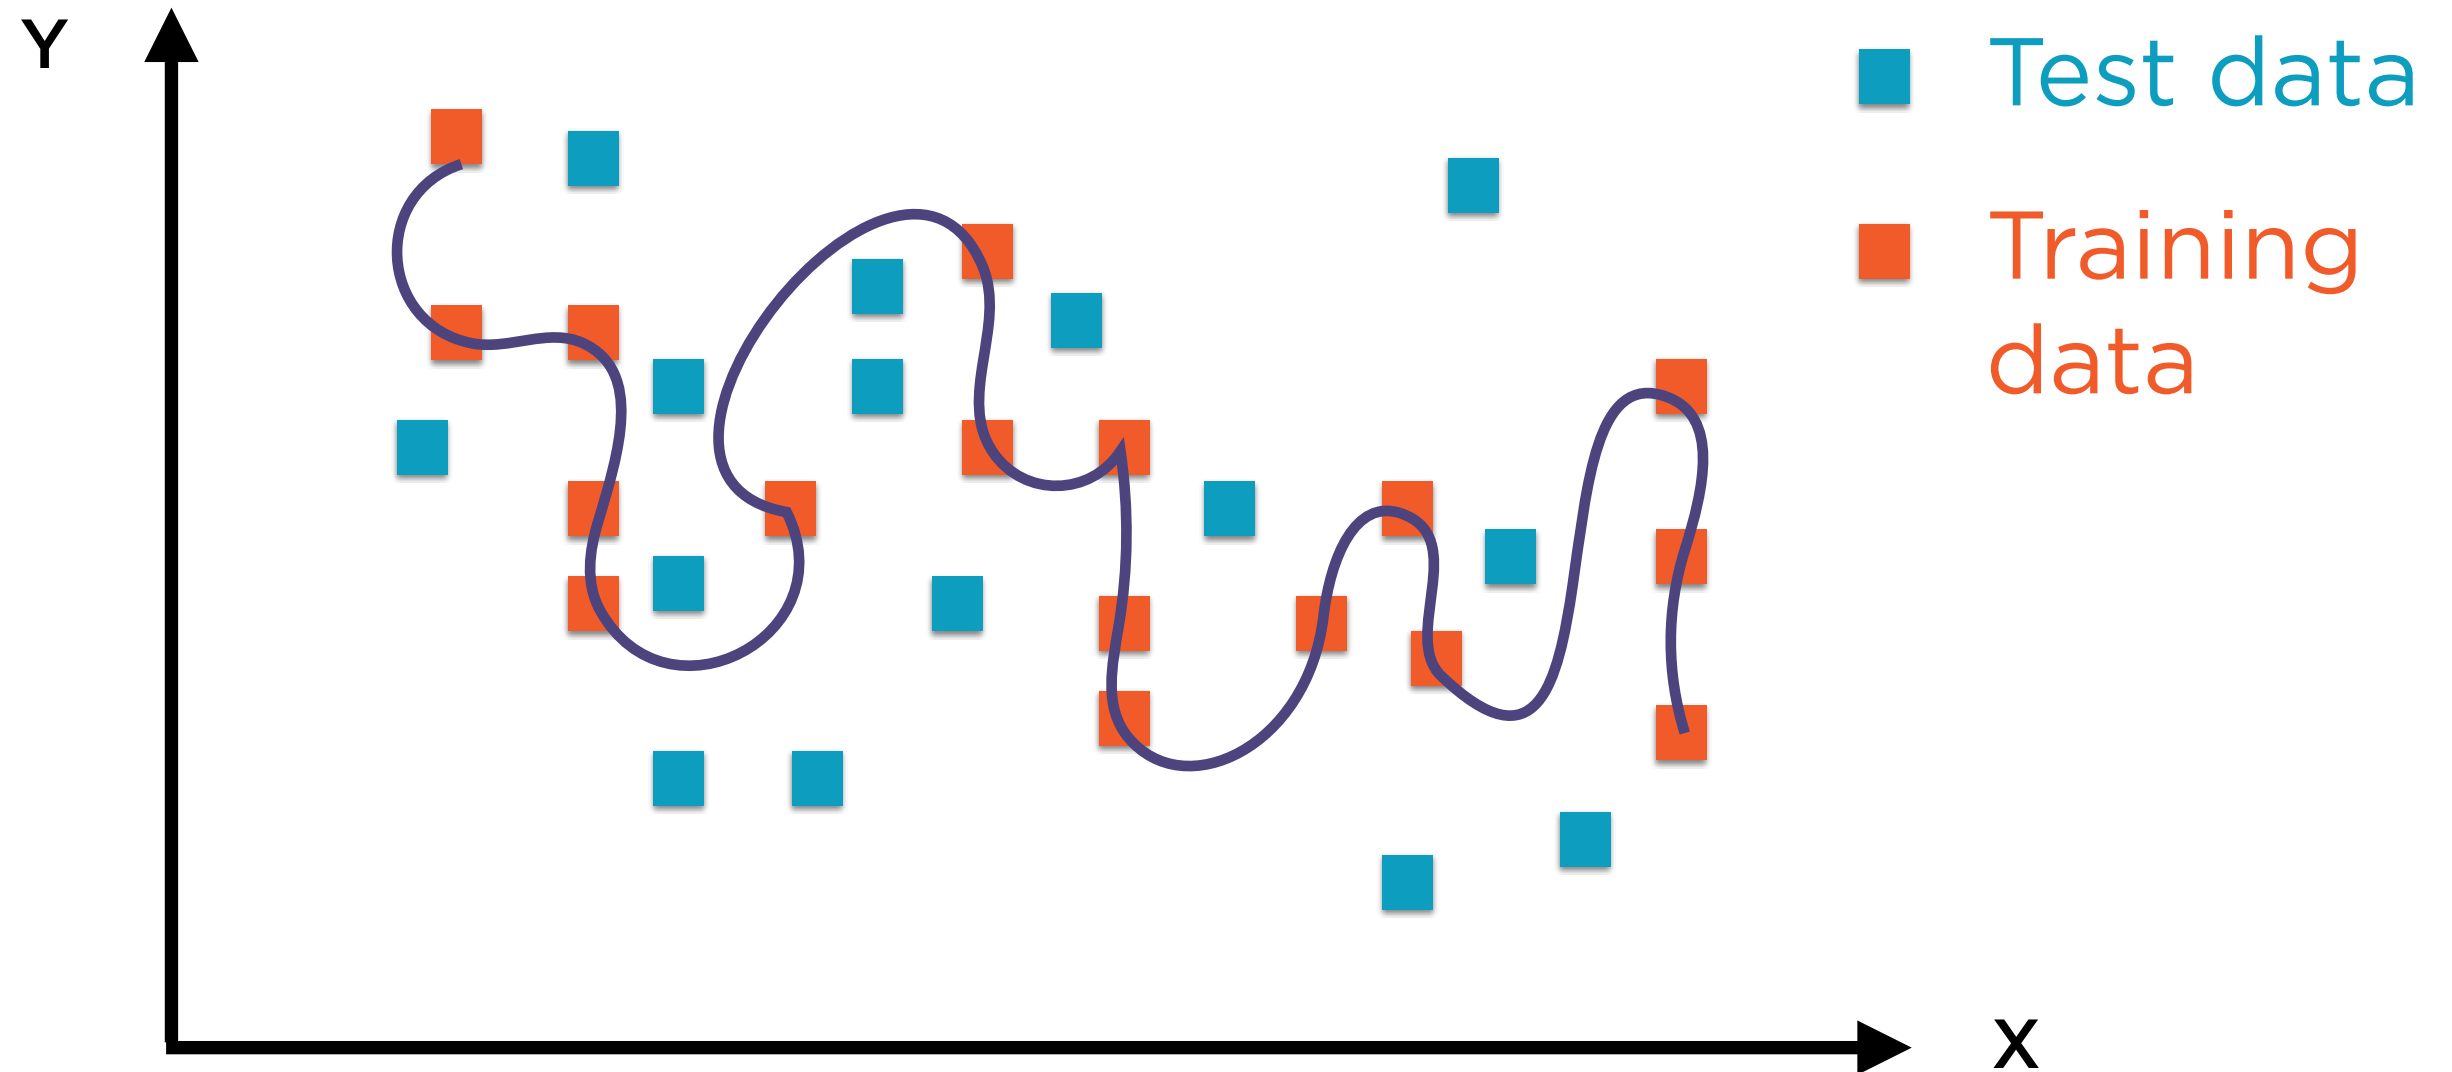
We can even make it pass through every single point

# Connecting the Dots



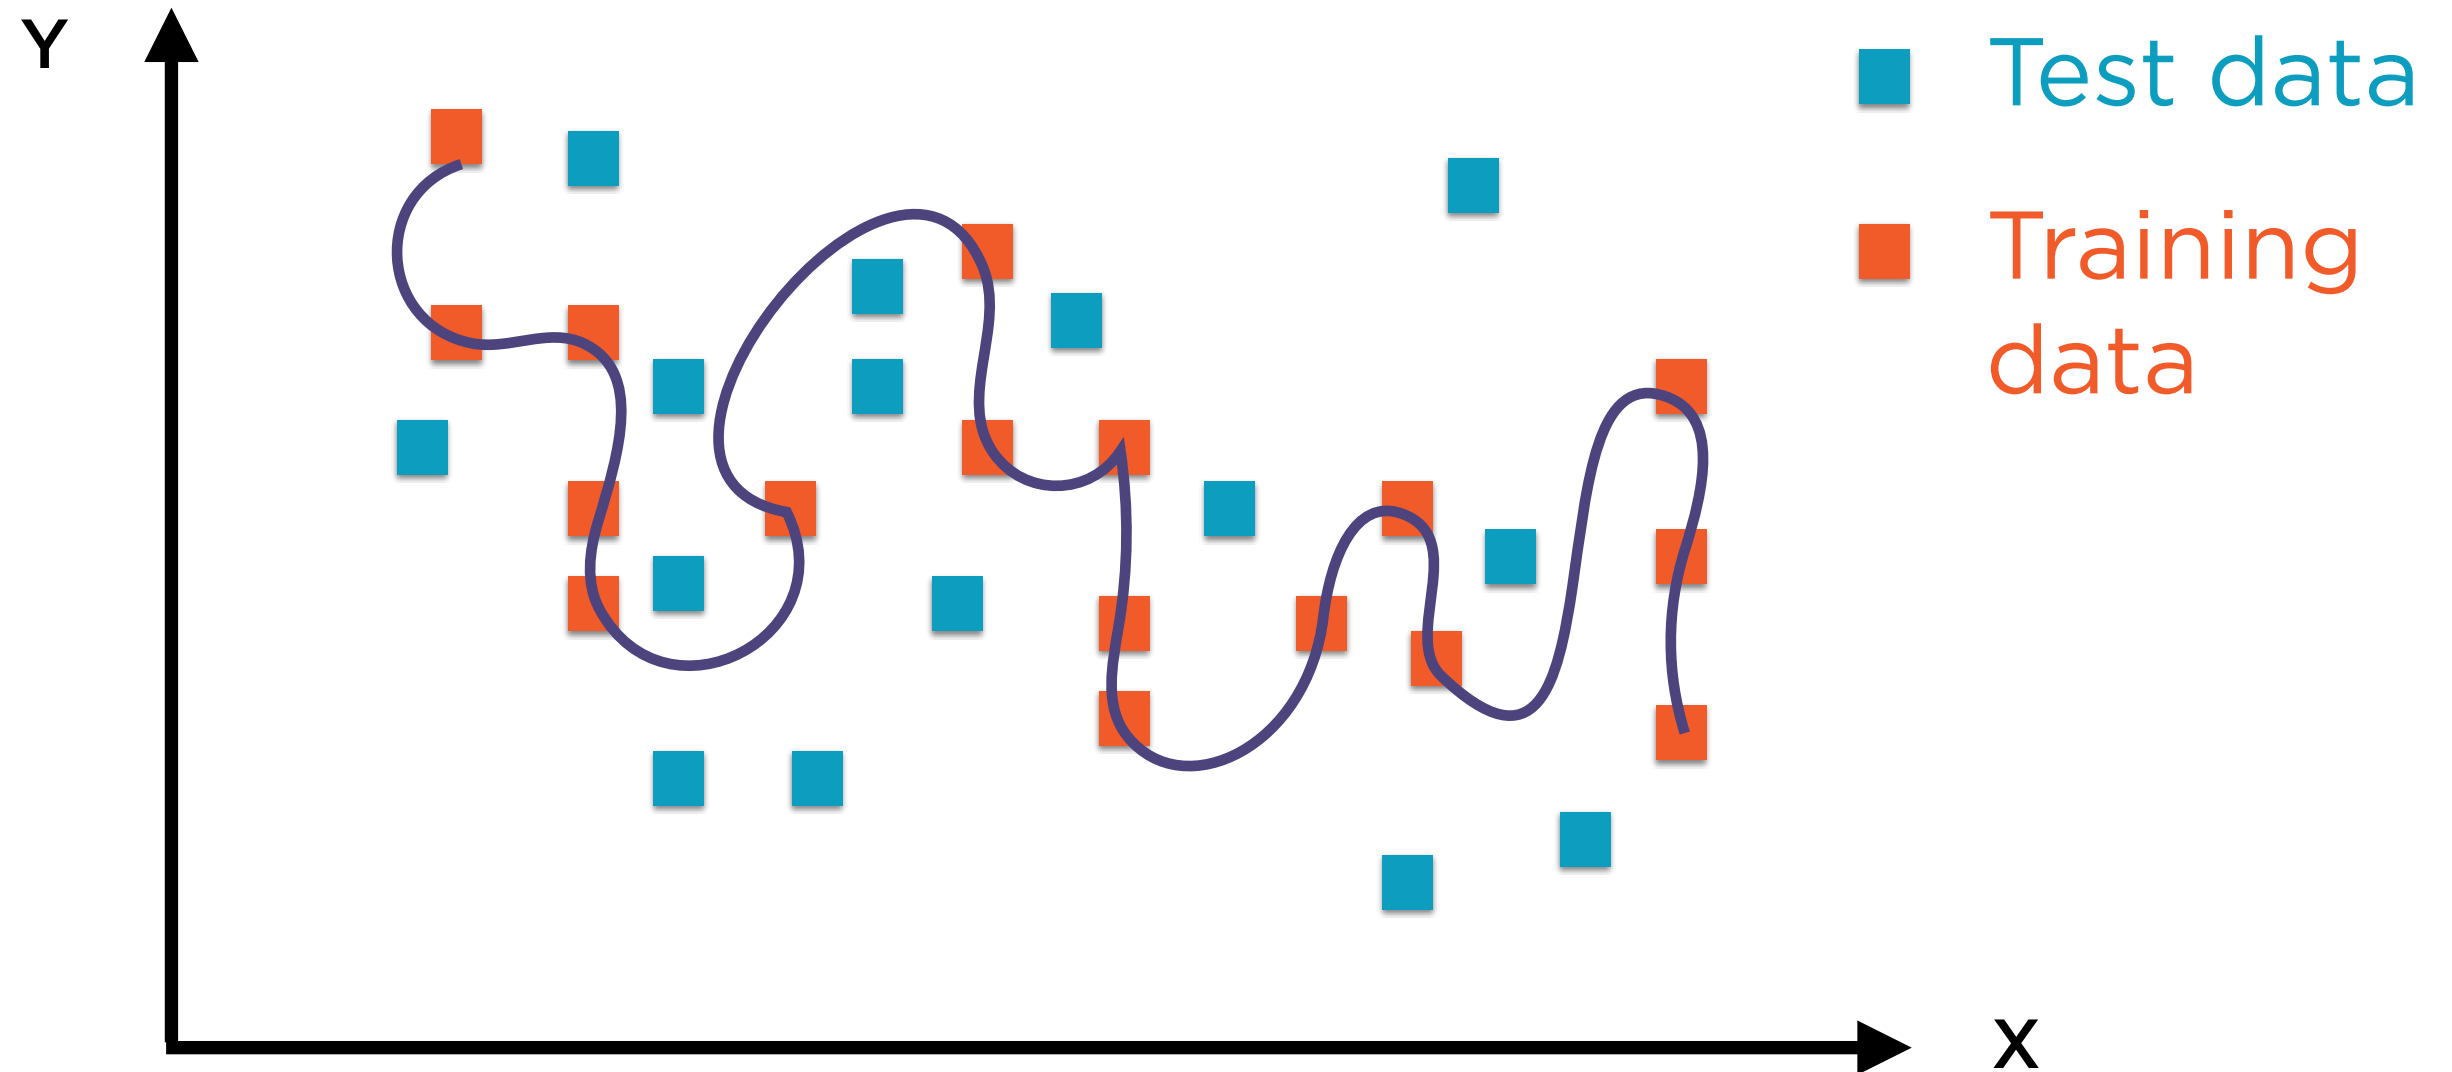
But given a new set of points, this curve might perform quite poorly

# Connecting the Dots



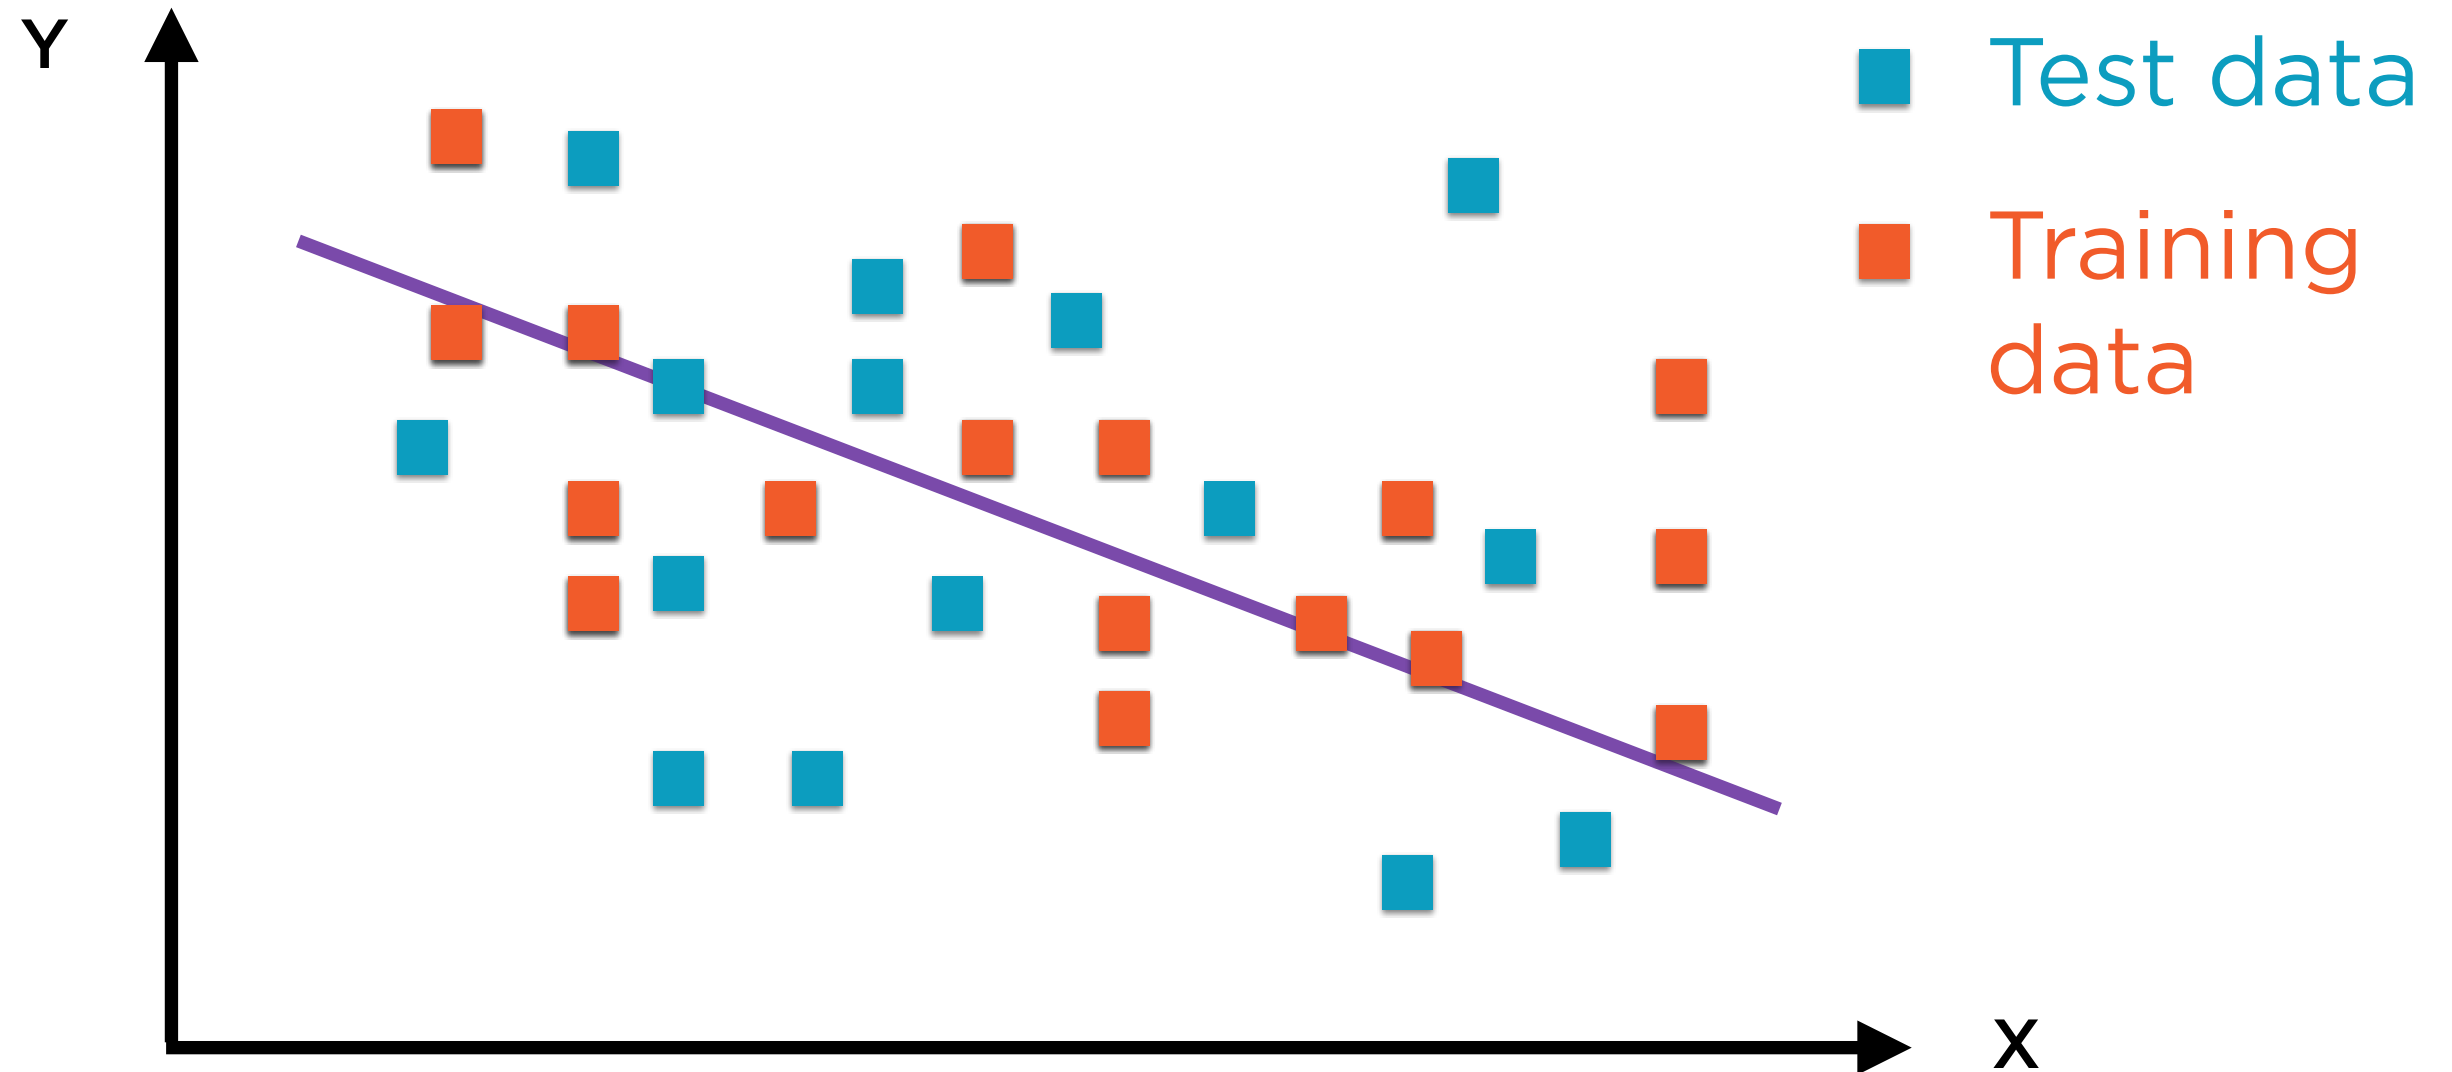
The original points were “training data”, the new points are “test data”

# Overfitting



Great performance in training, poor performance in real usage

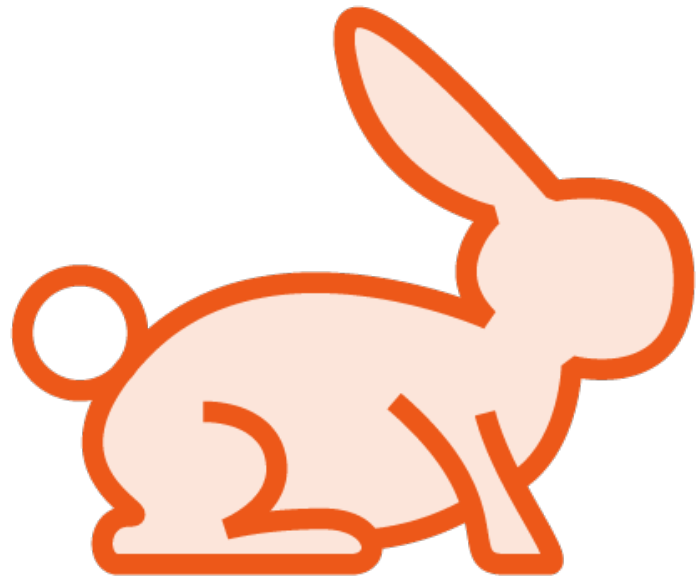
# Connecting the Dots



A simple straight line performs worse in training, but better with test data

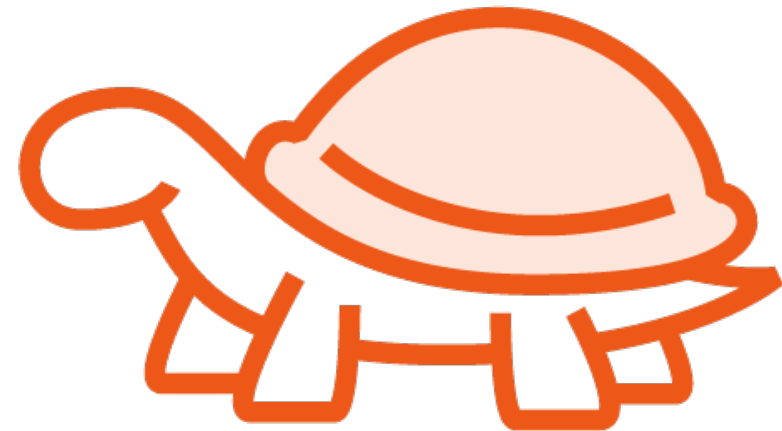


# Overfitting



**Low Training Error**

**Model does very well in training...**




**High Test Error**

**...but poorly with real data**

# Preventing Overfitting



**Regularization - Penalize complex models**



**Cross-validation - Distinct training and validation phases**



**Dropout (NNs only) - Intentionally turn off some neurons during training**

# Regularization



**Penalize complex models**

**Add penalty to objective function**

**Penalty as function of regression  
coefficients**

**Forces optimizer to keep it simple**

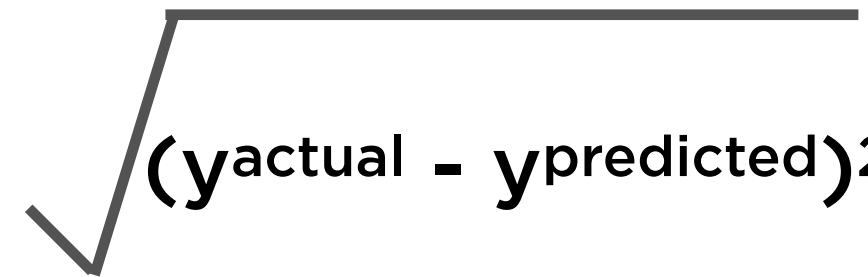
# Regularization



**Regularization reduces variance error**  
**But increases bias**

# Ordinary MSE Regression

**Minimize**


$$(y^{\text{actual}} - y^{\text{predicted}})^2$$

**To find**

**A, B**

**The value of A and B define the “best fit” line**

$$y = A + Bx$$

# Ridge Regression

ridge regression(岭回归)

Minimize

$$\sqrt{(y^{\text{actual}} - y^{\text{predicted}})^2} + \alpha (|A|^2 + |B|^2)$$

To find

A, B

L-2 Norm of regression  
coefficients

$\alpha$  is a hyperparameter

The value of A and B still define the “best fit” line

$$y = A + Bx$$

# Ridge Regression

Minimize

$$\sqrt{(y_{\text{actual}} - y_{\text{predicted}})^2}$$

To find

A, B

$$+ \alpha (|A|^2 + |B|^2)$$

L-2 Norm of regression  
coefficients

**$\alpha$  is a hyperparameter**

The value of A and B still define the “best fit” line

$$y = A + Bx$$

# Ridge Regression



**Add penalty for large coefficients**

**Penalty term is L-2 norm of coefficients**

**Penalty weighted by hyperparameter  $\alpha$**



Demo

**Implementing Ridge Regression**

# Summary

**Using linear regression for prediction**

**Linear regression using a single neuron**

**Hand-crafting an MSE regression model**

**Hand-crafting a Ridge regression model**

**Comparing to scikit-learn's linear regression estimator**