

Building a Regression Model Using PyTorch



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Training a neural network using forward and backward passes

Using optimizers to update model parameters

Using layers and activation functions to train and build neural networks


Understanding and using dropout to mitigate overfitting

Training A Neural Network

Training a neural network uses
Gradient Descent to find the
weights of the model parameters

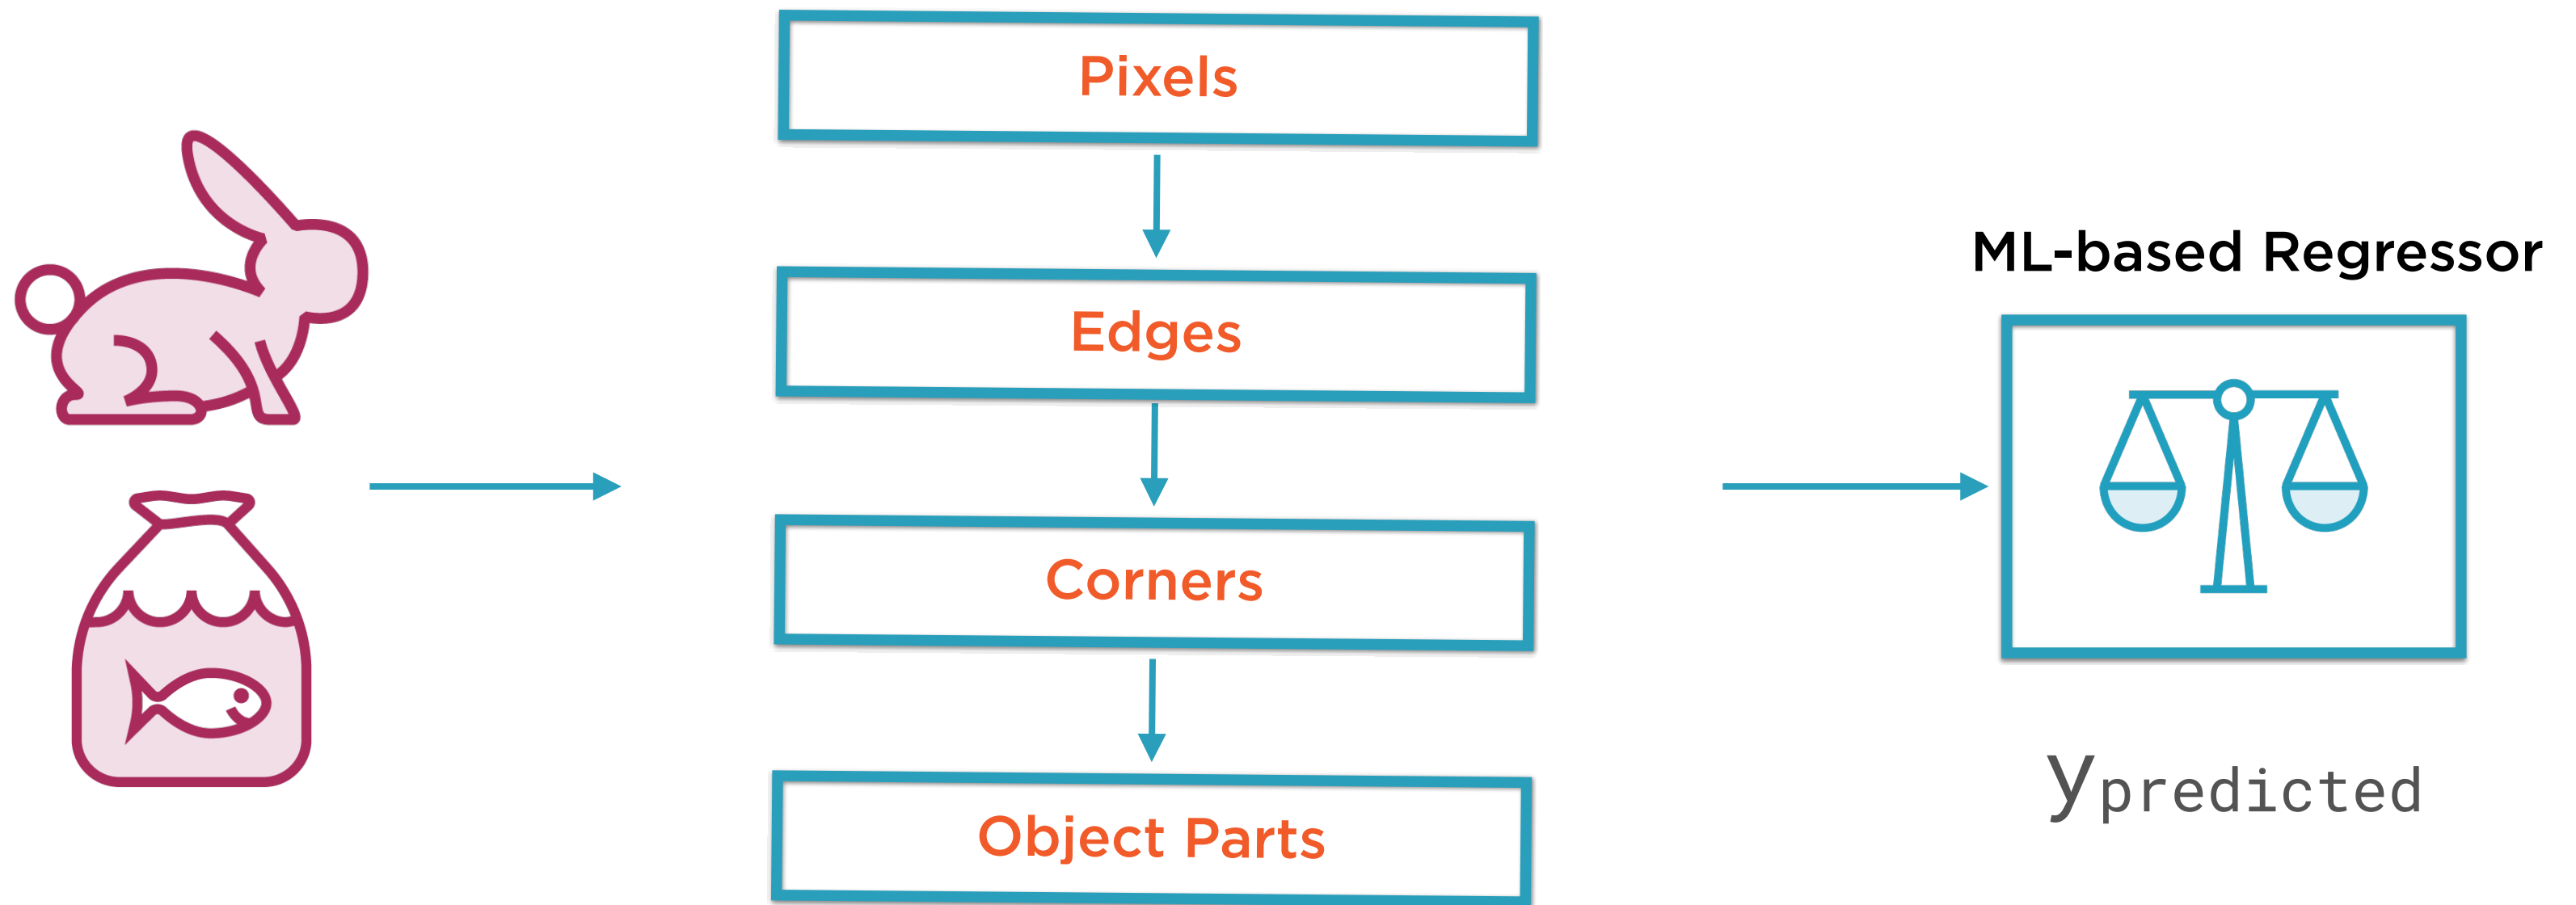
Gradient Descent

Converging on the “best”
value using an
optimization algorithm



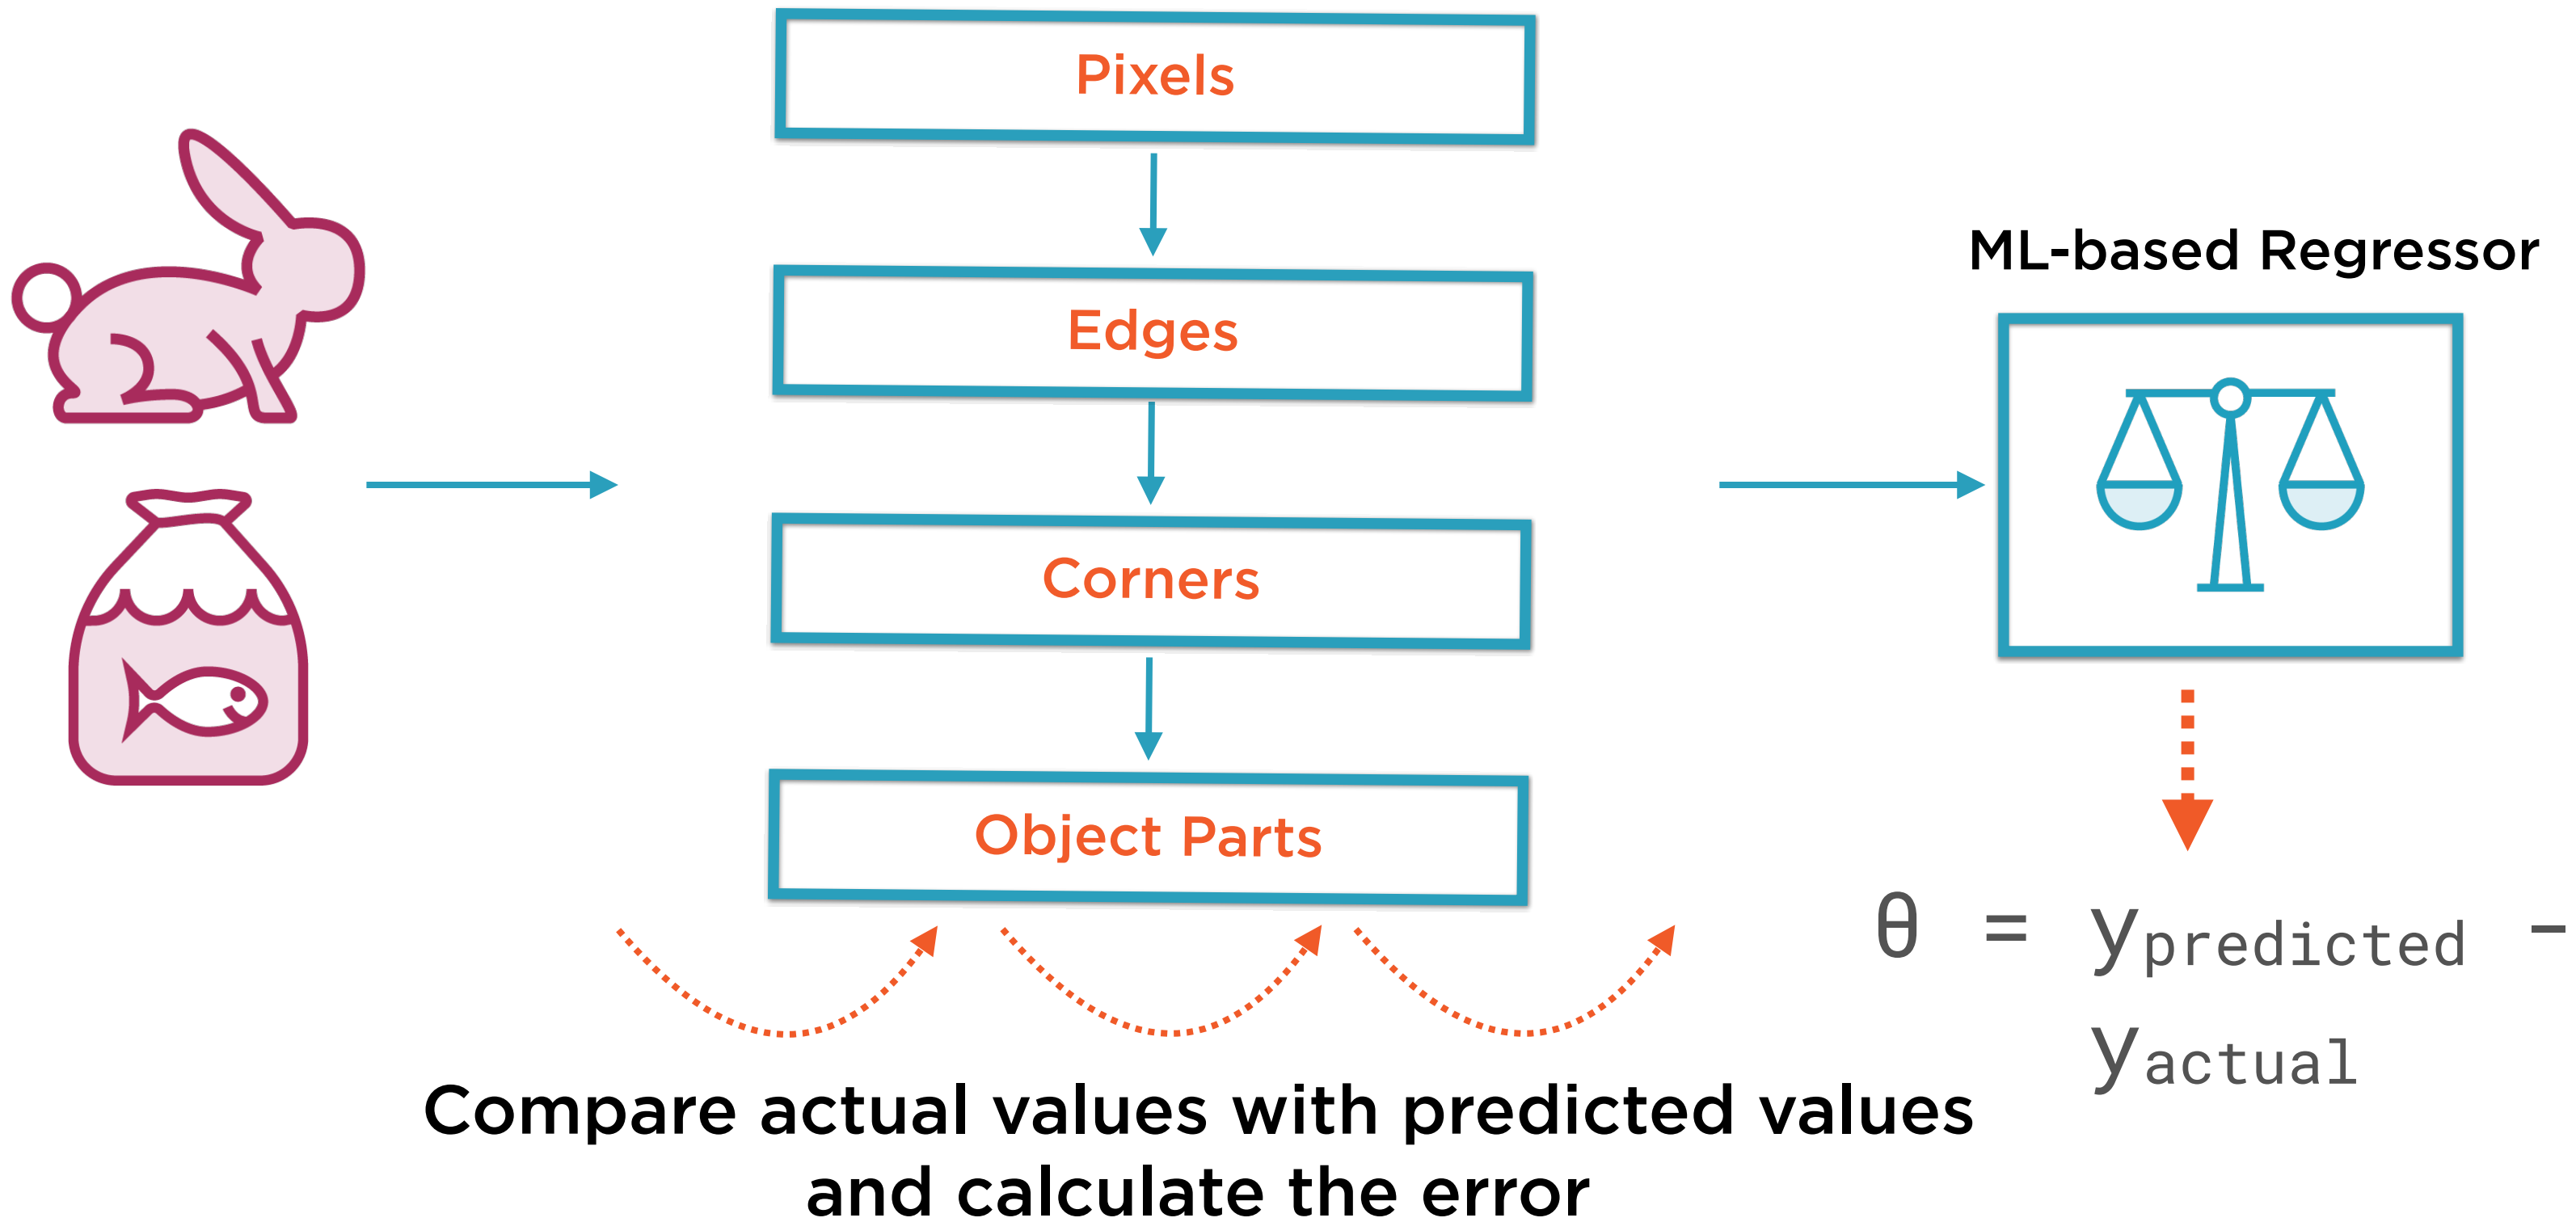
These gradients are used to
update the model parameters

Forward Pass

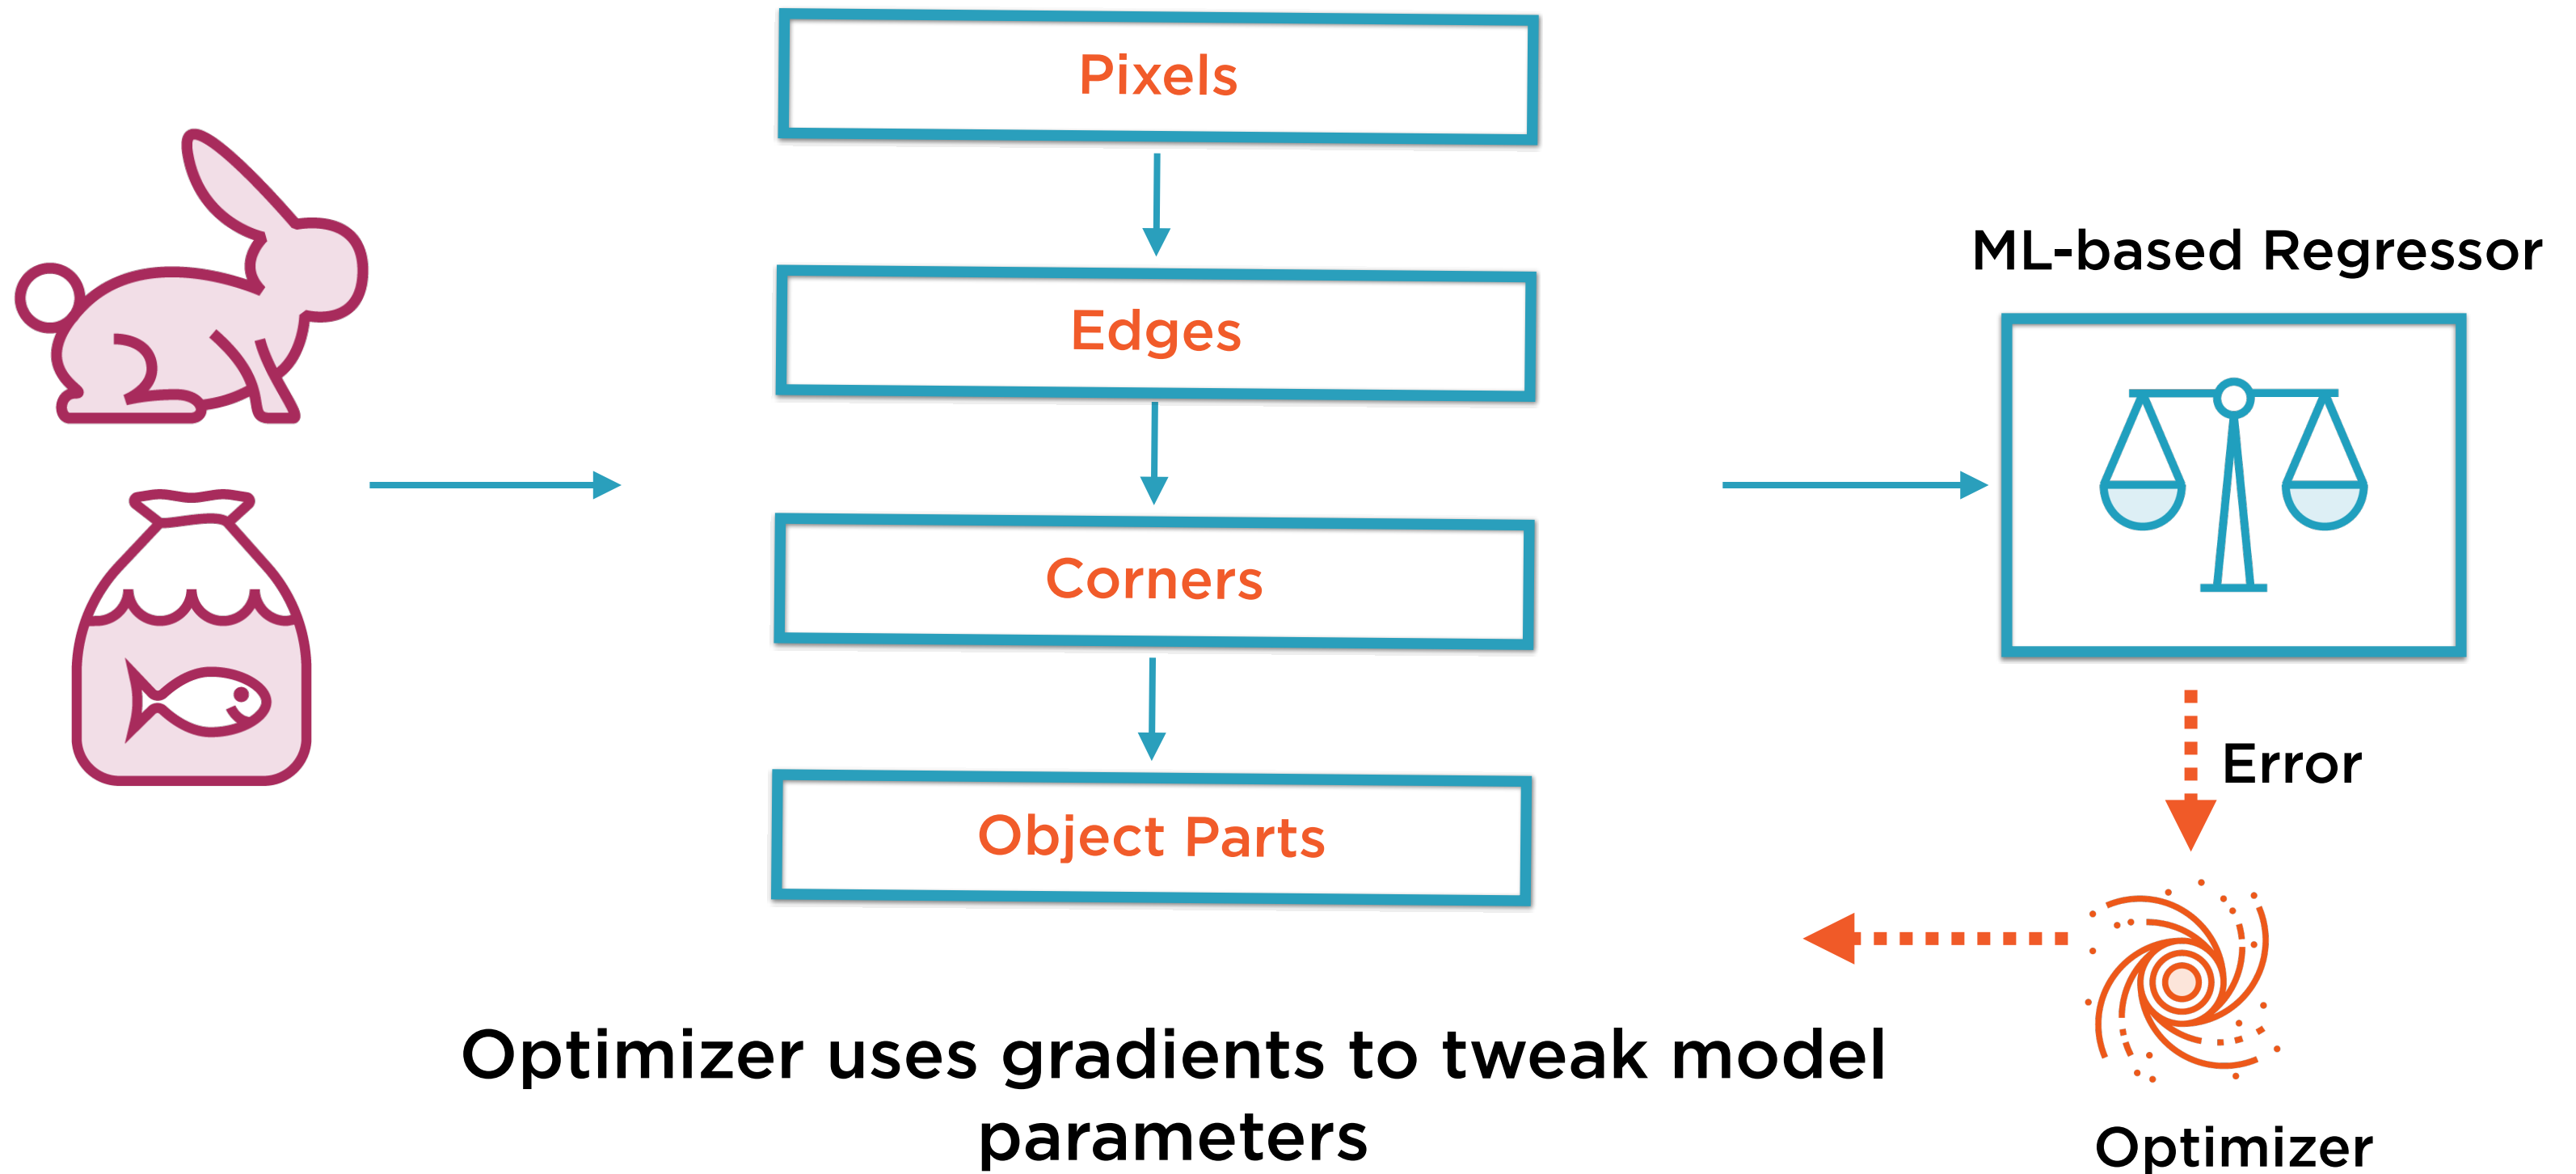


Use the current model weights and biases to
make a prediction


Forward Pass



Optimizer Calculates Gradients



Backward Pass



Pixels




Edges



Corners



Object Parts



ML-based Regressor



Error



Optimiser



PyTorch Optimizers

Linear Regression as an Optimization Problem



Objective Function


Minimize variance of
the residuals (MSE)



Constraints

Express relationship as
a straight line

$$y = Wx + b$$




Decision Variables

Values of W and b

Gradient Descent

Converging on the “best”
value using an
optimization algorithm



Using an Optimizer in PyTorch

Construct Optimizer Object

Pass iterable of all parameters

Each parameter should be a learnable tensor

Compute Gradients

Invoke `.backward()`

Autograd for reverse auto-differentiation

Specify Per-parameter Options

Pass iterable of dict objects

Each key a param, defines parameter group

Take an Optimization Step

Invoke `optimizer.step()`

Overloaded version takes in a closure (advanced)

torch.optim

torch.optim.Optimizer

torch.optim.Adadelta

torch.optim.Adagrad

torch.optim.Adam

Many others

$$\text{Parameters}^{t+1} = \text{Parameters}^t - \text{learning_rate} \times \text{Gradient}(\theta^t)$$

Basic SGD Optimizer

Move each parameter value in the direction of reducing gradient

$$\text{momentum_vec}^t = \text{momentum_coeff} + \text{learning_rate} \times \text{Gradient}(\theta^t)$$
$$\text{Parameters}^{t+1} = \text{Parameters}^t - \text{momentum_vec}$$

Momentum-based Optimizer

Momentum vector helps accelerate in the direction where gradient is decreasing

$$\text{momentum_vec}^t = \text{momentum_coeff} + \text{learning_rate} \times \text{Gradient}(\theta^t)$$
$$\text{Parameters}^{t+1} = \text{Parameters}^t - \text{momentum_vec}$$

Momentum-based Optimizer

Momentum vector helps accelerate in the direction where gradient is decreasing

$\text{momentum_vec}^t = \text{momentum_coeff} + \text{learning_rate} \times \text{Gradient}(\theta^t)$

$\text{Parameters}^{t+1} = \text{Parameters}^t - \text{momentum_vec}$

Momentum-based Optimizer

Momentum vector helps accelerate in the direction where gradient is decreasing

$$\text{momentum_vec}^t = \text{momentum_coeff} + \text{learning_rate} \times \text{Gradient}(\theta^t)$$
$$\text{Parameters}^{t+1} = \text{Parameters}^t - \text{momentum_vec}$$

Momentum-based Optimizer

Gradients at each step weighted by those in previous step

Benefit: Faster convergence

$$\text{momentum_vec}^t = \text{momentum_coeff} + \text{learning_rate} \times \text{Gradient}(\theta^t)$$

$$\text{Parameters}^{t+1} = \text{Parameters}^t - \text{momentum_vec}$$

Momentum-based Optimizer

Need a **momentum coefficient**, between 0 and 1 to prevent overshooting

Advanced Optimizers

Many variants of optimizers

Increasing complexity

More hyperparameters

Better performance

Adam ~ Adaptive Moment Estimation

Demo


**Regression Using NN Layers and
Optimizer**

Dropout to Mitigate Overfitting

Preventing Overfitting



Regularisation - Penalise complex models



Cross-validation - Distinct training and validation phases



Dropout - Intentionally turn off some neurons during training

Dropout



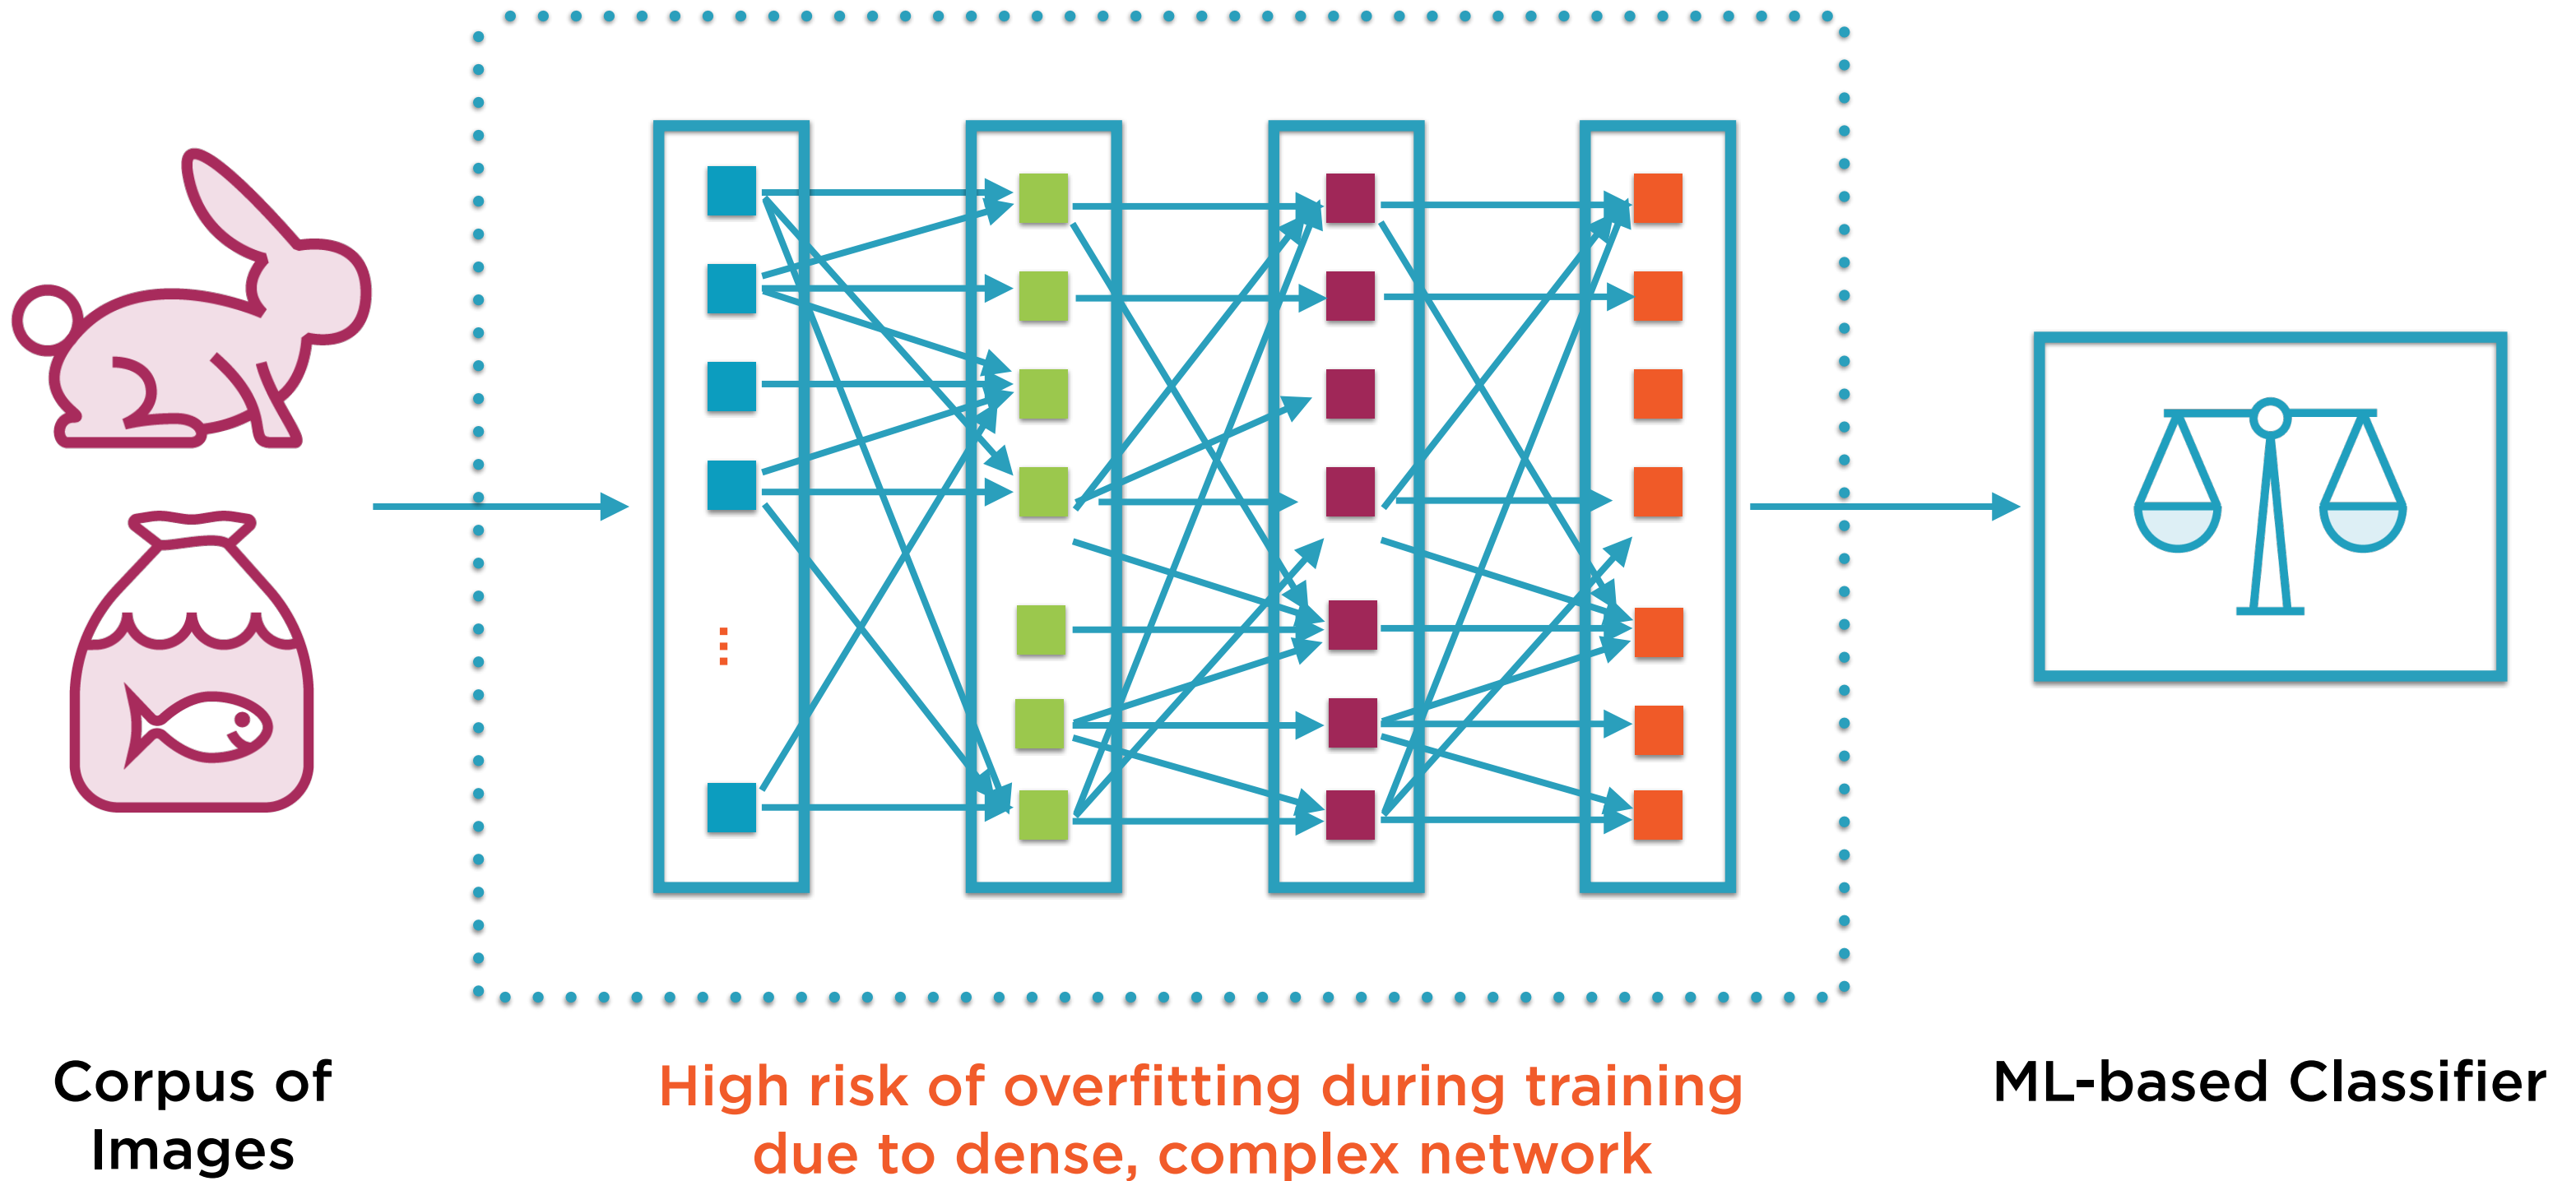
Specify a fraction of neurons that will stay off in each training step

“Dropout” neurons chosen at random


Different neurons off in each training step

In effect, each training step builds different network configuration


Densely Connected Neural Network



Dropout = 50%



Corpus of
Images




Randomly switch off say 50% of neurons
in each training step




ML-based Classifier

Dropout = 50%



Corpus of
Images




Training forced to rely on a much
simpler neural network




ML-based Classifier

Dropout = 50%



Corpus of
Images




Each training step will build a different
configuration




ML-based Classifier

Dropout = 50%



Corpus of
Images

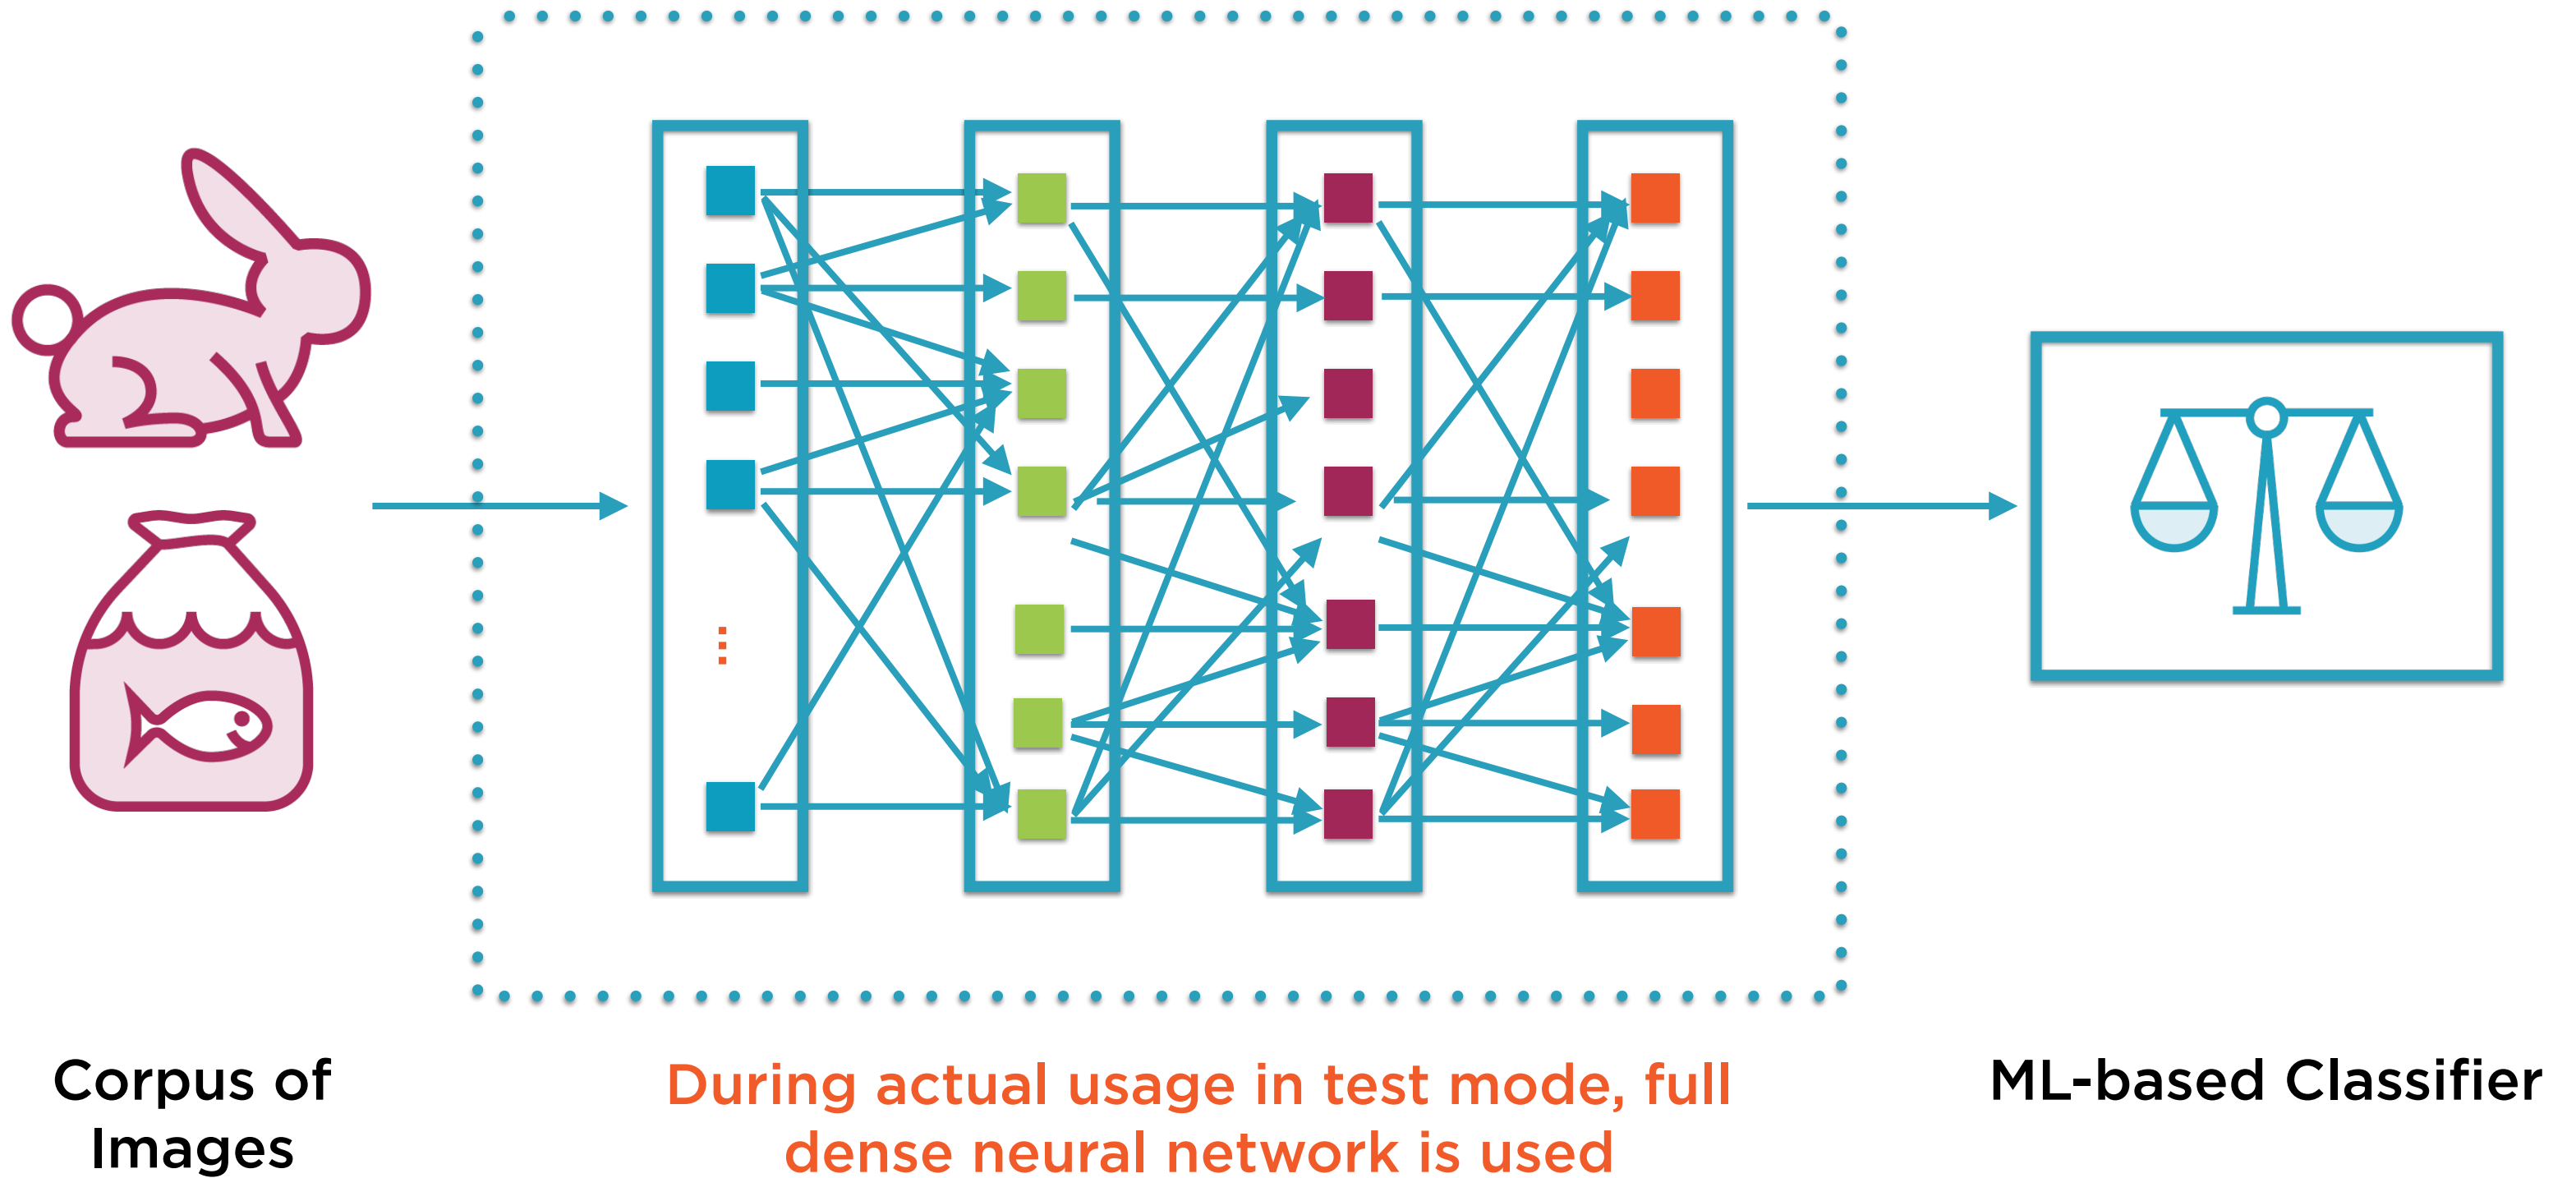


Each training step will build a different
configuration




ML-based Classifier

Dropout During Training Only

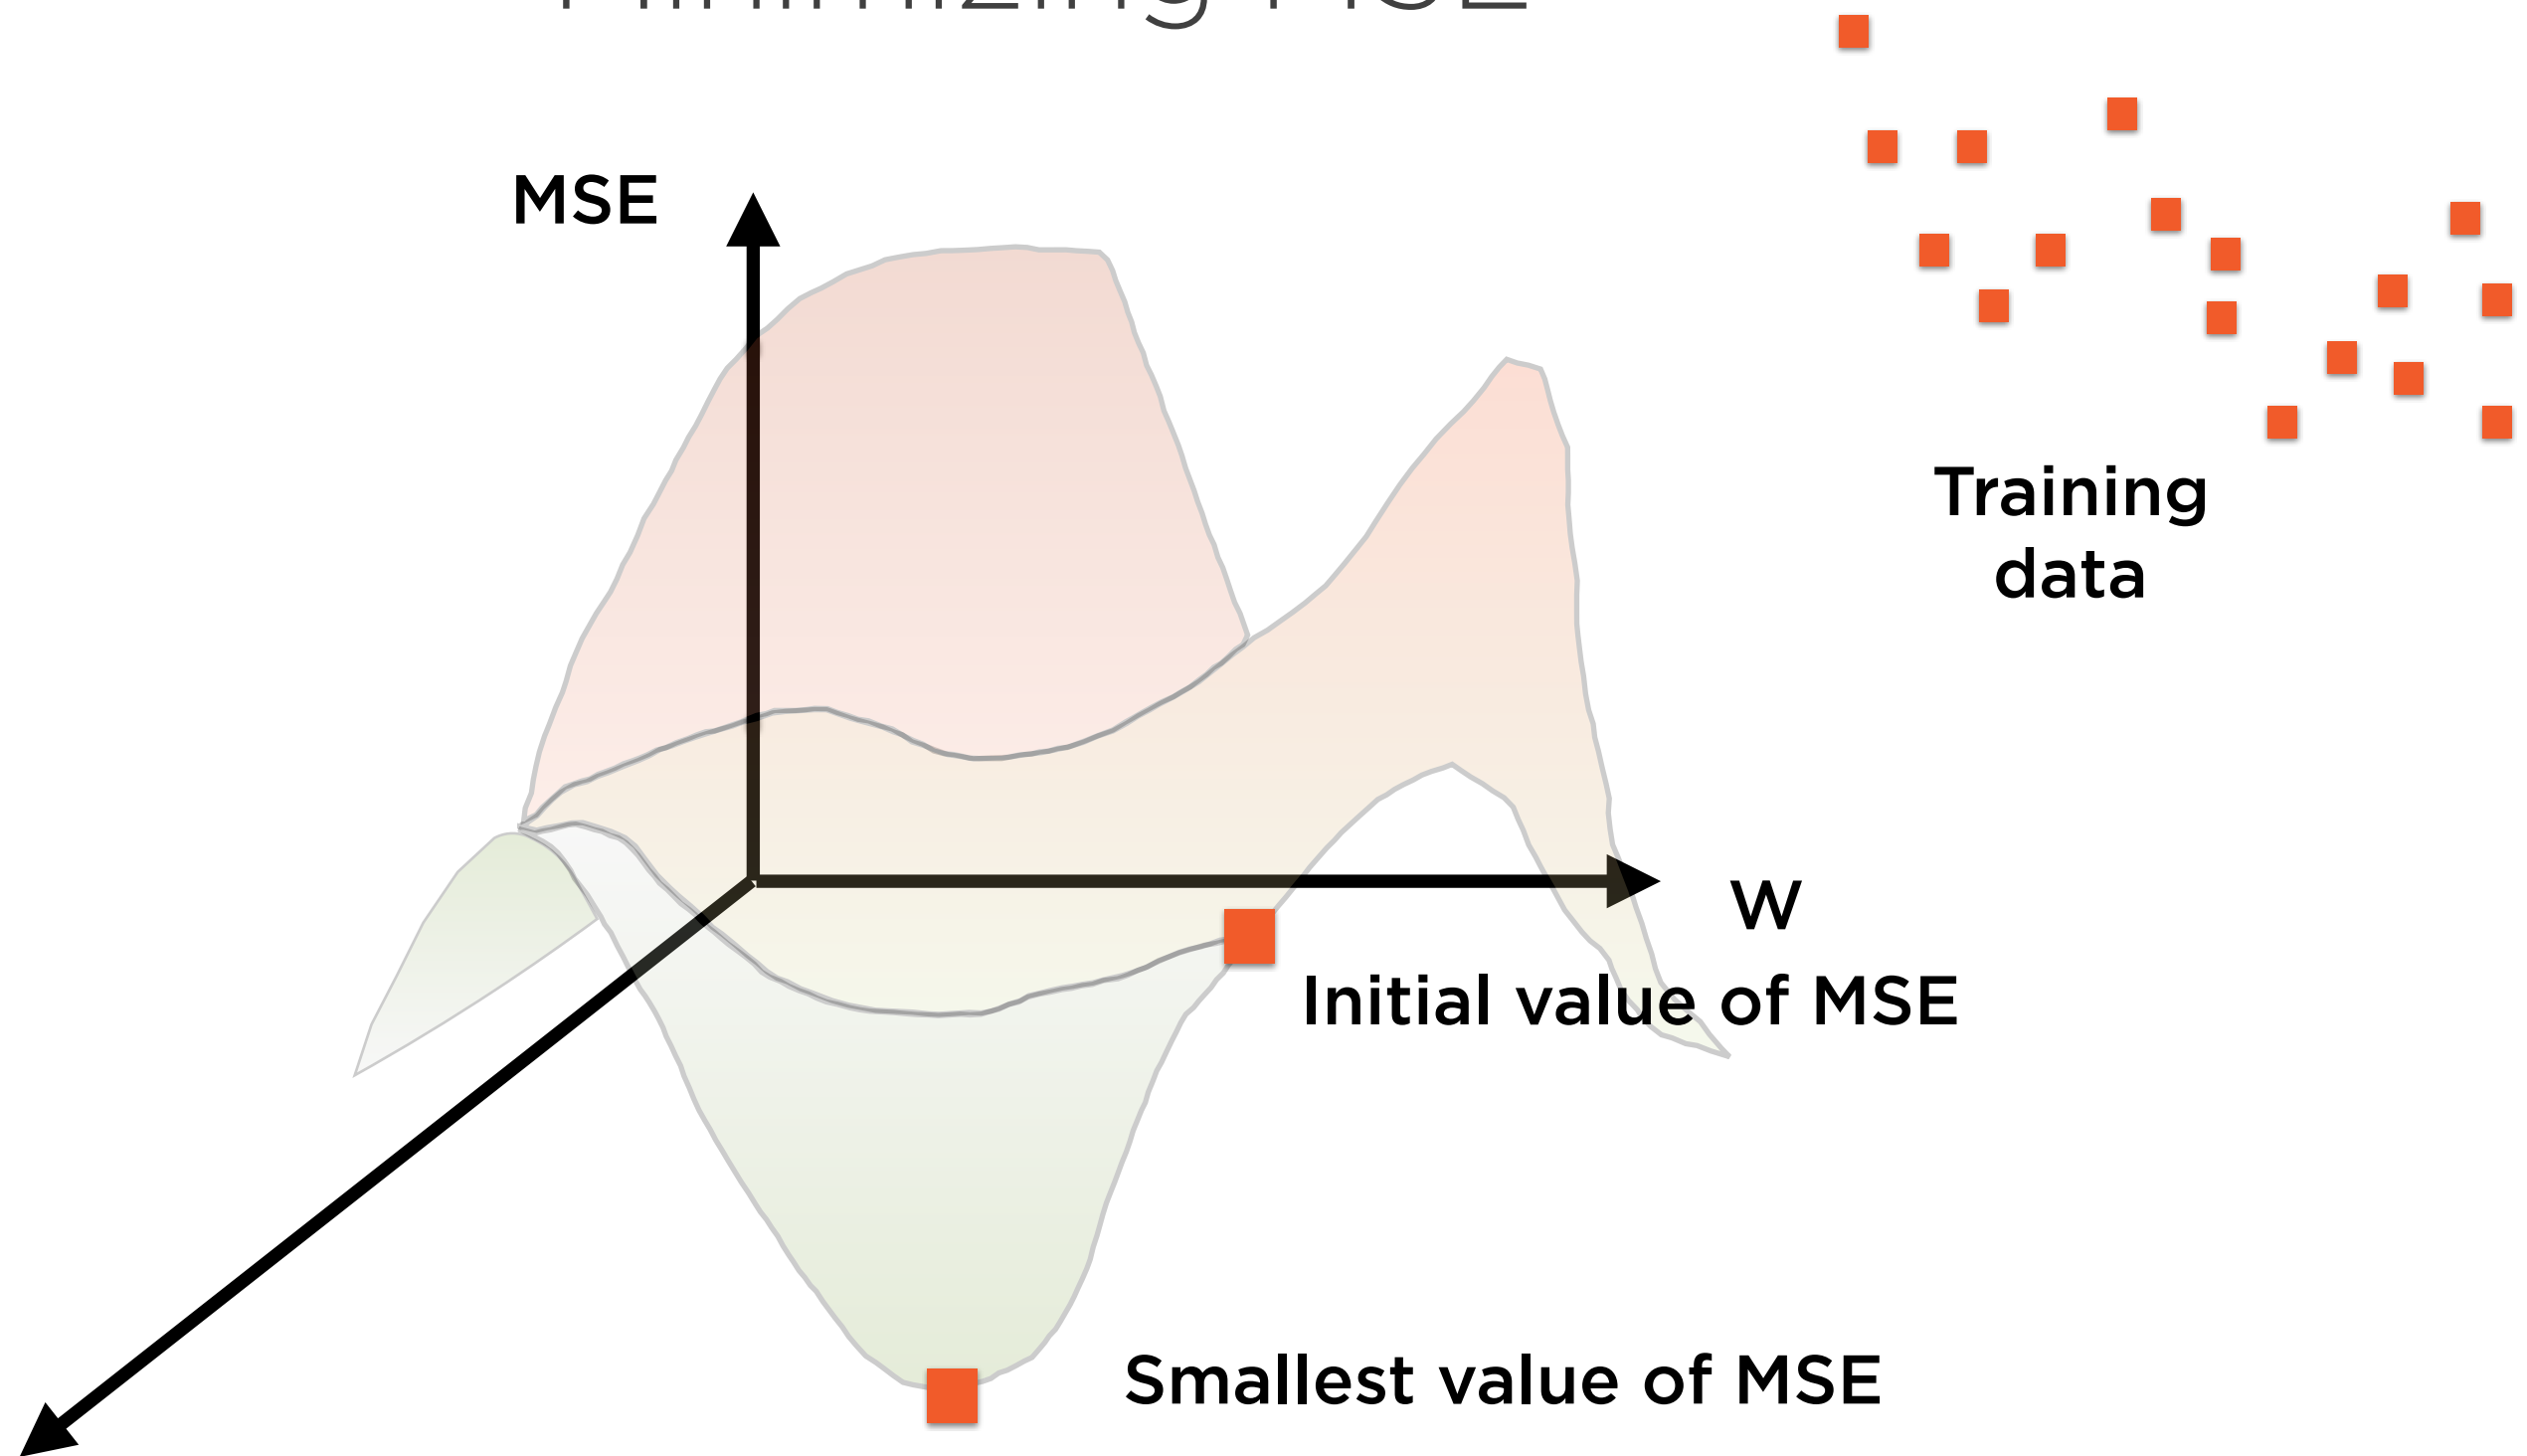


Gradient Descent and Batches

Regression: The Simplest Neural Network



Minimizing MSE



“Epoch”

MSE


1 pass through training
data = 1 epoch

Training
data


W

Initial value of MSE

Smallest value of MSE




Subset of Data in Each Iteration




Epoch ~ Multiple Batches

1 pass through training
data = 1 epoch




“Batch Size”

Number of training data
points considered =
batch size




“Batch Size”




**Stochastic Gradient
Descent**

1 point at a time



**Mini-batch Gradient
Descent**

Some subset in each
batch



**Batch Gradient
Descent**

All training data in
each batch

Demo

**Building a regression model for
demand prediction**

Summary

Training a neural network using forward and backward passes

Using optimizers to update model parameters

Using layers and activation functions to train and build neural networks

Understanding and using dropout to mitigate overfitting